

GRADUATION THESIS DEFENSE GUIDE

GRU-Based Handover Optimization in Open RAN (O-RAN) Systems

Student	Omar Farouk
Project	Open RAN Digital Twin Platform
Component	GRU Neural Network xApp + NS-3 Simulation
Team	Omar Farouk (xApp/Sim) + Fares (GRU Model Training)
Date	May 06, 2026
Institution	Faculty of Engineering
Pages	~75 pages Comprehensive Defense Reference

Abstract

This guide is your complete, exam-ready reference for defending your thesis on GRU-based handover optimization in Open RAN systems. Every concept is explained in simple language with analogies, every parameter is listed with its value and the reason it was chosen, and over 60 defense questions are answered in full detail. The guide covers: NS-3 network simulation, FlexRIC near-RT RIC, the C xApp (best2.c), the Python GRU inference service (gru_xapp.py), the 2D and 3D monitoring GUIs, the system orchestration layer, and a complete analysis of all simulation results.

What is O-RAN? — The Big Picture

1.1 The Problem O-RAN Was Created to Solve

Before O-RAN existed, mobile networks were built using equipment from a single vendor — for example, Ericsson or Nokia — and all the software inside that equipment was **closed and proprietary** (meaning: locked, secret, not shareable). This is called **vendor lock-in**. Imagine buying a smartphone where you can only install apps made by the phone manufacturer itself. If you want a better camera app, you cannot use it — you are stuck with whatever the manufacturer provides. That is exactly what happened to mobile operators: they could not mix equipment from different vendors, could not write their own optimization software, and had to pay whatever the vendor charged for upgrades.

The **O-RAN Alliance** (Open Radio Access Network Alliance), founded in 2018 by major operators including AT&T, Deutsche Telekom, NTT DoCoMo, Orange, and China Mobile, was created specifically to break this lock-in. Their solution: define open interfaces so that components from different vendors can talk to each other, and allow software (called xApps and rApps) to be written by anyone — universities, startups, operators themselves.

Note: Think of O-RAN as the Android of mobile networks. Just like Android opened smartphones to any developer, O-RAN opens the radio network to any software developer or researcher.

1.2 Traditional RAN vs O-RAN — Side-by-Side Comparison

Aspect	Traditional RAN (Closed)	O-RAN (Open)
Vendor	Single vendor (e.g., Nokia only)	Multi-vendor — mix any brands
Software	Proprietary, closed source	Open APIs, open interfaces
Optimization	Built-in, cannot be changed	Custom apps (xApps/rApps)
Intelligence	Fixed algorithms in hardware	AI/ML at RIC level
Cost	High (vendor controls pricing)	Lower (competition, open-source)
Innovation speed	Slow (wait for vendor release)	Fast (anyone can develop)
Interface standard	Proprietary (varies by vendor)	Standardized (O-RAN specs)
Radio Unit (RU)	Integrated, vendor-specific	O-RU — open spec
Baseband (DU/CU)	Single box (BBU)	Split into O-DU + O-CU
Control loop	ms range, inside hardware	10ms–1s (near-RT RIC) or >1s

1.3 The O-RAN Architecture — All the Components Explained

O-RAN splits a traditional base station (called gNB in 5G) into several parts that can run on standard servers and communicate over open interfaces. Here is what each part does:

- **O-RU (Open Radio Unit):** The physical antenna hardware. It converts digital signals to radio waves and back. This is the only part that must be near the antenna tower.
- **O-DU (Open Distributed Unit):** Handles the real-time radio processing — scheduling which UE gets radio resources in each 0.5ms slot. Runs on a standard server near the O-RU.
- **O-CU (Open Central Unit):** Handles higher-layer protocols (RRC, PDCP). Less time-critical, can run farther away in a data center.
- **Near-RT RIC (Near Real-Time RAN Intelligent Controller):** A software platform that runs xApps. It receives reports from O-DU/O-CU every 10ms–1s and can send control commands back. This is where FlexRIC runs in our system.
- **Non-RT RIC (Non-Real-Time RIC):** Runs rApps for long-term optimization (>1 second). Trains ML models, sets policies. Communicates with near-RT RIC via the A1 interface.
- **SMO (Service Management and Orchestration):** The overall management layer — deploys xApps, manages configuration, collects logs.

1.4 The O-RAN Interfaces — All 7 Explained Simply

Interface	Connects	Purpose	Timing
O1	SMO ↔ O-RAN nodes	Management: configure, monitor, fault	Minutes
O2	SMO ↔ Cloud/Infrastructure	Deploy software, manage compute resources	Minutes
A1	Non-RT RIC ↔ Near-RT RIC	Send AI policies from non-RT to near-RT	Seconds
E2	Near-RT RIC ↔ E2 Node (gNB)	KPM reports + RC control commands	10ms–1s
F1	O-CU ↔ O-DU	Internal gNB split interface	ms range
X2	gNB ↔ gNB (LTE legacy)	Inter-cell handover coordination (4G)	ms range
Xn	gNB ↔ gNB (5G)	Inter-cell handover coordination (5G)	ms range
Note: In our thesis, the most important interface is E2 . This is the channel through which NS-3 sends SINR/RSRP reports to FlexRIC (KPM) and FlexRIC sends handover commands back to NS-3 (RC).			

1.5 The Three RIC Types and Their Timing

RIC Type	Latency Window	Runs What	Our System
Non-RT RIC	> 1 second	rApps — train models, set policies	Not used directly

RIC Type	Latency Window	Runs What	Our System
Near-RT RIC	10 ms – 1 second	xApps — monitor KPMs, send RC commands	FlexRIC — runs our xApp

RT (embedded)	< 10 ms	Scheduler inside O-DU	NS-3 internal scheduler
---------------	---------	-----------------------	-------------------------

The word '**near-real-time**' means the system reacts within 10 milliseconds to 1 second. This is fast enough to make handover decisions (which take tens of milliseconds to execute) but not fast enough for slot-level scheduling (which needs sub-millisecond response). Our xApp evaluates each UE every **0.05 simulation-seconds**, which comfortably falls within the near-RT window.

Near-RT RIC and xApps — Deep Dive

2.1 What is FlexRIC?

FlexRIC is an open-source near-RT RIC developed by EURECOM (a French research institute). It implements the O-RAN E2 interface and provides a framework for writing xApps in C, C++, or Python. In our system, FlexRIC is the 'brain' that sits between NS-3 (the network simulator) and our GRU xApp. It receives raw measurements from NS-3, passes them to the xApp, and relays the xApp's handover decisions back to NS-3.

FlexRIC runs as a **background daemon** (meaning: a background process with no visible terminal output). All its log messages go to `/tmp/flexric.log` — this is by design, not a bug. If you want to see what FlexRIC is doing, you must run: `tail -f /tmp/flexric.log`

Important: FlexRIC ALWAYS logs to /tmp/flexric.log. Never expect output on the terminal. If you do not see activity there, FlexRIC has not started or NS-3 has not connected.

2.2 What is an xApp?

An **xApp** (pronounced 'ex-app') is a small application that runs inside the near-RT RIC and performs a specific network function. The 'x' stands for 'any' — there can be any number of xApps running simultaneously, each responsible for a different task (handover optimization, interference management, load balancing, etc.).

Think of xApps like browser extensions. Firefox (the RIC) runs independently, but you can install extensions (xApps) that add new capabilities — an ad blocker, a password manager, etc. Each extension has access to browser events (like KPM reports) and can influence behavior (like sending RC commands). Our xApp is called `xapp_handover_gru` and is built from the source file `best2.c`.

2.3 xApp Lifecycle — Step by Step

1. **xApp starts** → connects to FlexRIC via a local socket
2. **xApp subscribes** → sends E2SM-KPM subscription request: 'please send me SINR/RSRP every 0.05s'
3. **FlexRIC forwards subscription** → sends subscription to NS-3 via E2 interface
4. **NS-3 acknowledges** → confirms it will send KPM reports at the requested interval
5. **Reports arrive** → every 0.05 sim-seconds, NS-3 sends a KPM Indication message
6. **xApp processes report** → checks A3 condition, calls GRU service, decides handover
7. **xApp sends RC command** → if handover needed, sends E2SM-RC control to NS-3
8. **NS-3 executes HO** → UE connects to new cell

Important: The xApp MUST start BEFORE NS-3. If NS-3 starts first, it will wait for E2 connection. If the xApp is not ready, the subscription step (step 2-4) is missed and no KPM reports will ever be processed. This is why gru.sh starts FlexRIC+xApp first, then waits 3 seconds, then starts NS-3.

2.4 E2 Interface: KPM and RC Service Models

The E2 interface carries two types of messages in our system:

E2SM-KPM (Key Performance Measurement) — Monitoring

KPM is the 'listening' part. The xApp subscribes to receive measurement reports. Each KPM report contains a snapshot of the network state for all UEs. The fields in each KPM report that our xApp uses:

Field	Meaning	Unit	Used For
UE ID (RNTI)	Unique ID of each UE in the cell	Integer	Identify which UE the measurement belongs to
SINR	Signal-to-Interference-plus-Noise Ratio	dB	Primary metric for signal quality
RSRP	Reference Signal Received Power	dBm	Signal strength (power level)
Serving Cell ID	Which cell the UE is currently connected to	Integer 0-6	Know current attachment
Neighbor Cell SINRs	SINR from all other 6 cells	dB (array)	Find best alternative cell
Timestamp	Simulation time of measurement	Seconds	Order samples for rolling window

E2SM-RC (RAN Control) — Control

RC is the 'commanding' part. After deciding a handover is needed, the xApp sends an RC control message to NS-3 with the target cell ID. NS-3 then executes the handover procedure (RRC Reconfiguration), which involves:

1. NS-3 sends RRC Reconfiguration message to the UE
2. UE detaches from serving cell
3. UE attaches to target cell
4. Target cell sends RRC Reconfiguration Complete
5. Handover is recorded in the simulation log

2.5 Why 0.05 Simulation-Seconds per KPM Report?

The KPM report interval of **0.05 simulation-seconds** means the xApp receives **20 reports per simulation-second** per UE. This interval was chosen because:

- It provides enough temporal resolution to detect fast-changing channel conditions in mmWave
- The GRU rolling window needs 10 samples — at 0.05s per sample, the window covers 0.5 sim-seconds of history
- It matches the O-RAN near-RT RIC timing requirement (10ms–1s per action)
- Faster would increase computational load; slower would miss channel changes

NS-3 Network Simulator — How We Build the Virtual Network

3.1 What is NS-3 and Why Use It?

NS-3 is a free, open-source **discrete-event network simulator** used by thousands of researchers worldwide. 'Discrete-event' means the simulation jumps from one event to the next (a packet arrives, a timer fires, a UE moves) rather than running continuously. NS-3 is written in C++ and can simulate everything from simple packet networks to complex 5G mmWave systems.

We use NS-3 instead of a real network because:

- A real 7-cell mmWave network would cost millions of dollars to deploy
- NS-3 lets us control every parameter exactly (cell positions, UE speeds, channel model)
- NS-3 can run faster or slower than real-time — useful for experiments
- Results are repeatable — same random seed gives same result
- NS-3 has a built-in O-RAN E2 interface module that connects to FlexRIC

3.2 mmWave — What Is It and Why Use It for 5G?

mmWave stands for **millimeter wave**. These are radio frequencies above 24 GHz (traditional 4G LTE uses frequencies below 6 GHz). The millimeter refers to the physical wavelength — signals at 28 GHz have a wavelength of about 10mm, hence the name.

Property	mmWave (>24 GHz)	Sub-6 GHz (4G/5G)
Bandwidth available	Very large (400 MHz+)	Limited (20–100 MHz)
Data rate	Multi-Gbps possible	Hundreds of Mbps
Range	Short (~200m typical)	Long (km range)
Penetration	Poor (blocked by walls/rain)	Good
Handover frequency	High (UE exits coverage fast)	Low
Our simulation freq.	28 GHz (ns3-mmwave module)	N/A

The short range and poor penetration of mmWave is precisely WHY handover optimization is critical — UEs move in and out of mmWave coverage far more frequently than in 4G, making ping-pong handovers much more likely.

3.3 Our Simulation Setup — `gru_scenario.cc`

The main simulation script is at: `/home/omar_farouk/open-ran-clean/yousef_fathy/ns-0-RAN-flexric/mmwave-LENA-oran/gru_scenario.cc`

Cell Layout

The simulation uses **7 mmWave base stations (gNBs)** arranged in a grid or cluster pattern. Cell IDs are numbered 0 through 6. Each cell covers roughly a circular area, and cells overlap at their edges — this overlap zone is where UEs experience both cells' signals simultaneously and handovers are most likely.

Parameter	Value	Explanation
Number of cells (gNBs)	7	One central + 6 surrounding (hexagonal-like layout)
Number of UEs	20	20 user equipment nodes moving randomly
Carrier frequency	28 GHz	mmWave band — typical 5G NR mmWave frequency
Bandwidth	400 MHz	Full mmWave channel bandwidth
Simulation time (sim006/010)	60 seconds	60 sim-seconds = ~3 real hours to compute
Simulation time (sim011)	120 seconds	120 sim-seconds = ~6 real hours
KPM report interval	0.05 sim-seconds	20 reports/second per UE
E2 port	36421 (SCTP)	Standard O-RAN E2 interface port
UE mobility model	Random Walk 2D	UEs move in random directions, bounded area
UE speed	~3 m/s average	Pedestrian speed

3.4 How NS-3 Connects to FlexRIC — The E2 Interface

NS-3 includes an **E2 termination module** that implements the O-RAN E2 Application Protocol (E2AP). When the simulation starts, this module:

1. Opens a **SCTP connection** to FlexRIC at **127.0.0.1:36421** (loopback, same machine)
2. Sends an E2 Setup Request — introduces itself as an E2 Node
3. FlexRIC responds with E2 Setup Response — connection established
4. FlexRIC forwards the xApp's KPM subscription to NS-3
5. NS-3 starts generating KPM Indication messages every 0.05 sim-seconds

Note: *SCTP (Stream Control Transmission Protocol) is like TCP but designed for telecom signaling — it supports multiple streams and has better reliability properties for control-plane messages. Port 36421 is the IANA-registered port for S1AP/E2AP protocols.*

3.5 The A3 Event — How Traditional Handover Trigger Works

The **A3 event** is the standard 3GPP handover trigger. It fires when a neighboring cell's signal becomes stronger than the serving cell by a defined margin. In our xApp, we implement the A3 condition as follows:

A3 condition: $\text{neighbor_SINR} > \text{serving_SINR} + \text{A3_OFFSET} + \text{A3_HYSTERESIS}$
In numbers: $\text{neighbor_SINR} > \text{serving_SINR} + 1.0 \text{ dB} + 1.0 \text{ dB} = \text{serving_SINR} + 2.0 \text{ dB}$

Parameter	Value	Purpose
A3_OFFSET	1.0 dB	Minimum gain required — prevents HO for tiny improvements
A3_HYSTERESIS	1.0 dB	Extra margin — prevents ping-pong from noise fluctuations
Combined threshold	2.0 dB	Neighbor must be 2 dB better than serving cell to trigger A3

The A3 event alone (without GRU) would trigger a handover every time any neighbor exceeds the threshold. Our system uses A3 as a **gate** — it must be true before the GRU is even consulted. If A3 is not triggered, no handover evaluation happens. This prevents unnecessary GRU calls and avoids handovers when the UE is stable.

3.6 Simulation Pace — Why It Takes Hours

The simulation runs at approximately **0.005 simulation-seconds per real-second**. This means:

$60 \text{ sim-seconds} / 0.005 = 12,000 \text{ real-seconds} = 200 \text{ real-minutes} \approx 3.3 \text{ hours}$

This is not a bug — it is the computational cost of accurately simulating:

- 28 GHz mmWave channel propagation with ray-tracing-level accuracy
- 20 UEs each generating and receiving traffic simultaneously
- E2 message encoding/decoding (ASN.1 format)
- HTTP calls to the GRU Python service for each UE evaluation
- InfluxDB writes via `sim_data_pusher.py` running in parallel

GRU Neural Network — Theory Made Simple

4.1 What is a Neural Network? (Simple Analogy)

Imagine you are learning to recognize whether a fruit is an apple or an orange. At first you are bad at it. But every time someone tells you 'that was wrong,' you adjust your judgment slightly. After seeing thousands of fruits, you become very accurate. A **neural network** works exactly the same way.

A neural network is a mathematical system made of **layers of nodes** (called neurons). Each neuron receives numbers as input, multiplies them by learned **weights**, adds them up, and passes the result through an **activation function** (a mathematical curve that squashes the output to a useful range). The network is trained by adjusting these weights using real examples until predictions are accurate.

***Note:** Analogy: Think of each neuron as a person in a decision committee. Each person has an opinion (weight). They discuss (combine inputs), and the committee reaches a decision (output). Training = giving the committee feedback until their decisions match reality.*

4.2 Why Regular Neural Networks Fail for Time Series

Our task is to predict the best cell for a UE based on its **history of SINR measurements**. This is a time-series problem — the order of measurements matters. A UE moving toward Cell 3 will show rising SINR from Cell 3 over the last 10 samples. A regular feedforward neural network treats each sample independently — it has no memory of previous samples. It cannot detect trends.

4.3 RNN and the Vanishing Gradient Problem

A **Recurrent Neural Network (RNN)** has memory — each step passes a hidden state to the next step, like a person remembering the last thing they heard. However, RNNs suffer from the **vanishing gradient problem** (meaning: when training on long sequences, the mathematical signal used to update weights — called the gradient — becomes extremely tiny and the network effectively forgets the earliest parts of the sequence). For sequences longer than ~10 steps, vanilla RNNs fail badly.

4.4 LSTM and GRU — The Solutions

LSTM (Long Short-Term Memory) solved the vanishing gradient problem in 1997 by introducing **gates** — mathematical switches that control what information to remember and what to forget. **GRU (Gated Recurrent Unit)** was introduced in 2014 as a simpler version of LSTM with fewer parameters but similar performance.

Aspect	Vanilla RNN	LSTM	GRU (Our Choice)
Gates	None	3 gates (input, forget, output)	2 gates (update, reset)

Aspect	Vanilla RNN	LSTM	GRU (Our Choice)
Memory capacity	Short sequences only	Long sequences	Long sequences
Parameters	Fewest	Most	Middle (fewer than LSTM)
Training speed	Fast	Slow	Faster than LSTM
Vanishing gradient	Yes (problem)	Solved	Solved
Best for	Very short sequences	Very long, complex sequences	Moderate sequences (like ours)

4.5 How GRU Works — Plain English

A GRU cell has two gates. Think of them as two employees managing a memory file:

- **Update Gate (z):** Decides how much of the OLD memory to keep versus how much to update with new information. If $z=1$, keep everything old. If $z=0$, replace entirely with new. Analogy: A secretary who decides whether to update the filing cabinet or keep the old files.
- **Reset Gate (r):** Decides how much of the OLD memory is relevant to processing the CURRENT input. If $r=0$, ignore the old memory entirely. Analogy: A researcher who decides whether past notes are relevant to today's question.

At each time step (each 0.05s KPM report), the GRU cell:

1. Receives new SINR/RSRP/velocity measurements for one UE
2. Reset gate decides: how relevant is the previous hidden state?
3. Update gate decides: how much to blend old hidden state with new candidate
4. New hidden state is computed — it encodes the UE's recent movement pattern
5. After all 10 samples are processed, the final hidden state feeds a Dense layer
6. Dense layer + softmax outputs probabilities for each of the 7 cells
7. Cell with highest probability is returned as the recommendation

4.6 Our GRU Model Specifics

Parameter	Value	Explanation
Input window size	10 samples	Last 10 KPM reports = last 0.5 sim-seconds of history
Input features per sample	SINR + RSRP + velocity + serving_cell_id + neighbor SINRs (7)	~10 features per timestep

Parameter	Value	Explanation
GRU hidden units	Model-specific	Determines memory capacity — more units = more capacity
Output classes	7	One probability per cell (cells 0-6)
Output activation	Softmax	Converts raw scores to probabilities that sum to 1.0
Loss function	Categorical cross-entropy	Standard for multi-class classification
Optimizer	Adam	Adaptive learning rate — standard for RNN training
Model file format	.h5 or .keras	Keras/TensorFlow saved model format

4.7 Rolling Window Mechanism

The xApp maintains a **rolling window** (a list of the last 10 samples) per UE. Think of it as a 10-slot conveyor belt:

1. When a new KPM report arrives for UE #5, its SINR/RSRP values are added to the belt
2. If the belt already has 10 items, the oldest one is dropped
3. The 10-item belt is sent to the GRU service as one prediction request
4. The GRU processes all 10 items in sequence and returns the best cell ID

Note: If the belt has fewer than 10 items (UE just started), the GRU service returns a 'not enough data' response and the xApp falls back to pure A3 decision.

The Handover Problem in 5G mmWave Networks

5.1 Why Handovers Are Needed

A **handover** (also called handoff in American English) is the process of transferring an active connection from one base station (cell) to another without dropping the call or data session. Think of it like a relay race — the baton (connection) must pass smoothly from one runner (cell) to the next as the UE (user) moves away from the first cell.

Handovers are needed because:

- Radio signals weaken with distance (path loss)
- Obstacles (buildings, trees) block mmWave signals easily
- A UE that does not hand over will experience degrading SINR until connection drops
- There are 7 cells in our simulation — UEs are always near the boundary of 2+ cells

5.2 Ping-Pong Handover — The Main Problem

A **ping-pong handover** (PP) is when a UE hands over from Cell A to Cell B, then immediately hands back from Cell B to Cell A — bouncing back and forth like a ping-pong ball. This wastes network resources, increases signaling load, and can cause brief connection interruptions.

In our system, a ping-pong is defined as: **a UE performs HO from cell A to cell B, then performs HO from cell B back to cell A within 5 simulation-seconds.**

What causes ping-pong?

- **Cell-edge location:** UE is exactly between two cells — both cells have similar SINR
- **Fast movement:** UE moves quickly in and out of a cell's coverage area
- **Channel fluctuation:** mmWave SINR varies rapidly due to small-scale fading
- **Insufficient hysteresis:** If the threshold margin is too small, noise triggers HO

5.3 Anti-Ping-Pong Cooldown Mechanism

Our xApp implements a **cooldown period** of **5.0 simulation-seconds** per UE after any handover. During the cooldown:

1. The UE's last_handover_time is recorded
2. Before evaluating any new handover for this UE, the xApp checks: $\text{current_sim_time} - \text{last_HO_time} > 5.0$?
3. If NO (still in cooldown): skip handover evaluation entirely for this UE
4. If YES (cooldown expired): allow normal A3 + GRU evaluation

Cooldown condition: $(\text{current_time} - \text{last_HO_time}) > \text{COOLDOWN_TIME (5.0 sim-seconds)}$

***Note:** The 5.0s cooldown was chosen because: (1) it is long enough to prevent rapid oscillation, (2) it matches typical 3GPP recommended TTT (Time-to-Trigger) scales for 5G NR, (3) shorter cooldowns (e.g. 2s) still allowed some PPs in our experiments.*

5.4 GRU vs Pure A3 — The Improvement

Traditional 3GPP A3-only handover is purely **reactive** — it waits until the neighbor is already stronger and then switches. The GRU adds **proactive prediction** — it looks at the trend (is the UE moving toward or away from the candidate cell?) before deciding.

Aspect	Pure A3 Handover	A3 + GRU Handover (Our System)
Decision basis	Instantaneous SINR comparison	10-sample time-series trend analysis
Prediction	None — reactive only	Predicts future best cell from movement history
Ping-pong risk	Higher — may HO to briefly-better cell	Lower — GRU sees if improvement is sustained
Wrong-cell HO	Can happen (noise spike triggers HO)	GRU validates it is a real trend
Cooldown	Not standard	5.0 sim-second cooldown per UE
Typical PP rate (LTE)	5–15% (literature)	3–4% (our system)
Our PP rate	N/A (no baseline run)	1.72% – 3.65% across simulations

xApp Implementation — best2.c Explained

6.1 Overview of best2.c

The xApp source file is **best2.c**, located at: `/home/omar_farouk/open-ran-clean/yousef_fathy/ns-0-RAN-flexric/mmwave-LENA-ora/best2.c`. It is written in C and compiled into the executable **xapp_handover_gru**. The C language was chosen because FlexRIC's SDK is C-native and C gives the best performance for real-time processing. The file contains ~500 lines of code including: subscription setup, KPM indication callback, A3 logic, GRU HTTP call, RC command construction, and state management.

6.2 All Thresholds and Parameters

Constant Name	Value	Units	Meaning and Why This Value
A3_OFFSET	1.0	dB	Minimum gain before considering HO — prevents HO for tiny improvements
A3_HYSTERESIS	1.0	dB	Extra noise margin — prevents oscillation from measurement noise
A3_COMBINED	2.0	dB	= A3_OFFSET + A3_HYSTERESIS — actual trigger threshold
COOLDOWN_TIME	5.0	sim-seconds	After any HO, ignore this UE for 5s — prevents ping-pong
KPM_INTERVAL	0.05	sim-seconds	Report period — 20 reports/sim-second per UE
WINDOW_SIZE	10	samples	GRU rolling window — covers 0.5s of recent history
GRU_SERVICE_PORT	5000	TCP port	Python Flask GRU service listens here
GRU_SERVICE_HOST	127.0.0.1	IP	Loopback — GRU runs on same machine
NUM_CELLS	7	count	Total cells in simulation (IDs 0-6)
NUM_UES	20	count	Total UEs in simulation

6.3 UE State Machine

For each of the 20 UEs, the xApp maintains a state. The state transitions are:

- **IDLE:** UE is stable, no handover pending. xApp evaluates A3 condition on every KPM report.
- **COOLDOWN:** A handover just completed. xApp waits until $\text{sim_time} - \text{last_HO_time} > 5.0\text{s}$ before re-evaluating.
- **HO_REQUESTED:** xApp has sent RC handover command, waiting for execution confirmation from NS-3.

The per-UE state variables maintained in the xApp:

Variable	Type	Purpose
last_ho_time[ue_id]	double	Simulation time of last handover for this UE
serving_cell[ue_id]	int	Current serving cell ID (0-6)
sinr_window[ue_id][10]	float array	Rolling window of last 10 SINR values
rsrp_window[ue_id][10]	float array	Rolling window of last 10 RSRP values
window_idx[ue_id]	int	Current position in rolling window (0-9)
window_full[ue_id]	bool	True when window has 10 complete samples

6.4 Decision Logic — Step by Step

When a KPM Indication arrives (every 0.05 sim-seconds per UE):

1. **Extract data** from KPM report: UE ID, serving cell, SINR of all 7 cells, RSRP, timestamp
2. **Update rolling window** for this UE: append new SINR/RSRP to the 10-slot window
3. **Cooldown check**: if $(\text{sim_time} - \text{last_ho_time}[\text{ue}]) < 5.0 \rightarrow \text{SKIP}$, do nothing
4. **A3 check**: find the best neighbor cell. If $\text{best_neighbor_SINR} < \text{serving_SINR} + 2.0 \text{ dB} \rightarrow \text{SKIP}$
5. **Window check**: if window is not full (fewer than 10 samples) $\rightarrow$ use A3 best_neighbor directly
6. **GRU call**: send HTTP POST to `http://127.0.0.1:5000/predict` with the 10-sample window
7. **Parse response**: extract recommended cell_id and confidence score
8. **Validate GRU recommendation**: confirm recommended cell passes A3 threshold too
9. **Send RC command**: construct E2SM-RC handover command with target cell_id
10. **Update state**: set $\text{last_ho_time}[\text{ue}] = \text{sim_time}$, record new serving cell

6.5 HTTP Communication with GRU Service

Request Format (JSON sent to `http://127.0.0.1:5000/predict`)

```
POST /predict HTTP/1.1
Content-Type: application/json

{
  "ue_id": 5,
  "serving_cell": 2,
  "window": [
    {"sinr": -5.2, "rsrp": -85.1, "velocity": 2.3, "cell_sinrs":
    [-5.2, -8.1, -3.0, -12.5, -9.8, -7.3, -11.2]},
    ... (10 entries total)
  ]
}
```

```
}
```

Response Format (JSON returned from GRU service)

```
{  
  "cell_id": 2,  
  "confidence": 0.847,  
  "all_probs": [0.02, 0.05, 0.847, 0.01, 0.03, 0.02, 0.03]  
}
```

6.6 Fallback When GRU Service Is Down

If the HTTP call to port 5000 fails (connection refused, timeout, or malformed response), the xApp falls back to **pure A3 decision**: it picks the neighbor cell with the highest SINR among those that pass the A3 threshold. This ensures the handover system continues to function even if the Python GRU service crashes.

***Important:** In `gru.sh`, the GRU service (`gru_xapp.py`) is started **FIRST** (step 1), and the xApp (`best2.c` executable) starts **AFTER** a 2-second wait. This ensures `gru_xapp.py` is accepting connections before the first HTTP request arrives.*

GRU Python Service — gru_xapp.py Explained

7.1 Overview

The file **gru_xapp.py** is a Python Flask web service that runs the trained GRU model. It is completely separate from the C xApp — they communicate only via HTTP. This separation has important advantages: the ML model can be updated without recompiling the C xApp, Python's TensorFlow/Keras ecosystem is far easier to use than C for ML inference, and the service can be restarted independently if it crashes.

File location: /home/omar_farouk/open-ran-clean/yousef_fathy/ns-0-RAN-flexric/mmwave-LENA-oran/gru_xapp.py

7.2 Flask Server Setup

Flask is a lightweight Python web framework. In **gru_xapp.py**, Flask is used to create a simple HTTP server with one main route: `/predict`. The server listens on **port 5000** — a common convention for Flask development servers. Port 5000 was chosen because it does not require root privileges (ports below 1024 do) and does not conflict with other services in our stack (InfluxDB on 8086, Grafana on 3000, controller.py on 8001, Vite on 3001).

7.3 Model Loading

When **gru_xapp.py** starts, it immediately loads the trained GRU model from disk. The model is stored in Keras format (.h5 or .keras). Loading happens once at startup — not on every request — because loading a neural network from disk takes ~1 second, which would be unacceptable latency for each handover decision.

```
model = tf.keras.models.load_model('gru_model.h5') # loaded once at startup
```

7.4 The /predict Endpoint — Complete Flow

1. **Receive JSON POST request** from xApp with 10-sample window
2. **Parse JSON**: extract `ue_id`, `serving_cell`, and 10-element window array
3. **Validate**: confirm window has exactly 10 samples, each with required fields
4. **Build input array**: shape (1, 10, num_features) — batch size 1, 10 timesteps, N features
5. **Normalize features**: apply pre-computed mean/std scaling (StandardScaler or MinMaxScaler)
6. **Run model.predict()**: GRU processes the sequence, Dense+softmax outputs 7 probabilities
7. **Find argmax**: the cell index with highest probability = recommended cell
8. **Return JSON response**: {cell_id, confidence, all_probs}

7.5 Feature Normalization

Raw SINR values range from about -20 dB to +30 dB. Raw RSRP values range from -140 dBm to -44 dBm. Neural networks train and infer better when inputs are in a small range (e.g., -1 to +1 or 0 to 1). The normalization formula applied to each feature is:

$$\text{normalized_value} = (\text{raw_value} - \text{mean}) / \text{std_deviation}$$

The mean and std_deviation values are computed from the training dataset and saved alongside the model. They are loaded at startup and applied to every incoming sample.

7.6 Error Handling

Error Condition	xApp Response	Fallback Behavior
Connection refused (service not running)	HTTP error	xApp uses pure A3 decision
Window has < 10 samples	Returns HTTP 400 with 'insufficient_data'	xApp uses pure A3 decision
Model returns NaN	Service returns HTTP 500	xApp uses pure A3 decision
Timeout (> 100ms)	xApp times out HTTP call	xApp uses pure A3 decision

Data Pipeline — `sim_data_pusher.py` → InfluxDB → Grafana

8.1 Purpose of the Data Pipeline

While the simulation runs, a parallel process (`sim_data_pusher.py`) reads the NS-3 output and pushes it to **InfluxDB**, a time-series database. **Grafana** then reads from InfluxDB and displays live dashboards accessible at `http://localhost:8000`. This is the **2D GUI** of our system.

This pipeline is entirely separate from the control plane — it does not affect the xApp's handover decisions. It exists purely for visualization and monitoring. Think of it as the 'dashboard' in a car — it shows the driver what is happening without being part of the steering system.

8.2 What Data is Pushed to InfluxDB

Measurement Name	Tags	Fields	Push Rate
ue_sinr	ue_id, cell_id	sinr (dB)	Every KPM report (~0.05s)
ue_rsrp	ue_id, cell_id	rsrp (dBm)	Every KPM report
ue_position	ue_id	x, y (meters)	Every 0.5s
handover_event	ue_id	source_cell, target_cell, is_pingpong	On each HO
cell_load	cell_id	num_connected_ues	Every 1s

8.3 InfluxDB — What It Is

InfluxDB is a time-series database — meaning every data point has a timestamp, and the database is optimized for storing and querying data over time. Regular databases (like MySQL) store rows; InfluxDB stores measurements indexed by time. It runs as a Docker container, listening on port **8086**. The data is organized as:

- **Measurement:** Like a table name (e.g., 'ue_sinr')
- **Tags:** Indexed metadata for filtering (e.g., ue_id=5, cell_id=2)
- **Fields:** Actual numeric values (e.g., sinr=-5.2)
- **Timestamp:** Automatically recorded with nanosecond precision

8.4 Grafana Dashboard

Grafana is a visualization tool that connects to InfluxDB and creates real-time charts and dashboards. It runs as a Docker container on port **3000**, proxied through nginx to port **8000** for external access. Our Grafana

dashboard shows:

- SINR time-series plots for each UE (one line per UE, colored by serving cell)
- Cell load bar chart (how many UEs each cell serves over time)
- Handover event markers on the SINR plots
- UE position heatmap (if position data is available)
- Ping-pong event counter

8.5 Docker Compose Setup

The 2D GUI runs via Docker Compose in the **GUI/** directory. The containers are:

Container	Image	Port	Purpose
influxdb	influxdb:2.x	8086	Time-series database — receives pushed data
grafana	grafana/grafana	3000 (internal)	Visualization dashboards
nginx	nginx:alpine	8000 (external)	Reverse proxy — maps port 8000 → Grafana 3000

3D GUI, Orchestration, and Startup Sequence

9.1 The 3D GUI — Two-Component Architecture

The 3D visualization system has two parts that work together:

- **FastAPI Controller (controller.py)** — a Python API server running on port **8001**. It manages the simulation lifecycle (start, stop, status) and reads simulation output files.
- **Vite Frontend** — a React/Three.js web application running on port **3001**. It shows a 3D visualization of the cell layout, UE positions, and handover events in real time.

9.2 controller.py API Routes

Route	Method	Purpose	Returns
/api/status	GET	Current simulation state: running/stopped, sim_time, HO count	JSON
/api/start	POST	Launch gru.sh — starts full simulation stack	JSON confirmation
/api/stop	POST	Execute kill_sim.sh — stops all processes	JSON confirmation
/api/save	POST	Save current results to 3D_GUI_Sim_Results/ directory	JSON with save path
/api/results	GET	List all saved simulation result directories	JSON array
/api/ue_positions	GET	Latest UE positions and cell attachments	JSON array

9.3 Results Storage Structure

When a simulation ends, results are saved to a timestamped directory:

```
3D_GUI_Sim_Results/sim010_20260505_210259_gru_scenario/
■■■ decision_log.csv # Every HO decision with timestamp, UE, cells
■■■ summary.json # Total HOs, PPs, PP rate, accuracy
■■■ sinr_history.csv # SINR over time per UE
■■■ metadata.json # Sim parameters, git commit, run time
```

9.4 gru.sh — The Complete Startup Script

The file **gru.sh** is the master launcher for the entire simulation stack. It starts all processes in the correct order:

1. **Start GRU Python service:** `python3 gru_xapp.py &` — runs in background
2. **Wait 2 seconds:** ensures Flask is ready before xApp tries to connect
3. **Start FlexRIC:** `./nearRT-RIC &` — runs in background, logs to `/tmp/flexric.log`
4. **Wait 1 second:** ensures FlexRIC E2 socket is open
5. **Start xApp:** `./xapp_handover_gru &` — subscribes to FlexRIC immediately
6. **Wait 3 seconds:** ensures xApp has subscribed and FlexRIC is ready for NS-3
7. **Start NS-3:** `./ns3 run gru_scenario --simTime=60` — connects to FlexRIC, simulation begins
8. **Start sim_data_pusher.py:** runs in background to push data to InfluxDB

Important: The order in `gru.sh` is **CRITICAL**. If NS-3 starts before FlexRIC, the E2 connection will fail. If the xApp starts before `gru_xapp.py`, the first handover calls will fall back to pure A3. Any change to this order must be tested carefully.

9.5 kill_sim.sh — Shutdown Sequence

Stopping the simulation also requires a specific order to avoid orphaned processes or corrupted output files:

1. Kill NS-3 process first (stops new data from being generated)
2. Kill `sim_data_pusher.py` (stops InfluxDB writes — avoids partial data)
3. Kill xApp (stops handover decisions)
4. Kill FlexRIC (closes E2 interface)
5. Kill `gru_xapp.py` (stops GRU service)

Simulation Results and Analysis

10.1 Results Summary Table

Simulation	Sim Time (s)	Total HOs	Ping-Pong HOs	PP Rate (%)	Accuracy (%)
sim006	60	334	4	1.72%	98.28%
sim010	60	137	5	3.65%	96.35%
sim011	120	309	10	3.24%	96.76%

10.2 How Accuracy is Calculated

The accuracy metric measures how often the xApp made a 'good' handover decision — meaning a handover that did NOT lead to a ping-pong. The formula is:

$$\text{Accuracy} = ((\text{Total HOs} - \text{Ping-Pong HOs}) / \text{Total HOs}) \times 100\%$$

For sim011: Accuracy = $((309 - 10) / 309) \times 100\% = 299/309 \times 100\% = 96.76\%$

Note: This is a conservative accuracy measure. It does NOT measure whether the chosen cell was the theoretically optimal cell — only whether the handover avoided ping-pong. The actual 'correctness' of cell choice would require a ground-truth oracle comparing every decision to the best possible decision, which we do not have.

10.3 Why Ping-Pong Count Scales with simTime

Observe that sim011 (120s) has 10 PPs while both 60s simulations have ~4-5 PPs. This makes intuitive sense: if the PP rate per unit time is approximately constant (because it depends on how often UEs are at cell edges, which is determined by the random mobility model), then doubling the simulation time roughly doubles the total PP count.

$$\text{Expected PP count} \approx \text{PP_rate_per_second} \times \text{simTime}$$

$$\text{sim011: } \sim 10 \text{ PPs over 120s} \approx 0.083 \text{ PPs/second}$$

$$\text{sim006: } \sim 4 \text{ PPs over 60s} \approx 0.067 \text{ PPs/second (slightly fewer – different random seed)}$$

10.4 Why PP Rate Varies Between Simulations

The PP rate (%) varies more than the raw count because the **denominator (total HOs) also changes** with the random seed. sim006 had 334 total HOs but sim010 had only 137 — a huge difference for the same 60s sim time. Why?

- Different random seeds produce different UE movement patterns

- If UEs happen to spend more time at cell edges → more HOs → lower PP rate (same PP count, larger denominator)
- If UEs spend more time in cell centers → fewer HOs → higher PP rate (same PP count, smaller denominator)
- sim006's UEs apparently moved more (or in more boundary-crossing paths) than sim010's UEs

10.5 Comparison with Literature

System	HO Algorithm	Typical PP Rate	Source
Traditional LTE (4G)	A3 event only	5–15%	3GPP TS 36.331 studies
5G NR mmWave (baseline)	A3 event only	8–20%	mmWave handover literature
ML-assisted HO (various)	CNN/LSTM + A3	2–6%	IEEE papers 2020-2023
Our system (GRU + A3)	GRU + A3 + cooldown	1.72–3.65%	Our sim006/010/011

Our results sit at the better end of the ML-assisted HO range. The combination of GRU prediction AND the 5-second cooldown is what achieves this — the cooldown alone would reduce PPs but might block necessary handovers; the GRU alone without cooldown would still sometimes trigger rapid oscillation.

10.6 What 96.76% Accuracy Means in Practice

96.76% accuracy means that for every 100 handover decisions, approximately 97 of them resulted in a stable attachment (no ping-pong) and 3 resulted in a bounce-back. In a real network with 20 UEs over 2 minutes, this translates to roughly 3 extra handover signaling messages per 100 UE-minutes — a very low overhead. The remaining 3.24% of ping-pong events are those cases where the random walk trajectory was genuinely ambiguous (UE walked back toward original cell after the HO).

Complete System Architecture Summary

11.1 Full System Data Flow

Here is the complete data flow from simulation to visualization, described step by step:

1. **NS-3 gru_scenario.cc starts** → creates 7 gNBs + 20 UEs → opens SCTP connection to 127.0.0.1:36421
2. **FlexRIC nearRT-RIC** → accepts E2 Setup from NS-3 → relays xApp KPM subscription to NS-3
3. **NS-3 generates KPM reports** → every 0.05 sim-sec → sent via E2 to FlexRIC → FlexRIC delivers to xApp
4. **xApp (best2.c)** → checks cooldown → checks A3 → if triggered: HTTP POST to :5000/predict
5. **gru_xapp.py** → receives 10-sample window → normalizes → model.predict() → returns best cell_id
6. **xApp** → constructs E2SM-RC control message → sends via FlexRIC E2 to NS-3
7. **NS-3 executes HO** → UE disconnects from old cell → connects to target cell
8. **sim_data_pusher.py (parallel)** → reads NS-3 output → HTTP POST to InfluxDB :8086
9. **Grafana** → queries InfluxDB every 5s → updates dashboard at :8000
10. **controller.py** → reads sim log files → exposes /api/status at :8001
11. **Vite 3D frontend** at :3001 → polls controller.py /api/ue_positions → renders 3D view

11.2 All Ports and Services

Port	Protocol	Service	Purpose
36421	SCTP	FlexRIC E2 interface	NS-3 ↔ FlexRIC communication
5000	TCP/HTTP	gru_xapp.py (Flask)	GRU inference service
8086	TCP/HTTP	InfluxDB	Time-series database
3000	TCP/HTTP	Grafana (internal)	Dashboard renderer
8000	TCP/HTTP	nginx (external)	Reverse proxy to Grafana
8001	TCP/HTTP	controller.py (FastAPI)	Simulation orchestration API
3001	TCP/HTTP	Vite dev server	3D visualization frontend

11.3 Key Files Reference

File	Location (relative to project root)	Purpose
gru_scenario.cc	yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/	NS-3 simulation script

File	Location (relative to project root)	Purpose
best2.c	yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/	xApp source code (C)
gru_xapp.py	yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/	GRU Flask inference service
gru.sh	yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/	Master startup script
kill_sim.sh	yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/	Master shutdown script
sim_data_pusher.py	yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/	InfluxDB data pusher
controller.py	3D_GUI/backend/ (or project root)	FastAPI orchestration server
docker-compose.yml	GUI/	2D GUI Docker stack
decision_log.csv	3D_GUI_Sim_Results/simXXX.../	Per-HO decision log
flexric.log	/tmp/	FlexRIC runtime log (always here)

11.4 All Parameters in One Place

Parameter	Value	Where Defined
A3_OFFSET	1.0 dB	best2.c (compile-time constant)
A3_HYSTERESIS	1.0 dB	best2.c (compile-time constant)
A3_COMBINED threshold	2.0 dB	= A3_OFFSET + A3_HYSTERESIS
COOLDOWN_TIME	5.0 sim-seconds	best2.c (compile-time constant)
KPM report interval	0.05 sim-seconds	gru_scenario.cc (ns3 parameter)
Rolling window size	10 samples	best2.c + gru_xapp.py
GRU service port	5000	gru_xapp.py (Flask app.run)
E2 port	36421 SCTP	FlexRIC config + NS-3 E2 module
Number of cells	7 gNBs	gru_scenario.cc
Number of UEs	20	gru_scenario.cc
Carrier frequency	28 GHz	gru_scenario.cc
Simulation time	60s or 120s	gru.sh (--simTime argument)
UE mobility	Random Walk 2D	gru_scenario.cc
Simulation pace	~0.005 sim-s/real-s	Computational characteristic

Defense Q&A; — 60+ Questions with Detailed Answers

This chapter contains over 60 questions that a thesis committee might ask, organized into categories. Each answer is written in simple language with specific numbers and facts. Read every answer carefully — committees often ask follow-up questions based on your initial answer.

Category 1: O-RAN Basics (10 Questions)

Q: What does O-RAN stand for and what is its main purpose?

A: O-RAN stands for Open Radio Access Network. Its main purpose is to break the vendor lock-in that existed in traditional mobile networks, where operators were forced to buy all equipment from a single vendor. O-RAN defines open interfaces (like E2, A1, F1) so equipment from different manufacturers can interoperate. It was formally established by the O-RAN Alliance in 2018. The key benefit is that software intelligence (xApps, rApps) can be written by anyone — universities, startups, operators — and plugged into any O-RAN-compliant network.

Q: What is the difference between the near-RT RIC, non-RT RIC, and the E2 Node?

A: The non-RT RIC (Non-Real-Time RAN Intelligent Controller) operates on timescales greater than 1 second and handles long-term optimization, policy generation, and ML model training. The near-RT RIC operates between 10 milliseconds and 1 second and runs xApps that make real-time decisions like handover management. The E2 Node is the actual radio hardware (gNB or its components) that sends measurement reports to the near-RT RIC and receives control commands from it. In our system: FlexRIC = near-RT RIC, NS-3 gNBs = E2 Nodes.

Q: What is the E2 interface and what two functions does it serve?

A: The E2 interface connects the near-RT RIC to E2 Nodes (gNBs). It serves two functions: (1) Monitoring, via the E2SM-KPM service model, which allows the RIC to subscribe to measurement reports containing metrics like SINR, RSRP, and cell load. In our system, KPM reports arrive every 0.05 simulation-seconds per UE. (2) Control, via the E2SM-RC service model, which allows the RIC to send commands to the gNB, such as triggering a handover for a specific UE to a target cell. The E2 interface uses SCTP protocol on port 36421, which is the standard O-RAN port.

Q: Why does FlexRIC log to /tmp/flexric.log instead of the terminal?

A: FlexRIC is designed to run as a background daemon process in production environments, where there is no terminal attached. Logging to a file (/tmp/flexric.log) is the correct design for any service that might run unattended. The /tmp/ directory is used because it is always writable without special permissions. In our development environment, this means we must check this file explicitly to diagnose FlexRIC issues. The command 'tail -f /tmp/flexric.log' shows live log output. This is not a bug — it is intentional

production-grade logging behavior.

Q: What is an xApp and how is it different from a regular application?

A: An xApp is an application specifically designed to run inside a near-RT RIC and interact with the radio network via the E2 interface. Unlike a regular application, an xApp has a defined lifecycle: it registers with the RIC, subscribes to E2 service models (like KPM), processes real-time network events, and sends control commands via E2. Regular applications have no access to network control planes. Our xApp (`xapp_handover_gru`, built from `best2.c`) specifically subscribes to KPM reports at 0.05s intervals and sends RC handover commands when the GRU model recommends a cell change.

Q: What is the A1 interface and do you use it in your system?

A: The A1 interface connects the non-RT RIC to the near-RT RIC. It is used to pass high-level policies and ML model updates from the non-RT RIC (which has time to do complex training) to the near-RT RIC (which must act quickly). In our current system, we do NOT use the A1 interface. The GRU model was pre-trained offline (by Fares) and loaded directly into `gru_xapp.py` at startup. In a full production O-RAN deployment, the A1 interface would be used to push updated GRU model weights from a training server to the running xApp without downtime.

Q: What is the difference between O-DU, O-CU, and O-RU in O-RAN?

A: In traditional 4G/5G, the baseband processing was done in a single box called the BBU (Baseband Unit). O-RAN splits this into three parts: O-RU (Open Radio Unit) handles the radio frequency processing and connects to the antenna — it is physically near the tower. O-DU (Open Distributed Unit) handles real-time lower-layer processing including the scheduling of radio resources every 0.5ms slot. O-CU (Open Central Unit) handles higher-layer protocols like RRC and PDCP and can run in a central data center. In NS-3, these are simulated as a combined gNB entity, not separately split.

Q: Why is vendor lock-in a problem in traditional RAN?

A: Vendor lock-in means that once an operator deploys, say, Ericsson equipment, they cannot easily add components from Nokia or Huawei because the interfaces are proprietary (secret, non-standard). This forces operators to buy upgrades exclusively from the original vendor at whatever price they set. It also prevents innovation: operators cannot write their own optimization algorithms. Studies estimate vendor lock-in adds 20-40% cost premium compared to open competitive markets. O-RAN solves this by standardizing interfaces so any vendor's components can interoperate.

Q: What does the O-RAN Alliance do and who are its members?

A: The O-RAN Alliance is an industry consortium that defines the open specifications for O-RAN architecture, interfaces, and protocols. It was founded in 2018 by five founding operators: AT&T, Deutsche Telekom, NTT DoCoMo, Orange, and China Mobile. Today it has over 300 member organizations including Nokia, Ericsson, Samsung, Intel, and many universities. The Alliance publishes technical specifications that define exactly how components should communicate. FlexRIC implements the O-RAN Alliance E2AP specification, and NS-3's O-RAN module implements the E2 Node side of that same specification.

Q: What would you change about O-RAN architecture if you were designing from scratch?

A: One limitation of the current O-RAN architecture is the 10ms minimum latency floor of the near-RT RIC, which makes it unsuitable for ultra-low-latency use cases like industrial automation. I would add a fourth tier between the near-RT RIC and the O-DU scheduler with sub-1ms capability. I would also standardize the xApp lifecycle management API more rigorously — currently different near-RT RIC implementations (FlexRIC, OpenRAN Gym, etc.) have different xApp APIs, requiring code changes when porting. Finally, I would add built-in A1 interface support in FlexRIC for model updates without service restart.

Category 2: GRU and Machine Learning Basics (12 Questions)

Q: Why did you choose GRU over LSTM for your model?

A: GRU (Gated Recurrent Unit) was chosen over LSTM (Long Short-Term Memory) for three reasons. First, GRU has fewer parameters than LSTM (it has 2 gates vs LSTM's 3 gates), making it faster to train and infer — important for near-real-time handover decisions. Second, for our sequence length of 10 samples, research shows GRU performs comparably to LSTM while being computationally cheaper. Third, our input sequences are relatively short and predictable (SINR trends for pedestrian-speed UEs), which does not require LSTM's extra complexity. The accuracy achieved (96.76%) validates that GRU is sufficient for this task.

Q: What is the vanishing gradient problem and how does GRU solve it?

A: The vanishing gradient problem occurs during backpropagation through time in RNNs. When gradients (the error signals used to update weights) are multiplied at each time step, they either shrink toward zero (vanishing) or grow toward infinity (exploding). For vanilla RNNs, after more than ~10 steps, the gradient practically disappears and the network cannot learn from early timesteps. GRU solves this by using the **update gate**, which provides a direct path for gradient flow that bypasses the squashing activation functions. This allows gradients to propagate back through many timesteps without vanishing, enabling the network to learn from patterns across the full 10-sample window.

Q: What are the input features to your GRU model?

A: The GRU model receives a 10-timestep sequence per UE. At each timestep (a KPM report), the features include: (1) SINR of the serving cell in dB, (2) RSRP of the serving cell in dBm, (3) UE velocity estimate in m/s (derived from position change), (4) the serving cell ID encoded as an integer, and (5) SINR values from all 6 neighbor cells. This gives approximately 10 features per timestep, resulting in an input tensor of shape (1, 10, ~10) for each prediction request. All features are normalized using StandardScaler parameters computed from the training dataset.

Q: What does the softmax output of your GRU mean?

A: The softmax function is the final activation in our GRU model. It takes the 7 raw output scores (one per cell) and converts them to probabilities that sum to exactly 1.0. For example, if the raw output is [1.2, 0.3, 3.8, 0.1, 0.5, 0.2, 0.4], softmax might output [0.05, 0.02, 0.85, 0.01, 0.03, 0.02, 0.02] — meaning the model

is 85% confident that Cell 2 is the best choice. The `cell_id` returned is `argmax` of this array. The confidence value helps the xApp assess reliability: a confidence of 0.90+ is a strong recommendation, while 0.40 means the model is uncertain.

Q: How was the GRU model trained and what data was used?

A: The GRU model was trained by Fares (the ML team member) using labeled data collected from NS-3 simulation runs. The training data consists of (window, best_cell) pairs where best_cell is the cell that provides the highest sustained SINR after the handover point. Training used categorical cross-entropy loss and the Adam optimizer. The model was evaluated on a held-out test set to measure accuracy before deployment. The trained model was saved as a .h5 file and loaded by `gru_xapp.py` at runtime. One limitation is that training and testing used NS-3 data — a real-world deployment would need training on real network measurements.

Q: What is overfitting and how do you know your model does not overfit?

A: Overfitting occurs when a model memorizes the training data instead of learning general patterns — it performs well on training examples but poorly on new data. We assess overfitting by comparing training accuracy vs validation accuracy during training: if training accuracy is much higher than validation, the model is overfitting. In our case, the model achieves ~96-98% accuracy in the live simulation, which uses different random seeds than the training data, suggesting the model generalizes well. Regularization techniques such as dropout layers or L2 weight regularization could be added to further prevent overfitting if needed.

Q: Why use a rolling window of 10 samples specifically?

A: The window size of 10 samples was chosen based on experimentation. At 0.05s per sample, 10 samples cover 0.5 simulation-seconds of history. This window is long enough to detect a trend (a UE moving toward a cell will show consistently rising SINR from that cell over 5-10 steps) but short enough to respond to sudden direction changes. Shorter windows (e.g., 5 samples = 0.25s) might not capture stable trends; longer windows (e.g., 20 samples = 1.0s) would delay response and cover too much history for a pedestrian-speed UE. The 10-sample choice balances trend visibility with responsiveness.

Q: What is the difference between SINR and RSRP? Why use both?

A: SINR (Signal-to-Interference-plus-Noise Ratio) measures the quality of a signal — it compares signal strength to background interference and noise, expressed in dB. A high SINR means the signal is clean and strong relative to interference. RSRP (Reference Signal Received Power) measures only the absolute power level of the reference signal in dBm — it does not account for interference. We use both because they provide complementary information: RSRP tells us how strong the signal is, while SINR tells us how usable it is. A cell might have high RSRP but poor SINR if there is a lot of interference from neighboring cells. The GRU can use both metrics to make better decisions than with either alone.

Q: Could you replace GRU with a simpler model like a decision tree or SVM?

A: Technically yes, but GRU is more appropriate for this task. Decision trees and SVMs treat each sample independently — they cannot capture temporal patterns like 'SINR from Cell 3 has been rising for 10

consecutive steps.' A GRU explicitly models the sequence order and can detect trends. Additionally, SVMs require feature engineering (manually creating trend features like 'slope of last 5 SINR values'), while the GRU learns these features automatically. The sequential nature of the handover decision problem — where the history of measurements matters — naturally favors sequence models like GRU.

Q: What is the difference between classification and regression in your context?

A: Our GRU model is a **classification** model: it outputs discrete class labels (cell IDs 0-6). Each output class represents a specific cell to hand over to. We use softmax activation and categorical cross-entropy loss, which are the standard choices for multi-class classification. An alternative design would be to use **regression** — predict a continuous score for each cell and pick the maximum. The outputs would be similar in practice, but classification with softmax produces probabilities that are easier to interpret as confidence values. Our approach of returning 7 class probabilities is standard for multi-class neural network classification.

Q: How does the GRU handle a new UE that has fewer than 10 historical samples?

A: When a UE has fewer than 10 samples in its rolling window (which happens at the start of the simulation or if a UE is newly attached), `gru_xapp.py` returns an HTTP 400 response with the message 'insufficient_data'. The xApp detects this error response and falls back to a **pure A3 decision**: it picks the neighbor cell with the highest SINR that exceeds the A3 threshold (`serving_SINR + 2.0 dB`). This fallback ensures the system continues making reasonable handover decisions even before the GRU has enough context. Typically, the window fills within 0.5 simulation-seconds ($10 \text{ samples} \times 0.05\text{s/sample}$).

Q: Why is categorical cross-entropy used as the loss function?

A: Categorical cross-entropy is the standard loss function for multi-class classification problems with mutually exclusive classes. In our case, a UE should connect to exactly one cell at a time, making the classes mutually exclusive. The cross-entropy formula penalizes the model heavily when it is confident but wrong — for example, if the model says 90% probability for Cell 2 but the correct answer is Cell 3, the loss is very high. This encourages the model to be correctly confident, not just accurate. Mean squared error (MSE), the alternative, does not penalize confident wrong answers as heavily and is better suited for continuous regression outputs.

Category 3: System Implementation (12 Questions)

Q: Why did you use NS-3 instead of a real network or a different simulator?

A: We used NS-3 for three main reasons. First, cost: deploying a real 7-cell 28 GHz mmWave network would cost millions of dollars in hardware, licensing, and site acquisition. Second, control: NS-3 lets us precisely set every parameter — cell positions, UE speeds, mobility patterns, random seeds — and reproduce experiments exactly. Third, integration: NS-3 has a built-in O-RAN E2 interface module that connects natively to FlexRIC, saving months of integration work. Alternative simulators like MATLAB or OMNET++ also exist, but NS-3's mmWave module and E2 support made it the best fit for our specific scenario. The main limitation is simulation speed: 60 sim-seconds takes ~3 hours of real computation.

Q: Why does the simulation take so long (3 hours for 60 sim-seconds)?

A: The simulation pace of ~0.005 sim-seconds per real-second is caused by the computational cost of several simultaneous operations: (1) mmWave channel modeling at 28 GHz with accurate path loss, shadowing, and small-scale fading for 20 UEs across 7 cells, (2) encoding and decoding E2AP messages in ASN.1 format for every KPM report, (3) HTTP calls from the xApp to the GRU Python service for each UE evaluation, (4) Python model inference overhead (even though it is fast, it adds up for 20 UEs at 20 evaluations/second), and (5) InfluxDB writes from `sim_data_pusher.py`. Using a faster machine, disabling some logging, or reducing the number of UEs would reduce this time.

Q: What is SCTP and why does O-RAN use it instead of TCP for the E2 interface?

A: SCTP (Stream Control Transmission Protocol) is a transport protocol like TCP and UDP. It was designed specifically for telecom signaling. Unlike TCP, SCTP supports multi-homing (a connection can use multiple IP addresses for redundancy), multi-streaming (multiple independent message streams over one connection), and message-boundary preservation (messages are delivered as complete units, not byte streams). These properties are important for E2: if the RIC machine has two network interfaces, SCTP can use both for redundancy. The O-RAN Alliance chose SCTP for E2 because 3GPP S1AP (used in LTE for similar control-plane signaling) also uses SCTP. Port 36421 is the IANA-registered port for this type of application.

Q: How does the xApp know which UE ID maps to which RNTI in NS-3?

A: The RNTI (Radio Network Temporary Identifier) is the unique ID assigned to each UE by the serving cell. In NS-3, each UE is assigned an RNTI when it attaches to its first cell. This RNTI is included in every KPM report. The xApp maintains a mapping array indexed by RNTI so that when a KPM report arrives with RNTI=5, the xApp knows to update the state for UE #5. In our simulation with 20 UEs, RNTIs are typically 1-20. If a UE performs a handover and gets a new RNTI from the target cell, the xApp must handle this ID change — this is one complexity in real O-RAN deployments that NS-3 simplifies by keeping RNTI assignments consistent.

Q: What happens if FlexRIC crashes during the simulation?

A: If FlexRIC crashes: (1) NS-3 loses the SCTP connection to 127.0.0.1:36421 and will log an E2 connection error. NS-3 continues simulating but stops sending KPM reports since there is no receiver. (2) The xApp loses its connection to FlexRIC and enters an error state. (3) Handover decisions stop. The

simulation continues producing UE movement and channel data, but no handovers are executed. The result would be UEs that degrade to very poor SINR without being handed over. In our setup, FlexRIC stability has been good in practice. To handle crashes gracefully in production, one would use a process supervisor like systemd to automatically restart FlexRIC.

Q: Why is the GRU service written in Python while the xApp is written in C?

A: The GRU service uses TensorFlow/Keras for model inference, which is overwhelmingly implemented in Python. Writing ML inference in C would require either implementing a custom neural network from scratch or using C bindings to TensorFlow (which are complex and poorly documented). Python's Keras/TensorFlow makes loading and running a trained model trivially easy (two lines of code: `load_model + model.predict`). The xApp is in C because FlexRIC's native SDK is C-based, and C provides the lowest latency for processing E2 messages. The HTTP interface between the two components is a clean separation — each component uses the best language for its task.

Q: How does the cooldown timer interact with the simulation clock vs real clock?

A: The cooldown timer uses the **simulation clock**, not the real clock. When the xApp records `last_ho_time[ue] = sim_time`, the `sim_time` is the NS-3 simulation timestamp (e.g., 32.5 sim-seconds), not the wall-clock time. The cooldown check is: `current_sim_time - last_ho_time[ue] > 5.0` sim-seconds. Since the simulation runs 200x slower than real time, 5 sim-seconds of cooldown corresponds to ~1000 real-seconds (about 16 minutes) of wall-clock waiting. This is intentional — the cooldown is meaningful in the simulated network time, not in the researcher's real time. All timing thresholds in our system are in simulation-seconds unless otherwise specified.

Q: What is the decision_log.csv file and what does it contain?

A: `decision_log.csv` is a comma-separated log file written by the xApp or a post-processing script. Each row represents one handover event. Columns include: `sim_time` (when the HO was decided), `ue_id` (which UE), `source_cell` (previous cell), `target_cell` (new cell), `gru_confidence` (confidence score from the GRU prediction), `a3_triggered` (boolean, was A3 condition met), `is_pingpong` (boolean, was this HO a ping-pong return to a previous cell within 5s). The ping-pong detection script reads this file and counts rows where `is_pingpong=True`, grouped per UE, to compute the final PP statistics shown in Chapter 10.

Q: How do you detect ping-pong events in post-processing?

A: Ping-pong detection works by analyzing the `decision_log.csv` file. For each UE, the log is sorted by `sim_time`. For consecutive handover pairs (HO1: A→B at time `t1`, HO2: B→A at time `t2`), we check if $(t2 - t1) < 5.0$ sim-seconds AND `target_cell` of HO2 equals `source_cell` of HO1. If both conditions are true, HO2 is flagged as a ping-pong. The total PP count is the number of flagged HOs. The PP rate = `flagged_HOs / total_HOs`. An important subtlety: a handover is flagged as PP even if it was made by the GRU service correctly — what matters is the outcome (did the UE bounce back?), not the decision method. This is why the 5-second cooldown is added in the xApp: to prevent this situation.

Q: Why is port 5000 used for the GRU service?

A: Port 5000 is Flask's default development server port — it was chosen for simplicity. It does not require root privileges (which ports below 1024 need). It does not conflict with other services in our stack: InfluxDB uses 8086, Grafana uses 3000 (internal), nginx proxy uses 8000, FastAPI controller uses 8001, and Vite frontend uses 3001. In a production deployment, the GRU service would likely be placed behind a proper WSGI server (like gunicorn) on a non-default port, but port 5000 is perfectly adequate for our simulation environment.

Q: What would happen if the `A3_OFFSET` was set to 0.0 dB instead of 1.0 dB?

A: Setting `A3_OFFSET` to 0.0 dB would make the combined threshold $0.0 + 1.0 = 1.0$ dB (only hysteresis remains). This means a neighbor cell only needs to be 1 dB better to trigger A3 evaluation. In practice, SINR measurements fluctuate by several dB due to small-scale fading, especially in mmWave channels. A 1 dB margin would cause many false A3 triggers from normal channel noise. The GRU would then be called much more frequently, and even with the 5s cooldown, the system would likely see higher ping-pong rates because more marginal handover opportunities would be evaluated. The 1.0 dB offset provides a meaningful 'the neighbor is actually better' gate.

Q: How does `sim_data_pusher.py` know what data to push? Does it run in real time?

A: `sim_data_pusher.py` reads NS-3 simulation output files (or a named pipe if configured for real-time output). NS-3 writes trace files including mobility traces (UE positions), KPM report logs, and handover logs. `sim_data_pusher.py` parses these files and sends the data to InfluxDB using the InfluxDB Python client library. It runs in parallel with the simulation (as a background process), sleeping briefly between batches to avoid overwhelming InfluxDB. The push rate is typically every 0.5-1 simulation-second of data. Because the simulation runs ~200x slower than real time, the pusher has ample real-time to process and push data before the next batch arrives.

Category 4: Results Interpretation (8 Questions)

Q: Why did sim006 have 334 handovers but sim010 only had 137, for the same 60 seconds?

A: Both simulations ran for 60 sim-seconds with 20 UEs, but different random seeds produced very different UE movement patterns. In NS-3's Random Walk 2D mobility model, each run with a different seed creates different trajectories. In sim006, UEs happened to traverse more cell boundaries — perhaps due to higher-speed random walks or paths that crossed the cell overlap zones repeatedly. In sim010, UEs may have spent more time in cell centers or had fewer boundary crossings. The total HO count is very sensitive to the mobility realization. This is a known property of random mobility models and is why multiple simulation runs with different seeds are needed to obtain statistically meaningful averages. With 3 simulation runs, our results show the trend but would benefit from more runs.

Q: Is a PP rate of 3.65% in sim010 worse than the 1.72% in sim006? Which is better?

A: A lower PP rate is better (fewer wasted handovers), so sim006's 1.72% is technically better than sim010's 3.65%. However, the difference is largely explained by the different denominators: sim006 had 334 total HOs vs sim010's 137. With only 4-5 PPs each, the raw PP count is similar (indicating similar underlying system behavior), but the rate differs because it divides by a much larger number for sim006. The more meaningful comparison is the number of PPs per simulation-second: $\text{sim006} = 4/60 = 0.067/\text{s}$, $\text{sim010} = 5/60 = 0.083/\text{s}$ — comparable. Therefore, both simulations demonstrate similar system quality; the rate difference is a statistical artifact of different mobility realizations.

Q: What does 96.76% accuracy specifically mean? What is the other 3.24% doing?

A: 96.76% accuracy means that 299 out of 309 handovers in sim011 did NOT result in a ping-pong within 5 sim-seconds. These 299 handovers were 'successful' in the sense that the UE stayed at the target cell for at least 5 seconds. The remaining 3.24% (10 HOs) were cases where the UE bounced back within 5 seconds. These 10 ping-pong events occurred despite the cooldown because the UE was caught in a genuine physical boundary situation: the UE's random walk path brought it back toward the original cell faster than the 5-second cooldown expired, or the channel conditions changed sharply after the handover. These cases represent the fundamental limit of our approach for fast-moving cell-edge UEs.

Q: How do your results compare to the state of the art in the literature?

A: Traditional A3-only handover in LTE/5G systems typically achieves ping-pong rates of 5-15% depending on UE speed and cell density (from 3GPP study item reports and academic publications). ML-assisted handover systems reported in 2020-2023 IEEE papers typically achieve 2-6% PP rates in similar setups. Our system achieves 1.72-3.65% across three runs, placing it at or near the better end of the ML-assisted range. The key advantage of our approach is the combination of GRU sequential prediction AND explicit cooldown timer — many published systems use ML alone without a cooldown, leading to higher rates. One limitation is that we do not have a baseline (pure A3, no GRU) simulation to directly quantify the GRU's contribution.

Q: Why did you only run 3 simulations? Is that statistically sufficient?

A: Three simulations provide a preliminary indication of system performance but are not statistically sufficient for rigorous validation. For publication-quality results, 10-30 simulation runs with different seeds would be needed to compute mean PP rate with confidence intervals. We ran only 3 because each run takes 3-6 real hours, making 30 runs impractical for a thesis timeline. The three runs we have show consistent behavior (3-4% PP range) which suggests the system performance is stable, but the exact numbers could shift with more runs. Future work would address this by either running more simulations on a cluster computer or optimizing the simulation speed (e.g., using faster channel models).

Q: What does the simulation tell you that a real network deployment would not?

A: NS-3 simulation provides perfect ground-truth information that is unavailable in real networks: exact UE positions at all times, perfect SINR measurements without real hardware noise, deterministic channel conditions with a known seed, and the ability to replay the exact same scenario with different algorithms for fair comparison. In a real network, measurements are noisy, UE positions are uncertain, and channel conditions are non-repeatable. The simulation results therefore represent an optimistic upper bound on performance — real deployment would likely see somewhat higher PP rates due to measurement noise and unpredictable channel conditions. This is a standard limitation of simulation-based thesis work.

Q: If you ran sim011 for 180 seconds, what results would you predict?

A: Based on the linear scaling observed: sim011 (120s) had 309 HOs and 10 PPs. Extrapolating linearly: at 180s, we would expect approximately $309 \times (180/120) = 464$ HOs and $10 \times (180/120) = 15$ PPs. The PP rate would stay around $15/464 \approx 3.2\%$, consistent with our observed 3.24% for sim011. This linear prediction assumes the same random seed and mobility pattern, which would be true if the simulation seed is the same. However, the actual numbers would depend on whether the extended simulation time brings UEs into new cell-edge situations. The PP rate should remain in the 3-4% range.

Q: What is the significance of the 5-second cooldown specifically for mmWave?

A: The 5-second cooldown is particularly important for mmWave systems because mmWave channel conditions can change dramatically in very short distances due to blockage effects. A pedestrian walking past a parked car can experience a sudden 20 dB SINR drop on one cell within a fraction of a second. Without a cooldown, this could trigger a HO to a neighbor, then immediately back (once the car is passed). At pedestrian speed (~ 3 m/s) and 5-second cooldown, the UE travels 15 meters during the cooldown period — typically enough to move clearly into or out of a blockage zone. For vehicular UEs at 30 m/s, 5 seconds would be excessive (150 meters) and a shorter cooldown like 1 second would be more appropriate.

Category 5: Comparisons and Tradeoffs (8 Questions)

Q: What are the tradeoffs of a larger A3 threshold (e.g., 5 dB instead of 2 dB)?

A: A larger A3 threshold means a neighbor must be significantly better before a handover is even considered. Tradeoffs: (Pros) Fewer total handovers, lower signaling overhead, lower ping-pong rate since only clearly-better cells trigger evaluation. (Cons) UEs stay connected to degrading cells longer, experiencing poor SINR for more time before handover. In mmWave, where SINR can drop sharply near cell edges, a 5 dB threshold might mean a UE suffers very low throughput for several seconds before the

handover fires. The 2 dB threshold (1 dB offset + 1 dB hysteresis) is a balance: it ensures the neighbor is meaningfully better, but does not wait for the serving cell to become severely degraded.

Q: What are the tradeoffs of a longer cooldown (10 seconds) vs shorter (2 seconds)?

A: Longer cooldown (10s): Fewer ping-pongs (the UE cannot bounce back within 10 sim-seconds). But: if a UE genuinely needs to hand back (its trajectory loops), it is forced to stay on a degrading cell for up to 10 seconds. 10 sim-seconds at pedestrian speed means the UE has moved 30 meters — possibly far into another cell's territory entirely. Shorter cooldown (2s): More responsive to genuine trajectory reversals. But: more vulnerable to noise-induced ping-pong since the UE can bounce within 2s if SINR fluctuates. The 5-second value was chosen empirically from our experiments as the best balance for pedestrian-speed UEs in a 28 GHz 7-cell layout.

Q: How does your system compare to a Reinforcement Learning (RL) based approach?

A: Reinforcement Learning (RL) for handover would train an agent that directly learns when to hand over based on a reward signal (e.g., +1 for staying connected, -1 for ping-pong). RL advantages: it can adapt online during deployment, learns the optimal policy for specific cell layouts, and can discover non-obvious strategies. Our GRU approach advantages: it is supervised (trained on labeled data, more stable), interpretable (we can inspect GRU outputs), and has deterministic inference. RL disadvantages: requires careful reward shaping, can converge to suboptimal policies, and has exploration phase where bad decisions are made. For a first thesis prototype, GRU provides a solid, well-understood baseline. RL would be an excellent extension.

Q: Why use FlexRIC specifically instead of OpenRAN Gym or SD-RAN?

A: FlexRIC was chosen for three reasons: (1) It has native NS-3 E2 interface support — the NS-3 mmWave module already has FlexRIC E2 integration code written and tested by the same research groups (EURECOM, NYU). Using OpenRAN Gym or SD-RAN would require writing or adapting the NS-3 E2 integration from scratch. (2) FlexRIC is lightweight and easy to compile — no Docker required, runs as a single process. (3) FlexRIC has active documentation and community support from EURECOM. OpenRAN Gym is better for ML training pipelines but adds complexity for our simulation-focused use case.

Q: What would change if you replaced Random Walk mobility with a realistic urban model?

A: Random Walk 2D (our current model) moves UEs in random directions with random speeds, bounded within a rectangle. Real urban pedestrians follow streets, go around buildings, and have purposeful destinations. A realistic model (like Gauss-Markov or SUMO-based waypoint mobility) would produce more structured movement patterns. Expected changes: (1) Fewer random cell-boundary crossings → fewer total HOs. (2) More predictable movement → GRU predictions might be more accurate (consistent patterns to learn). (3) Some UEs might follow streets that run parallel to cell boundaries → could cause sustained cell-edge situations → more PPs. Overall the PP rate would likely be similar but the absolute HO count would decrease.

Q: What is the computational cost of running your system?

A: The full system requires: (1) One machine capable of compiling and running NS-3 (any modern quad-core Linux system, ~8GB RAM recommended). (2) Python 3.8+ with TensorFlow, Flask, InfluxDB client (all pip installable). (3) Docker for InfluxDB and Grafana containers. (4) ~3-6 hours per 60-120 second simulation run. CPU usage during simulation: NS-3 uses 1 core at ~100%, gru_xapp.py uses ~10% (model inference is fast), FlexRIC uses ~5%, sim_data_pusher.py uses ~5%. Total disk space for one simulation run: ~100-500MB of trace files. No GPU is needed for inference (our model is small enough for CPU inference with acceptable latency).

Q: What is the advantage of separating the GRU service from the xApp process?

A: Separating the GRU service (Python/Flask) from the xApp (C/FlexRIC) has four key advantages: (1) Independent updates: the GRU model can be replaced with a new version by restarting gru_xapp.py without recompiling the C xApp. (2) Language flexibility: ML in Python is far more practical than ML in C — TensorFlow has no official C API. (3) Independent testing: the GRU service can be tested with curl commands independently of the full simulation stack. (4) Scalability: in future, the GRU service could be deployed on a separate server or as a microservice, with the xApp calling it over a real network rather than localhost. The only downside is the HTTP latency overhead (<5ms on localhost), which is acceptable for near-RT RIC operations.

Q: Could your system work with a real 5G network instead of NS-3?

A: Yes, with modifications. The key components that would need to change: (1) NS-3 would be replaced by real O-RAN-compliant gNBs (e.g., from O-RAN Alliance compliant vendors) that support the E2 interface. (2) FlexRIC would connect to the real gNBs via SCTP on a real network instead of loopback. (3) The GRU model would need to be retrained on real SINR/RSRP measurements from the target deployment environment (different channel characteristics from NS-3 simulation). (4) The cooldown and A3 thresholds might need tuning for the actual cell density and UE speeds. The xApp code (best2.c) and gru_xapp.py would require minimal changes since they interface with FlexRIC's standard API and HTTP respectively.

Category 6: Future Work and Improvements (6 Questions)

Q: What is the most important improvement you would make to this system?

A: The most impactful improvement would be adding a **baseline comparison**: running the same simulations with pure A3 (no GRU, no ML) and comparing PP rates directly. Currently we can only compare to literature values, not our own system's A3-only mode. This would definitively quantify the GRU's contribution. The second most important improvement would be increasing the number of simulation runs (from 3 to 20+) to compute statistically valid mean PP rates with confidence intervals. Third, adding online learning capability — updating the GRU weights during simulation based on observed outcomes — would make the system adaptive to changing conditions.

Q: How could you extend this to support load balancing in addition to handover?

A: Load balancing (LB) ensures no single cell is overloaded while others are underused. An LB extension would add a second xApp (or extend best2.c) that monitors cell load (number of connected UEs per cell) via KPM reports. When a cell exceeds a load threshold (e.g., >15 UEs), the LB xApp identifies 'movable' UEs — those near a cell boundary with acceptable SINR on a less-loaded neighbor — and triggers handovers to redistribute load. The GRU could be extended to incorporate cell_load as an additional input feature, allowing the model to consider both signal quality AND cell congestion when recommending handover targets.

Q: How would you deploy this system in a production O-RAN network?

A: Production deployment would involve: (1) Containerize all components — xApp as a Docker container deployable via SMO's O2 interface. (2) Replace Flask with a production WSGI server (gunicorn + nginx) for the GRU service. (3) Use the A1 interface to push model updates from a central training server without service restart. (4) Add authentication/authorization to the GRU service HTTP endpoint. (5) Implement high availability: multiple FlexRIC instances with failover. (6) Instrument with proper monitoring (Prometheus metrics, not just InfluxDB). (7) Retrain the GRU model on real network data from the target deployment. (8) Conduct field testing with a small pilot deployment before full rollout.

Q: What other ML architectures could work better than GRU for this problem?

A: Several architectures are worth exploring: (1) **Transformer with attention**: The self-attention mechanism would allow the model to identify which of the 10 samples is most informative for the prediction, potentially outperforming GRU on complex patterns. (2) **Temporal Convolutional Network (TCN)**: 1D convolutions over the time axis can capture local patterns efficiently and train faster than RNNs. (3) **LSTM with attention**: Classic LSTM enhanced with an attention layer to weight the importance of each timestep. (4) **Graph Neural Network**: Model the 7-cell layout as a graph where cells are nodes and inter-cell relationships are edges — captures cell topology that GRU ignores. GRU was chosen as a solid, well-understood starting point; the above are natural extensions for a follow-up research paper.

Q: How would you handle UEs moving at vehicular speed (100 km/h) instead of pedestrian?

A: At 100 km/h (~28 m/s), a UE moves 28 meters per second. With a 0.05s KPM interval, the UE moves 1.4 meters between samples — covering a 7-meter range in the 10-sample window. At 28 GHz, the channel

coherence time is on the order of milliseconds at vehicular speed, meaning channel conditions change too fast for our 0.05s sampling. Adaptations needed: (1) Reduce KPM interval to 0.01s or 0.005s for faster sampling. (2) Reduce cooldown from 5.0s to 0.5-1.0s (vehicular UE moves out of cell in seconds). (3) Retrain GRU on vehicular-speed data (patterns are different from pedestrian). (4) Consider a separate model or branch for vehicular vs pedestrian UEs.

Q: What are the ethical or social implications of AI-controlled handovers in 5G?

A: AI-controlled handovers raise several considerations: (1) **Fairness:** A GRU model trained on certain UE patterns might perform poorly for UEs with unusual mobility (e.g., wheelchair users, cyclists) that are underrepresented in training data. (2) **Explainability:** Regulators may require network operators to explain why a specific handover decision was made — GRU is somewhat opaque, unlike rule-based A3. (3) **Reliability:** A ML system that fails catastrophically (e.g., model corruption) could affect thousands of users. The fallback to pure A3 mitigates this. (4) **Privacy:** UE trajectory data used to train the model could reveal location information — proper anonymization is essential. These considerations motivate combining ML with robust fallback mechanisms as we have done.

Category 7: Tricky and Hard Questions (6 Questions)

Q: Your PP detection checks within 5 seconds. But your cooldown is also 5 seconds. Isn't that circular?

A: This is a perceptive question. The cooldown and the PP detection threshold are indeed both set to 5.0 sim-seconds, but they serve different purposes and work at different stages. The **cooldown** is a proactive prevention mechanism in the xApp — it prevents the xApp from even evaluating a new handover for 5 seconds after the last one. The **PP detection** in post-processing is a retrospective measurement — it flags any HO that returned to the previous cell within 5 seconds. The two work together: if the cooldown is perfectly effective, PP detection would never flag anything (since no new HO can be issued within 5 seconds). The PPs we observe (4-10 per run) occur in edge cases where the channel degrades so severely that a new A3 trigger fires just after the 5-second cooldown expires — the cooldown expired but the UE is still near the original cell boundary.

Q: The GRU recommends a cell, but what if that cell is at full capacity?

A: This is a valid limitation of the current design. The GRU model only considers SINR/RSRP quality metrics and UE velocity — it does NOT consider cell load (how many UEs are already connected). It is theoretically possible that the recommended cell is already serving 20 UEs and has no scheduling resources available. In our NS-3 simulation with 20 UEs total across 7 cells, this is unlikely since cells average ~3 UEs each. But in a real deployment with higher UE density, this matters. The fix would be to add `cell_load` as a feature to the GRU model, or to add a post-processing step in the xApp that checks cell load (available from KPM reports) before accepting the GRU recommendation.

Q: If the GRU model gives 85% confidence for Cell 2 but only 14% for Cell 3 — when should you NOT trust it?

A: Several situations where the 85% confidence should not be trusted: (1) **Training distribution mismatch:** If the current UE movement pattern is fundamentally different from anything in the training data (e.g., a UE moving in a straight line at high speed vs the random-walk training data), the model may be confidently wrong — 85% confidence does not guarantee correctness. (2) **Recently corrupted features:** If the SINR measurements are noisy or missing due to a temporary blockage, the GRU processes bad input and outputs unreliable predictions. (3) **Near cell boundaries:** When 3+ cells have similar SINR, the true best cell is ambiguous — even 85% may be unreliable. A robustness check would be: if the A3 condition does not also agree with the GRU's choice, do not execute the handover.

Q: What is the theoretical maximum accuracy your system could achieve, and why can't it be 100%?

A: The theoretical maximum accuracy (minimum achievable PP rate) is NOT 100% for fundamental physical reasons. A UE following a random walk can genuinely need to hand back to the original cell after 5.1 seconds (just after cooldown expires) if its random trajectory turns around. No algorithm can predict a truly random walk perfectly — only a UE with a deterministic, predictable path could guarantee 100% correct decisions. The remaining ~3% PP rate in our system represents cases where the UE's random trajectory was genuinely ambiguous or reversed direction shortly after handover. Improving the GRU model, extending the rolling window, or predicting future position might reduce this to ~1%, but eliminating it entirely is physically impossible for random mobility.

Q: Why is your simulation rate 0.005 sim-seconds/real-second and not faster? Could you speed it up?

A: The 0.005 ratio is determined by the slowest component in the simulation pipeline. Profiling would show which step is the bottleneck: likely the mmWave channel computation (ray-tracing based propagation at 28 GHz for 20 UEs $\times$ 7 cells = 140 link computations per timestep) or the E2 message encoding/decoding (ASN.1 is computationally expensive). Speed improvements are possible: (1) Use a simplified channel model (e.g., log-distance instead of full ray-tracing): 10-50 $\times$ speedup but reduced accuracy. (2) Run on a multi-core server with NS-3 parallel mode. (3) Reduce number of UEs (20 $\rightarrow$ 10 would roughly halve simulation time). (4) Use ns3-ai (machine learning integration without HTTP overhead). (5) Compile NS-3 with higher optimization flags. For our thesis purposes, the current speed is acceptable since simulations run overnight.

Q: How would you validate that your GRU model is not simply learning the cell layout rather than UE movement?

A: This is a critical question about generalization. If the GRU simply memorizes 'UEs at position X always use Cell Y' based on the fixed cell layout, it would fail in a different cell arrangement. To test this: (1) Train on one cell layout (e.g., 7 cells in configuration A) and test on a different layout (configuration B with different cell positions). If accuracy drops significantly, the model learned the layout, not the movement patterns. (2) Use cross-validation: train on random seeds 1-8, test on seeds 9-10. Since seeds produce different UE trajectories in the same layout, this tests generalization to new movements. (3) Visualize attention weights (if using an attention mechanism): do the important features relate to movement trends (rising/falling SINR) or absolute values (fixed geographic positions)?

Glossary — Every Technical Term Explained

This glossary defines every technical term used in this thesis in plain language. Use this as a quick reference when preparing for your defense — the committee may ask you to define any of these terms.

5G NR: 5th Generation New Radio — the latest mobile network standard, operating at both sub-6 GHz and mmWave frequencies above 24 GHz.

A3 Event: A 3GPP-defined handover trigger. Fires when a neighboring cell's SINR exceeds the serving cell's SINR by a defined margin (in our case, 2.0 dB). Named 'A3' in the 3GPP measurement event classification.

Adam Optimizer: An adaptive learning rate optimization algorithm used during neural network training. It adapts the learning rate for each parameter individually, combining ideas from RMSprop and momentum. Standard choice for training GRU/LSTM models.

ASN.1: Abstract Syntax Notation One — a data serialization format used for encoding E2AP messages. It is compact and efficient but computationally expensive to encode/decode, contributing to simulation slowness.

Categorical Cross-Entropy: The loss function used to train our GRU model. It penalizes the model heavily when it is confidently wrong about the correct cell class.

Cooldown Period: A 5.0 sim-second window after each handover during which no new handover is evaluated for that UE. Prevents ping-pong by blocking rapid re-evaluation.

Docker: A container platform that packages software (like InfluxDB and Grafana) with all its dependencies so it runs identically on any machine. Used for our 2D GUI stack.

E2 Interface: The O-RAN interface connecting the near-RT RIC (FlexRIC) to E2 Nodes (NS-3 gNBs). Carries KPM reports (monitoring) and RC commands (control) over SCTP port 36421.

E2AP: E2 Application Protocol — the signaling protocol used over the E2 interface. Handles subscription setup, indication delivery, and control message exchange.

E2SM-KPM: E2 Service Model for Key Performance Measurements. Defines the format of monitoring reports (SINR, RSRP, cell load, etc.) sent from E2 Nodes to the RIC.

E2SM-RC: E2 Service Model for RAN Control. Defines the format of control commands (like handover triggers) sent from the RIC to E2 Nodes.

FastAPI: A modern Python web framework for building APIs. Used in controller.py to expose the 3D GUI backend API on port 8001.

Flask: A lightweight Python web framework used in gru_xapp.py to create the HTTP inference server on port 5000.

FlexRIC: An open-source near-RT RIC implementation from EURECOM. Implements the O-RAN E2 interface and provides an SDK for writing xApps in C/C++/Python.

GRU: Gated Recurrent Unit — a type of recurrent neural network with two gates (update and reset) that solve the vanishing gradient problem. Used to predict the best handover target cell from a sequence of SINR measurements.

gNB: Next Generation Node B — the 5G base station. In our simulation, we have 7 gNBs forming the cell layout.

Grafana: An open-source data visualization tool. Reads from InfluxDB and displays real-time dashboards. Accessible at port 8000 in our system.

Handover (HO): The process of transferring an active UE connection from one cell (gNB) to another. Needed when a UE moves and the current cell's signal degrades.

Hysteresis: An extra margin added to handover thresholds to prevent oscillation. Our $A3_HYSTERESIS = 1.0$ dB means the neighbor must be 1 dB better than just the offset.

InfluxDB: A time-series database optimized for storing timestamped data (like SINR measurements over time). Runs as a Docker container on port 8086.

KPM Report: Key Performance Measurement report. Sent by NS-3 to FlexRIC every 0.05 sim-seconds per UE, containing SINR, RSRP, cell ID, and other metrics.

LSTM: Long Short-Term Memory — a type of RNN with three gates that solves the vanishing gradient problem. Similar to GRU but more complex. We chose GRU for its lower computational cost.

mmWave: Millimeter Wave — radio frequencies above 24 GHz (wavelength ~ 1 -10mm). Used in 5G for high bandwidth. Short range and easily blocked.

Near-RT RIC: Near Real-Time RAN Intelligent Controller — runs xApps on 10ms–1s timescales. FlexRIC is our near-RT RIC implementation.

Non-RT RIC: Non-Real-Time RIC — runs rApps on >1 second timescales for long-term optimization and policy generation.

NS-3: Network Simulator 3 — a free, open-source discrete-event network simulator written in C++. Used to simulate our 7-cell mmWave network with 20 UEs.

O-RAN Alliance: Industry consortium that defines open specifications for O-RAN architecture, founded in 2018 by AT&T, Deutsche Telekom, NTT DoCoMo, Orange, and China Mobile.

O-RAN: Open Radio Access Network — a mobile network architecture with open, standardized interfaces that enables multi-vendor deployments and custom software intelligence.

Ping-Pong (PP): A handover event where a UE moves from Cell A to Cell B, then back to Cell A within 5 simulation-seconds. Indicates an unnecessary handover that wasted resources.

Random Walk 2D: The NS-3 mobility model used for UEs. Each UE moves in a random direction at a random speed, bouncing off boundaries. Produces unpredictable but statistically uniform coverage of the simulation area.

RC Command: RAN Control command — an E2SM-RC message sent by the xApp to NS-3 commanding a specific UE to perform a handover to a target cell.

RIC: RAN Intelligent Controller — software platform for running control and optimization applications (xApps/rApps) in an O-RAN network.

Rolling Window: A fixed-size buffer (10 samples in our system) that always holds the most recent N measurements. When a new sample arrives, the oldest one is dropped.

RNN: Recurrent Neural Network — a neural network with feedback connections, enabling it to process sequences. Suffers from vanishing gradient for long sequences.

RNTI: Radio Network Temporary Identifier — a unique ID assigned to each UE by its serving cell. Used to identify which UE a KPM report or RC command refers to.

RSRP: Reference Signal Received Power — measures the absolute power level of a cell's reference signal in dBm. Indicates signal strength but not quality.

SCTP: Stream Control Transmission Protocol — a transport protocol used for the E2 interface. Similar to TCP but supports multi-homing and message boundaries. Port 36421.

SINR: Signal-to-Interference-plus-Noise Ratio — the ratio of desired signal power to interference + noise power, in dB. Primary metric for channel quality.

SMO: Service Management and Orchestration — the O-RAN management layer that deploys xApps, manages configuration, and interfaces with cloud infrastructure.

Softmax: A mathematical function that converts a vector of raw scores into probabilities that sum to 1.0. Used as the final activation in our GRU model to produce cell probabilities.

UE: User Equipment — any device connected to the mobile network. In NS-3, we simulate 20 UEs moving randomly.

Vendor Lock-in: The situation where an organization is dependent on a single vendor and cannot easily switch to alternatives. O-RAN was created specifically to eliminate vendor lock-in in mobile networks.

Vite: A fast JavaScript build tool and development server. Used to serve the 3D visualization frontend on port 3001.

xApp: A small application running inside the near-RT RIC that performs a specific network function using E2 interface subscriptions and control commands. Our xApp (`xapp_handover_gru`) handles handover optimization.

System Troubleshooting — Common Problems and Solutions

This chapter describes the most common problems encountered when running the simulation and their solutions. Understanding these issues demonstrates deep system knowledge and may come up during your defense.

14.1 Problem: Simulation runs but no handovers happen

- Symptom: NS-3 runs, FlexRIC is running, but `decision_log.csv` shows zero entries.
- Cause 1: xApp did not start before NS-3. The subscription was never established. Solution: Follow `gru.sh` startup order exactly.
- Cause 2: GRU service (`gru_xapp.py`) crashed at startup. xApp receives no HTTP responses and falls back, but if A3 never triggers due to low UE speeds, no HO fires.
- Cause 3: `A3_OFFSET+HYSTERESIS` is too high. No neighbor ever exceeds threshold. Check `/tmp/flexric.log` for 'A3 threshold not met' messages.
- Diagnosis: Run `'tail -f /tmp/flexric.log'` during simulation. Should see 'KPM indication received' lines every few seconds.

14.2 Problem: `gru_xapp.py` fails to load model

- Symptom: 'Error loading model' in Python output. GRU service exits immediately.
- Cause: The `.h5` or `.keras` model file is missing or in the wrong directory.
- Solution: Ensure the model file is in the same directory as `gru_xapp.py`, or update the model path in the `load_model()` call.
- Also check: TensorFlow version compatibility. A model saved with TF 2.10 may not load in TF 2.5.

14.3 Problem: FlexRIC cannot bind to port 36421

- Symptom: 'bind: Address already in use' in `/tmp/flexric.log`.
- Cause: A previous FlexRIC instance is still running.
- Solution: Run `kill_sim.sh` to clean up all processes. Then check: `'ss -ltn | grep 36421'` to confirm port is free.

14.4 Problem: High ping-pong rate (>10%)

- Symptom: `decision_log.csv` shows many `is_pingpong=True` entries.
- Cause 1: Cooldown too short. Increase `COOLDOWN_TIME` in `best2.c` (requires recompile).
- Cause 2: A3 threshold too low. Increase `A3_OFFSET` to 1.5 dB.

-
- Cause 3: GRU model is performing poorly. Check model accuracy on validation set.
 - Cause 4: UEs moving too fast (check mobility speed parameter in `gru_scenario.cc`).

14.5 Problem: Simulation very slow (> 5 hours for 60s)

- Normal pace: ~3.3 hours for 60 sim-seconds. If significantly slower, investigate:
- Check 1: Is `gru_xapp.py` spending excessive time on model inference? Profile with Python cProfile.
- Check 2: Is InfluxDB overwhelmed? Reduce push rate in `sim_data_pusher.py`.
- Check 3: NS-3 compiled in debug mode? Recompile with `--build-profile=optimized`.
- Check 4: Other CPU-intensive processes running? Check with 'htop'.

14.6 Problem: Grafana shows no data at port 8000

- Check 1: Is Docker running? 'docker ps' should show influxdb, grafana, nginx containers.
- Check 2: Is `sim_data_pusher.py` running? It feeds InfluxDB.
- Check 3: Is the InfluxDB datasource configured in Grafana? Check Grafana datasource settings.
- Check 4: Check InfluxDB health: 'curl http://localhost:8086/health' should return {status: pass}.

Mathematical Background — Key Equations Explained

15.1 Path Loss Model for mmWave

The NS-3 mmWave module uses a realistic path loss model for 28 GHz propagation. Path loss determines how much signal power is lost as the radio wave travels from the gNB antenna to the UE. The basic model is:

$$PL(d) = PL(d_0) + 10 \times n \times \log_{10}(d/d_0) + X_{\sigma} \text{ [dB]}$$

Where:

- **PL(d)**: Path loss at distance d (in dB)
- **PL(d₀)**: Free-space path loss at reference distance d₀ = 1 meter
- **n**: Path loss exponent (~2 in free space, 3-4 in urban areas)
- **X_σ**: Shadow fading term — a random variable (Gaussian) representing signal blocking from obstacles
- **At 28 GHz**: PL(1m) ≈ 61.4 dB (much higher than 2.4 GHz which is ~40 dB)

15.2 SINR Calculation

SINR (Signal-to-Interference-plus-Noise Ratio) is calculated for each UE from each cell:

$$SINR_i = P_{\text{received_from_cell_i}} / (\text{Sum}(P_{\text{interference_from_other_cells}}) + P_{\text{noise}})$$

$$SINR_i \text{ [dB]} = 10 \times \log_{10}(SINR_i \text{ [linear]})$$

In our simulation with 7 cells, for a UE connected to Cell 0:

- **Signal**: power received from Cell 0 (serving cell)
- **Interference**: sum of power received from Cells 1-6 (all neighbors transmit simultaneously)
- **Noise**: thermal noise power (kTB, where k=Boltzmann's constant, T=temperature, B=bandwidth)

15.3 GRU Gate Equations

For completeness, the mathematical equations of the GRU cell:

$$\text{Update gate: } z_t = \text{sigmoid}(W_z \times x_t + U_z \times h_{t-1} + b_z)$$

$$\text{Reset gate: } r_t = \text{sigmoid}(W_r \times x_t + U_r \times h_{t-1} + b_r)$$

$$\text{Candidate: } h_{t\sim} = \tanh(W_h \times x_t + U_h \times (r_t \blacksquare h_{t-1}) + b_h)$$

$$\text{New state: } h_t = (1 - z_t) \blacksquare h_{t-1} + z_t \blacksquare h_{t\sim}$$

Where: x_t = input at timestep t, h_t = hidden state, W/U = weight matrices, b = bias vectors, $\text{sigmoid} = 1/(1+e^{-x})$, $\tanh$ = hyperbolic tangent, $\blacksquare$ = element-wise multiplication.

Note: You do NOT need to memorize these equations for your defense. But if a committee member asks 'how does GRU work mathematically,' you can refer to them and explain that z controls the blend of old vs new state, and r controls how much old memory is relevant to computing the new candidate.

15.4 Softmax Function

$$\text{softmax}(x_i) = e^{(x_i)} / \text{Sum}(e^{(x_j)} \text{ for all } j)$$

For our 7-cell output: if raw scores are [1.2, 0.3, 3.8, 0.1, 0.5, 0.2, 0.4], softmax converts these to probabilities summing to 1.0. The largest score (3.8 for Cell 2) gets the highest probability after exponentiation.

15.5 Ping-Pong Rate Formula

$$\text{PP Rate (\%)} = (\text{Number of Ping-Pong HOs} / \text{Total Handovers}) \times 100$$

$$\text{Accuracy (\%)} = 100\% - \text{PP Rate} = ((\text{Total HOs} - \text{PP HOs}) / \text{Total HOs}) \times 100$$

Simulation	Total HOs	PP HOs	PP Rate	Accuracy
sim006	334	4	$(4/334) \times 100 = 1.72\%$	98.28%
sim010	137	5	$(5/137) \times 100 = 3.65\%$	96.35%
sim011	309	10	$(10/309) \times 100 = 3.24\%$	96.76%

Technical Deep Dives — Extended Explanations

16.1 How NS-3 mmWave Channel Model Works

The NS-3 mmWave module uses the **3GPP TR 38.901 channel model**, which is the standard channel model for 5G NR. At 28 GHz, channel propagation is fundamentally different from 4G sub-6 GHz channels. Here is a step-by-step explanation of what happens when NS-3 computes the channel between a gNB and a UE:

- 1. Large-scale path loss:** Based on distance and LOS/NLOS (Line-of-Sight / Non-Line-of-Sight) determination. If no obstacle is between gNB and UE → LOS (lower path loss). Otherwise → NLOS (higher path loss).
- 2. Shadow fading:** A log-normal random variable (typically 4-8 dB standard deviation for mmWave) is added to account for large obstacles that cause random variations in received power.
- 3. Small-scale fading:** Multiple reflections and diffractions create many copies of the signal arriving at slightly different times and angles. Their combination causes rapid SINR fluctuations over small distances (~cm at 28 GHz).
- 4. Beamforming:** mmWave gNBs use phased arrays to focus signal energy in a narrow beam toward the UE (beamforming). NS-3 models beam alignment and updates it periodically.
- 5. Blockage:** mmWave signals are blocked by the human body, cars, and buildings. NS-3 can model dynamic blockage with moving obstacles.

The result of all this computation is a SINR value for each (gNB, UE) pair, updated every simulation timestep. This computation is the main reason the simulation is slow.

16.2 The E2AP Message Exchange in Detail

The E2AP (E2 Application Protocol) governs how messages are exchanged between FlexRIC and NS-3. Here is the complete sequence of E2AP messages during a typical simulation run:

Step	Message	Direction	Content
1	E2 Setup Request	NS-3 → FlexRIC	gNB ID, list of supported E2 service models (KPM, RC)
2	E2 Setup Response	FlexRIC → NS-3	RIC ID, list of accepted service models
3	RIC Subscription Request	FlexRIC → NS-3	KPM subscription: report interval=0.05s, UE IDs
4	RIC Subscription Response	NS-3 → FlexRIC	Subscription ID, admission: accepted

Step	Message	Direction	Content
5	RIC Indication (KPM)	NS-3 → FlexRIC	UE measurements: SINR, RSRP, cell IDs (every 0.05s)
6	RIC Control Request (RC)	FlexRIC → NS-3	UE ID + target cell ID for handover
7	RIC Control Acknowledge	NS-3 → FlexRIC	Confirmation that HO command was accepted
8	E2 Connection Update	Either direction	Heartbeat / keep-alive (periodic)

16.3 Why Softmax Probabilities Can Be Misleading (Calibration)

A common misconception is that if the GRU outputs 90% probability for Cell 2, it is correct 90% of the time. This is **only true if the model is well-calibrated**. Many neural networks (especially those trained with cross-entropy on imbalanced data) are **overconfident** — they output high probabilities even when wrong.

In our case, if the GRU was trained predominantly on clear-cut handover situations (UE clearly moving toward Cell 2), it may output high confidence even for ambiguous cell-edge situations. Calibration can be checked using a reliability diagram: plot the model's confidence (x-axis) against actual accuracy (y-axis). A perfectly calibrated model shows a diagonal line. Techniques like temperature scaling can improve calibration post-training without retraining.

16.4 The Role of Velocity in GRU Features

UE **velocity** is included as a feature because it helps the GRU distinguish between a UE that is moving quickly (and will soon leave the current cell) vs a stationary or slow-moving UE (that might have a momentary SINR dip due to channel fading, not actual movement). In NS-3, velocity is computed from the difference in UE position between consecutive KPM reports:

$$\text{velocity}_t = \sqrt{(\text{x}_t - \text{x}_{\{t-1\}})^2 + (\text{y}_t - \text{y}_{\{t-1\}})^2} / \text{delta}_t$$

A high velocity (>2 m/s) combined with rising SINR from a neighbor cell is a strong indicator that the UE is genuinely approaching that cell and a stable handover is warranted. A high velocity combined with fluctuating SINR from multiple cells equally might indicate the UE is at a high-interference corner, where the GRU should be more cautious.

16.5 Thread Safety in the C xApp

The xApp (best2.c) must handle KPM indications that arrive rapidly — potentially 20 indications per second for each of 20 UEs = 400 indications/second. If multiple indications are processed simultaneously in different threads, shared state (like the rolling window arrays and last_ho_time values) could be corrupted. FlexRIC's SDK handles this by calling the indication callback from a single thread, serializing all indication processing. This means our rolling window updates and cooldown checks are inherently thread-safe. The HTTP call to gru_xapp.py is blocking — the xApp waits for the response before processing the next indication. This is acceptable because the response is fast (<5ms on localhost) and we do not need to process 400

indications/second instantaneously (the 0.05s real-time gap between indications is $\gg$ 5ms).

16.6 How Grafana Reads from InfluxDB — Flux Query Language

Grafana panels are configured with **Flux** queries (InfluxDB's query language) to retrieve and visualize data. A typical query for a SINR panel:

```
from(bucket: "oran_metrics")
|> range(start: -5m)
|> filter(fn: (r) => r._measurement == "ue_sinr")
|> filter(fn: (r) => r.ue_id == "5")
|> aggregateWindow(every: 1s, fn: mean)
|> yield(name: "mean_sinr_ue5")
```

This query: reads from the 'oran_metrics' bucket, filters the last 5 minutes, selects the 'ue_sinr' measurement for UE #5, computes 1-second averages, and returns the result for plotting. Grafana refreshes this query every 5 seconds, creating a near-real-time live chart.

Pre-Defense Checklist and Examination Tips

17.1 One Week Before Defense

- Re-read all simulation results (Chapter 10) and memorize the key numbers: 334/4/1.72%, 137/5/3.65%, 309/10/3.24%
- Practice explaining the full system flow out loud, without notes, in under 3 minutes
- Make sure you can draw the architecture diagram on a whiteboard: NS-3 ↔ E2 ↔ FlexRIC ↔ xApp ↔ GRU service
- Know every file name, every port number, every threshold — the committee will test specifics
- Prepare 2-3 analogies you are comfortable with (GRU as short-term memory, O-RAN as Android, HO as relay race)
- Practice answering 'What would you do differently?' — have 3 honest improvements ready

17.2 Common Committee Traps — How to Avoid Them

Trap	What It Looks Like	How to Handle
False precision	Quoting 96.76% as absolute truth	Say: 'based on 3 runs with different seeds, our results consistently show ~3-4% PP rate'
Over-claiming	Saying 'our system is better than everything'	Say: 'our results are in the better range of ML-assisted HO literature'
Not knowing a number	Forgetting the A3 threshold	If you forget: 'I know the combined threshold is 2 dB — 1 dB offset plus 1 dB hysteresis'
Scope creep question	'Why didn't you implement X?'	'That is excellent future work. We focused on X because of time constraints. I have noted it as future work in Chapter 12.'
Gotcha question	'What if the model is wrong?'	'We have a fallback to pure A3 decision if the GRU service fails or returns low confidence.'
Simulation criticism	'Your results are just simulation'	'Yes — simulation allows controlled experiments impossible in real networks. Real deployment would require retraining on field data.'

17.3 Things to Say vs Things to Avoid

SAY:

- 'The system achieves 96.76% accuracy in sim011, meaning only 3.24% of handovers resulted in ping-pong.'
- 'We chose GRU over LSTM because GRU achieves comparable accuracy with fewer parameters.'

- 'The 5-second cooldown was determined empirically from our experiments.'
- 'FlexRIC is an open-source near-RT RIC from EURECOM that implements the O-RAN E2 interface.'
- 'A limitation of our work is that we only ran 3 simulations — more runs would improve statistical confidence.'

AVOID:

- 'I don't know' without any attempt. ALWAYS say what you DO know, then acknowledge the limit.
- 'The model just works.' — Always explain WHY it works.
- 'We got 96.76% which is great.' — Always contextualize against literature.
- 'NS-3 is the same as a real network.' — Always acknowledge simulation limitations.

17.4 The 5-Minute System Overview You Should Memorize

Practice delivering this overview in exactly 5 minutes — it answers the first question 'Can you briefly describe your system?' completely:

Key Point: 'Our system implements GRU-based handover optimization in an Open RAN environment. We simulate a 5G mmWave network with 7 base stations and 20 users using NS-3. NS-3 connects to FlexRIC — an open-source near-RT RIC — via the O-RAN E2 interface on SCTP port 36421. A custom xApp written in C (best2.c) subscribes to receive SINR and RSRP measurements every 0.05 simulation-seconds. When a user's signal degrades and a neighbor cell is 2 dB better (A3 event), the xApp sends the last 10 measurements to our GRU Python service on port 5000. The GRU model predicts which of the 7 cells will provide the best signal based on the movement pattern. The xApp then sends a handover command via E2 back to NS-3. An anti-ping-pong cooldown of 5 simulation-seconds prevents rapid oscillation. We ran three simulations of 60-120 seconds each, achieving ping-pong rates of 1.72-3.65%, corresponding to 96-98% accuracy — better than the 5-15% typical for A3-only systems.'

Additional Defense Questions — Supplementary Q&A;

These additional questions cover aspects of the system that a thorough committee might ask beyond the main Q&A; in Chapter 12. Each answer is written to demonstrate depth of understanding.

18.1 Questions About the Project Team Division

Q: Who built what in this project? What was your specific contribution?

A: The project was divided between two students. Omar Farouk (this thesis) was responsible for: (1) Setting up the NS-3 mmWave simulation environment (`gru_scenario.cc`) with 7 cells and 20 UEs, (2) Integrating NS-3 with FlexRIC via the E2 interface, (3) Developing and compiling the xApp (`best2.c`) including A3 logic, rolling window, cooldown timer, and HTTP communication with the GRU service, (4) Setting up the 2D GUI (InfluxDB + Grafana via Docker), (5) Implementing `sim_data_pusher.py` for real-time data visualization, (6) Running all three simulations and analyzing results, (7) Developing the 3D GUI controller and frontend. Fares was responsible for: training the GRU model, preparing the training dataset from simulation data, implementing `gru_xapp.py` (the Flask inference service), and model validation. The division follows a natural system split: Fares owns the ML model, Omar owns the O-RAN integration.

Q: How did you and Fares coordinate the interface between the C xApp and the Python GRU service?

A: We defined a clear API contract early in the project: the JSON format for `/predict` requests and responses. Omar's xApp (`best2.c`) would always send exactly: a JSON object with `ue_id` (int), `serving_cell` (int), and a 'window' array of 10 objects, each with `sinr`, `rsrp`, `velocity`, and `cell_sinrs` (array of 7 floats). Fares's `gru_xapp.py` would always return: `cell_id` (int 0-6), `confidence` (float 0-1), and `all_probs` (array of 7 floats). We wrote test scripts to validate the interface before integrating. This approach — defining the API contract first and implementing independently — is standard software engineering practice for team projects.

18.2 Questions About 5G Standards and 3GPP

Q: What is 3GPP and what role does it play in your system?

A: 3GPP (3rd Generation Partnership Project) is the international standards body that defines mobile network specifications — 3G (UMTS), 4G (LTE), and 5G (NR). The O-RAN Alliance builds on top of 3GPP specifications. Specifically in our system: (1) The A3 event is defined in 3GPP TS 36.331/38.331 as a measurement event for handover triggering. (2) The mmWave channel model we use (3GPP TR 38.901) is a 3GPP technical report. (3) The RRC (Radio Resource Control) procedure that NS-3 executes to perform a handover is specified in 3GPP TS 38.331. (4) The RNTI assignment is specified in 3GPP TS 38.321. O-RAN is an industry initiative that extends 3GPP specifications with open interfaces — it does not replace 3GPP but complements it.

Q: What is the TTT (Time-to-Trigger) in standard 3GPP handovers, and do you use it?

A: TTT (Time-to-Trigger) is a timer in standard 3GPP A3 handovers. The A3 condition must be true continuously for the entire TTT duration before the handover is triggered. This prevents triggering on momentary SINR spikes. Typical TTT values are 40ms, 64ms, 80ms, 256ms, up to 1024ms. In our implementation, we do NOT explicitly implement a TTT timer in the traditional sense. Instead, the 10-sample rolling window feeding the GRU serves a similar purpose — the GRU must see a sustained trend over 0.5 sim-seconds before recommending a handover. Additionally, the 0.05s KPM interval means the A3 check happens every 0.05 sim-seconds, so any single-sample spike would only trigger if the GRU also confirms it. The cooldown further suppresses short-duration reversals. This multi-layer approach is arguably more robust than a simple TTT timer.

Q: What is the difference between intra-frequency and inter-frequency handovers?

A: Intra-frequency handovers occur between cells operating on the same carrier frequency. Inter-frequency handovers occur between cells on different frequencies. In our simulation, all 7 cells operate on the same 28 GHz carrier frequency — so all handovers are intra-frequency. This simplifies the implementation significantly: measurements for all cells are directly comparable since they use the same frequency. Inter-frequency handovers require additional measurements (the UE must briefly scan other frequencies), adding complexity and delay. For a thesis scope, intra-frequency is the correct starting point. A real 5G network might have cells on multiple bands (e.g., 28 GHz mmWave AND 3.5 GHz sub-6 GHz), requiring inter-frequency handovers — a natural extension.

18.3 Questions About Software Engineering

Q: Why did you use HTTP (REST) instead of a more efficient IPC mechanism for xApp ↔ GRU communication?

A: HTTP was chosen for pragmatic reasons. The alternative would be a shared-memory interface, Unix socket, or gRPC. Each has tradeoffs: HTTP advantages: universally supported, easy to debug (curl can test the endpoint), language-agnostic (any language can implement either side), and well-documented. HTTP disadvantages: ~1ms overhead per request vs ~0.01ms for shared memory. In our case, the overhead is acceptable: we make one HTTP call per handover decision, and handovers happen at most 20 times per sim-second (once per UE per KPM report). Even $1\text{ms} \times 20\text{ calls} = 20\text{ms}$ real-time per sim-second is negligible compared to the 3-hour total simulation time. If we needed to process 1000s of requests per second, gRPC or shared memory would be justified.

Q: How is the simulation reproducible? What ensures two researchers get the same result?

A: Reproducibility is ensured by: (1) NS-3 random seeds — the simulation uses fixed RngSeed and RngRun values set in gru_scenario.cc. Same seed = same UE mobility patterns and channel realizations. (2) Fixed GRU model weights — the trained .h5 file is version-controlled and its weights are deterministic. (3) Fixed threshold parameters (A3_OFFSET=1.0, COOLDOWN=5.0, etc.) are compile-time constants in best2.c. (4) The git repository captures all code at each simulation run via git commits. Limitations: floating-point arithmetic can differ between CPU architectures, and the HTTP timing between xApp and GRU service

introduces slight non-determinism in which UE is evaluated first when multiple indications arrive simultaneously. In practice, results are reproducible within ± 1 handover event across machines.

18.4 Questions About the Research Contribution

Q: What is the novel contribution of your thesis?

A: The novel contributions of this thesis are: (1) **Integration of GRU prediction with the O-RAN E2 interface** — we demonstrate, end-to-end, how a GRU model can be deployed as an O-RAN xApp receiving live KPM reports and issuing RC control commands. Most ML-for-handover papers stop at algorithmic comparison; we build the full working system. (2) **Anti-ping-pong cooldown combined with ML prediction** — the combination of GRU trend prediction AND explicit cooldown has not been widely studied together. Our results (1.72-3.65% PP rate) demonstrate the effectiveness of this combination. (3) **Open-source, reproducible simulation platform** — the full stack (NS-3 + FlexRIC + xApp + GRU service + 2D/3D GUIs) is available and reproducible. This can serve as a platform for future O-RAN xApp research.

Q: What would you publish from this thesis, and in which venue?

A: The most publishable contribution is the end-to-end O-RAN xApp implementation with quantitative ping-pong reduction results. Suitable venues: (1) IEEE Wireless Communications Letters or IEEE Communications Letters (short papers, 5 pages) — focus on the GRU+cooldown combination and PP rate results. (2) IEEE VTC (Vehicular Technology Conference) — appropriate for handover optimization work. (3) IEEE ICC (International Conference on Communications) — broader audience. For publication, we would need to: add a baseline comparison (pure A3, no GRU), increase runs to 10-20 for statistical confidence, and potentially add comparison against LSTM and pure A3 in the same simulation environment.

Worked Examples — Walking Through Real Scenarios

This chapter walks through concrete scenarios step by step, showing exactly what happens inside the system. These examples will help you explain the system to the committee with specific, concrete details.

19.1 Scenario A: A Clean Successful Handover

UE #5 is moving from Cell 2 toward Cell 4. Here is exactly what happens:

Background: Simulation time = 25.00 sim-seconds. UE #5 is connected to Cell 2 (SINR = +2.0 dB). It has been in Cell 2 for 8 sim-seconds. Cooldown has expired.

1. **t=25.00:** NS-3 generates KPM report for UE #5. Cell 2 SINR = +2.0 dB (serving). Cell 4 SINR = +1.5 dB (neighbor). A3 check: $1.5 < 2.0 + 2.0 = 4.0$. **A3 NOT triggered.** No action.
2. **t=25.05:** UE has moved 0.15m toward Cell 4. Cell 2 SINR = +1.8 dB. Cell 4 SINR = +2.5 dB. A3 check: $2.5 < 1.8 + 2.0 = 3.8$. **A3 NOT triggered.**
3. **t=25.30:** UE has moved another 1.5m closer to Cell 4. Cell 2 SINR = -1.0 dB. Cell 4 SINR = +3.5 dB. A3 check: $3.5 > -1.0 + 2.0 = 1.0$ dB. **A3 TRIGGERED! Cell 4 is 4.5 dB better.**
4. **Rolling window:** 10 samples for UE #5 are ready. The window shows consistent trend: Cell 4 SINR rising for last 6 samples, Cell 2 falling.
5. **HTTP POST to :5000/predict:** xApp sends window with 10 samples.
6. **GRU processes:** Update gate retains trend memory. Reset gate confirms Cell 4 trend relevant. Output: Cell 4 = 91% probability, Cell 2 = 4%, others = 5%.
7. **xApp validates:** GRU says Cell 4. A3 also says Cell 4. Cooldown check: $25.30 - 17.00 = 8.3s > 5.0$. OK.
8. **RC command sent:** xApp sends E2SM-RC to FlexRIC → NS-3: 'Move UE #5 to Cell 4'.
9. **NS-3 executes:** RRC Reconfiguration at t=25.31. UE #5 now attached to Cell 4.
10. **State update:** xApp records last_ho_time[5] = 25.30. Cooldown starts.
11. **Result:** UE #5 stays in Cell 4 for the next 12 sim-seconds. No ping-pong. SUCCESS.

19.2 Scenario B: A Ping-Pong Handover Despite Protections

UE #12 is at the exact boundary between Cell 1 and Cell 3, executing a curved random walk that briefly enters Cell 3's territory then returns. Here is what happens:

1. **t=40.00:** UE #12 in Cell 1 (SINR = +0.5 dB). Cell 3 SINR = +4.0 dB. A3 triggered: $4.0 > 0.5 + 2.0 = 2.5$. GRU window shows rising Cell 3 trend. GRU: 88% Cell 3.
2. **HO executed at t=40.00:** UE #12 moves to Cell 3. last_ho_time[12] = 40.00.
3. **t=40.5–44.9:** Cooldown active for UE #12. No handover evaluation regardless of SINR.

4. **t=44.9:** UE #12's random walk has curved back toward Cell 1. Cell 3 SINR = -2.0 dB (degrading). Cell 1 SINR = +5.0 dB. The UE has walked 15m back.
5. **t=45.05:** Cooldown expires: $45.05 - 40.00 = 5.05 > 5.0$. Evaluation resumes.
6. **A3 check:** Cell 1 SINR = +5.0 dB > Cell 3 SINR (-2.0 dB) + 2.0 = 0.0 dB. TRIGGERED.
7. **GRU window:** Shows clear declining Cell 3 SINR for last 6 samples, Cell 1 rising. 92% Cell 1.
8. **RC command:** UE #12 handed over to Cell 1 at t=45.05.
9. **PP detection:** This is HO from Cell 3 → Cell 1 at t=45.05, after HO from Cell 1 → Cell 3 at t=40.00. Time difference = 5.05 sim-seconds > 5.0. TECHNICALLY NOT a ping-pong by our 5-second definition. Edge case!

***Note:** This scenario illustrates why a 5-second cooldown is imperfect. The UE genuinely needed to return to Cell 1. The 5.05-second gap barely avoids the PP classification. If the UE had curved back faster and the return HO happened at 44.9s (within cooldown), the return handover would have been blocked and the UE would suffer poor SINR in Cell 3 until the cooldown expired.*

19.3 Scenario C: GRU Falls Back to A3 (Insufficient Window)

UE #18 just started the simulation at t=0.0. It takes 0.5 sim-seconds to accumulate 10 samples. During this warmup period:

1. **t=0.05–0.40:** xApp receives KPM reports for UE #18. Window has 1-8 samples. Each time A3 is checked: if triggered, xApp sends HTTP POST to :5000/predict.
2. **gru_xapp.py response:** Returns HTTP 400 — {error: 'insufficient_data', samples: 7}.
3. **xApp fallback:** Uses pure A3 decision. Picks neighbor with highest SINR above threshold.
4. **t=0.50:** Window now has 10 samples. GRU is ready to predict normally.
5. **From t=0.50 onward:** Normal GRU predictions are used for all handover decisions.

***Note:** In a 60-second simulation, only the first 0.5 sim-seconds (0.8% of total time) use the A3-only fallback. This has negligible impact on overall results. If a handover happens in the first 0.5s, it was decided by pure A3 — this explains why some early-simulation handovers might have slightly higher PP risk.*

19.4 Scenario D: GRU and A3 Disagree

A3 says Cell 5, but GRU says Cell 2. What does the xApp do?

1. **A3 evaluation:** UE #7 serving Cell 0 (SINR = +1.0 dB). Cell 5 SINR = +3.5 dB (passes A3: $3.5 > 1.0 + 2.0 = 3.0$). Cell 2 SINR = +2.8 dB (fails A3 marginally: $2.8 < 3.0$).
2. **GRU evaluation:** Rolling window shows Cell 2's SINR has been RISING steadily for 8 of 10 samples. Cell 5's SINR has been FLAT for 10 samples (just happens to be high now). GRU output: Cell 2 = 76%, Cell 5 = 18%.
3. **xApp validation logic:** GRU recommends Cell 2, but Cell 2 does NOT pass A3 (2.8 dB < threshold 3.0 dB). The xApp checks if GRU recommendation also passes A3.
4. **Decision:** In our implementation, the xApp requires the GRU-recommended cell to ALSO pass A3. Since Cell 2 narrowly fails A3 (by 0.2 dB), the xApp falls back to the A3-best cell (Cell 5) for this decision.

5. Outcome: HO to Cell 5. UE #7 stays in Cell 5 for 6 sim-seconds. No PP.

Note: This scenario shows the value of using A3 as a safety gate even when GRU gives a different recommendation. A tiny A3 margin (0.2 dB) might be noise — the combined system is conservative, preferring not to hand over to a cell that barely fails the quality threshold.

Understanding the Simulation Output Files

After a simulation run completes, several output files are generated. This chapter explains what each file contains and how to interpret it. The committee may ask you to explain specific numbers from your results — knowing how to read these files is essential.

20.1 decision_log.csv — The Primary Results File

This file contains one row per handover event. It is the primary source for computing PP rate and accuracy. Column definitions:

Column	Type	Example Value	Meaning
sim_time	float	25.300	Simulation time when HO decision was made (seconds)
ue_id	int	5	RNTI of the UE that performed the handover
source_cell	int	2	Cell ID the UE was leaving
target_cell	int	4	Cell ID the UE is moving to
gru_confidence	float	0.912	GRU model confidence for target_cell
a3_triggered	bool	True	Was A3 condition met? Always True in our system
serving_sinr	float	−1.0	Serving cell SINR at decision time (dB)
target_sinr	float	+3.5	Target cell SINR at decision time (dB)
is_pingpong	bool	False	Was this HO a return to previous cell within 5s?
prev_ho_time	float	17.000	Time of last HO for this UE (for PP calculation)

20.2 summary.json — The Results Summary

```
{
  "total_handovers": 309,
  "pingpong_handovers": 10,
  "pingpong_rate": 0.0324,
  "accuracy": 0.9676,
  "sim_time": 120.0,
  "num_ues": 20,
  "num_cells": 7,
  "a3_offset": 1.0,
  "a3_hysteresis": 1.0,
  "cooldown_time": 5.0,
```

```
"run_timestamp": "2026-05-05T21:02:59",
"random_seed": 11
}
```

20.3 How to Compute PP Statistics from decision_log.csv

To verify the PP numbers manually from the CSV:

1. Read all rows sorted by (ue_id, sim_time)
2. For each UE, find consecutive HO pairs: (HO_n: source=A→target=B) followed by (HO_{n+1}: source=B→target=A)
3. If $\text{HO}_{\{n+1\}}.\text{sim_time} - \text{HO}_n.\text{sim_time} < 5.0$: mark HO_{n+1} as ping-pong
4. Count total rows = total_handovers. Count PP-marked rows = pingpong_handovers
5. PP rate = $\text{pingpong_handovers} / \text{total_handovers} \times 100\%$

Python code to reproduce this:

```
import pandas as pd
df = pd.read_csv('decision_log.csv')
pp_count = df['is_pingpong'].sum()
total = len(df)
print(f'PP rate: {pp_count/total*100:.2f}%, Accuracy: {(total-pp_count)/total*100:.2f}%')
```

20.4 /tmp/flexric.log — What to Look For

When diagnosing simulation issues, /tmp/flexric.log is your first stop. Key log entries to look for:

Log Message	What It Means
E2 Setup Request received from gNB_id=1	NS-3 successfully connected to FlexRIC
KPM subscription confirmed for xApp	xApp's subscription was accepted
RIC Indication received: ue_id=5, sinr=-1.2	KPM report arriving normally
RIC Control Request sent: ue_id=5, target_cell=4	HO command was sent to NS-3
Connection reset by peer	NS-3 or xApp crashed — check other logs
Bind error: port 36421 in use	Previous FlexRIC instance still running

Design Decisions — Why Each Choice Was Made

Every number, every file, every architecture choice in this system was made for a reason. This chapter documents the reasoning behind key design decisions. Being able to defend WHY you made each choice is as important as knowing WHAT you built.

21.1 Why These Specific Parameter Values?

Parameter	Chosen Value	Why This Value (Not Higher/Lower)
A3_OFFSET	1.0 dB	Lower (0.5 dB) would trigger on measurement noise. Higher (2.0 dB) would delay necessary HOs. 1.0 dB represents a meaningful signal improvement in mmWave where SINR varies by ± 2 -3 dB.
A3_HYSTERESIS	1.0 dB	Same reasoning as offset. Total 2.0 dB combined threshold is consistent with 3GPP recommended values for intra-frequency 5G NR handovers in urban environments.
COOLDOWN_TIME	5.0 s	At 3 m/s pedestrian speed, 5 seconds = 15 meters of movement. Enough to move clearly to or from a cell coverage zone. Shorter (2s): still allows some PPs. Longer (8s): blocks necessary return HOs for slow-walking UEs near cell edges.
KPM_INTERVAL	0.05 s	20 samples/second provides sufficient temporal resolution for pedestrian mobility. Gives 10 samples over 0.5s for the rolling window. Faster would increase computational load; slower would miss channel dynamics.
WINDOW_SIZE	10	Covers 0.5 sim-seconds of history at 0.05s/sample. Enough to distinguish a sustained trend from a momentary spike. Smaller (5): may catch noise spikes. Larger (20): delays responsiveness to direction changes.
NUM_CELLS	7	Represents a realistic small urban cell cluster. 1 cell too simple (no inter-cell decisions). 19+ cells would increase simulation time prohibitively. 7 provides meaningful topology complexity for handover optimization research.
NUM_UES	20	Enough to generate statistically meaningful handover events (~ 2 -5 HOs/sim-second). Fewer UEs (5-10) would produce too few events for statistical analysis. More UEs (50+) would increase simulation time significantly.

21.2 Why Flask for the GRU Service?

Flask was chosen over alternatives (Django, FastAPI, gRPC) for three reasons:

- **Simplicity:** A Flask server with one route takes ~20 lines of code. Django and FastAPI require more boilerplate for a single-endpoint service.
- **Familiarity in ML community:** Flask is the most commonly used web framework in Python ML tutorials and examples — Fares was already familiar with it.
- **Adequate performance:** Flask's single-threaded default mode handles our ~20 requests/second load comfortably. We would only need FastAPI/gRPC for thousands of requests/second.
- **Easy debugging:** Flask auto-restarts on code changes (debug=True), and errors show detailed stack traces. Essential during development.

21.3 Why SCTP on Port 36421?

SCTP (Stream Control Transmission Protocol) was not chosen by us — it is mandated by the O-RAN Alliance specification for the E2 interface. Port 36421 is the IANA-registered port for S1AP/E2AP protocol (Signaling Transport over SCTP, registered as 'e2ap' service). This means any O-RAN-compliant E2 implementation (whether FlexRIC, OpenAirInterface, or a commercial vendor) must use SCTP on 36421. We cannot change this without violating the standard.

21.4 Why Separate GUI Components (2D and 3D)?

The 2D GUI (Grafana+InfluxDB) and 3D GUI (Vite+FastAPI) serve different purposes and were developed by different team members:

Aspect	2D GUI (Grafana)	3D GUI (Vite + FastAPI)
Primary purpose	Real-time operational monitoring	Research visualization and result export
Data source	InfluxDB (pushed continuously)	NS-3 output files (read at end of sim)
View type	Time-series plots, charts, tables	3D spatial: cell positions, UE movements, HO arcs
Update rate	Live (every 5 seconds)	Polling-based during simulation
Target user	Operator monitoring the simulation	Researcher analyzing spatial patterns
Technology	Docker: InfluxDB + Grafana + nginx	Python FastAPI + React/Three.js/Vite

21.5 Why the Specific Startup Order in gru.sh?

Every 1-second and 3-second wait in gru.sh exists for a specific reason:

Step	Command	Wait After	Why This Wait is Necessary
1	Start gru_xapp.py	2 seconds	Flask needs 1-2 seconds to load TensorFlow model into memory and start HTTP listener. If xApp starts before Flask is ready, first HTTP call fails.
2	Start FlexRIC nearRT-RIC	1 second	FlexRIC needs ~0.5s to bind to SCTP port 36421. If NS-3 connects before binding, it gets 'Connection refused'.
3	Start xApp (xapp_handover_gru)	3 seconds	xApp must complete the E2AP subscription handshake with FlexRIC (2-3 messages). NS-3 must NOT connect until subscription is established — otherwise KPM reports arrive with no subscriber to receive them.
4	Start NS-3 simulation	None	NS-3 connects to FlexRIC (now ready), receives subscription forwarding, starts generating KPM reports. Full pipeline is active.
5	Start sim_data_pusher.py	None	Can start any time after NS-3 since it reads output files. Running in background.

21.6 Trade-offs We Consciously Accepted

Every engineering system involves trade-offs. Here are the ones we made explicitly:

Trade-off	We Chose	What We Gave Up	Justification
Speed vs Accuracy	Full 3GPP TR 38.901 channel model (accurate, slow)	Faster simplified models	Research credibility requires realistic channel — this is a thesis, not a production system
Complexity vs ML power	GRU (moderate complexity)	Transformer (more powerful)	GRU is well-understood, proven for time-series, and sufficient for our window size
Safety vs Performance	Strict 5s cooldown (safe)	Faster response to genuine return HOs	Reducing PP is the primary goal; slight delay in genuine return HOs is acceptable
Integration vs Isolation	HTTP between xApp and GRU (isolated)	In-process call (faster)	HTTP allows independent development, debugging, and replacement of the GRU service
Runs vs Depth	3 deep runs (60-120s each)	30 short runs for statistics	Time constraint: 3×3 hours = 9 hours of compute vs $30 \times 3 = 90$ hours

Exam-Day Preparation — Self-Test and Final Review

22.1 15 Flash Facts — Memorize These Before You Sleep

If you remember nothing else, remember these 15 facts. The committee will almost certainly ask about some of them:

1	O-RAN Alliance founded 2018 by AT&T, Deutsche Telekom, NTT DoCoMo, Orange, China Mobile
2	Near-RT RIC timing window: 10 milliseconds to 1 second
3	E2 interface: SCTP port 36421, connects FlexRIC to NS-3
4	KPM report interval: 0.05 sim-seconds = 20 reports per sim-second
5	A3 threshold: $\text{neighbor_SINR} > \text{serving_SINR} + 2.0 \text{ dB}$ (1.0 offset + 1.0 hysteresis)
6	Cooldown: 5.0 sim-seconds per UE after any handover
7	GRU window: 10 samples covering 0.5 sim-seconds of history
8	GRU service: Flask on port 5000, gru_xapp.py, loads model once at startup
9	FlexRIC log: ALWAYS /tmp/flexric.log — NEVER stdout
10	Simulation: 7 cells, 20 UEs, 28 GHz mmWave, Random Walk 2D mobility
11	Pace: 60 sim-seconds $\approx$ 3.3 real hours (0.005 sim-s per real-s)
12	sim006: 334 HOs, 4 PPs, 1.72% PP rate, 98.28% accuracy
13	sim010: 137 HOs, 5 PPs, 3.65% PP rate, 96.35% accuracy
14	sim011: 309 HOs, 10 PPs, 3.24% PP rate, 96.76% accuracy (120 seconds)
15	Literature baseline: A3-only LTE/5G systems achieve 5-15% PP rate

22.2 Self-Test — Answer These Without Looking

Cover this column and try to answer each question. Then uncover and check:

Question	Answer
What port does FlexRIC use for E2?	36421 (SCTP)
What is the A3 combined threshold?	2.0 dB (1.0 + 1.0)

Question	Answer
How many cells in the simulation?	7 gNBs
How many UEs?	20
What is the KPM report interval?	0.05 sim-seconds
What is the GRU window size?	10 samples
What port does gru_xapp.py use?	5000 (Flask HTTP)
Where does FlexRIC log to?	/tmp/flexric.log
What is the cooldown time?	5.0 sim-seconds
What carrier frequency is used?	28 GHz (mmWave)
What is sim011's accuracy?	96.76%
What is sim006's PP rate?	1.72%
What does GRU stand for?	Gated Recurrent Unit
What are the two GRU gates?	Update gate and Reset gate
What does SINR measure?	Signal quality vs interference + noise (dB)
What does RSRP measure?	Absolute signal power (dBm)
What is E2SM-KPM?	E2 Service Model for Key Performance Measurements (monitoring)
What is E2SM-RC?	E2 Service Model for RAN Control (sending commands)
What is the simulation pace?	~0.005 sim-seconds per real-second
How long does sim006 take in real time?	~3.3 hours (60 sim-s / 0.005)
What is the Grafana port?	8000 (via nginx proxy)
What is the controller.py port?	8001 (FastAPI)
What is the Vite frontend port?	3001
What does 'ping-pong' mean?	HO from A→B then B→A within 5 sim-seconds
What file contains handover decisions?	decision_log.csv

22.3 How to Handle Questions You Don't Know the Answer To

It is normal not to know everything. Here is a professional framework for handling knowledge gaps in a defense:

1. **State what you DO know:** 'I know that the E2 interface uses SCTP, and I know the specific port is 36421, but I am not certain of the exact ASN.1 schema for the KPM indication message.'
2. **Acknowledge the limit:** 'This is an implementation detail I would need to look up in the O-RAN Alliance E2SM-KPM specification.'
3. **Show you know WHERE to find it:** 'The exact format is defined in O-RAN.WG3.E2SM-KPM-R003-v03.00.'
4. **Pivot to what you implemented:** 'In our FlexRIC-based implementation, the relevant fields are extracted by the FlexRIC SDK callback and we receive them as a C struct in the xApp callback function.'

This approach shows intellectual honesty, self-awareness, and depth of knowledge even for things you don't know by heart.

22.4 The 3-Minute Elevator Pitch

If someone asks 'what is your thesis about?' in 3 minutes or less:

Key Point: Mobile networks need to hand phones from one tower to the next as users walk around. Traditional systems do this with a simple rule: switch if the new tower is at least 2 dB stronger. The problem is this causes 'ping-pong' — the phone bounces back and forth between towers. We solved this using a GRU neural network that looks at the last 10 signal measurements to detect movement trends and predict the BEST tower to switch to. The system runs inside an Open RAN architecture — specifically a near-RT RIC called FlexRIC — which is the industry-standard platform for exactly this kind of intelligent network control. We tested it with NS-3, simulating a 7-tower 5G mmWave network with 20 moving users. Our system achieved ping-pong rates of 1.72 to 3.65%, compared to 5-15% for traditional systems. That means 96-98% of handover decisions were correct and stable.

Related Work — How This Thesis Fits in the Literature

A committee will often ask 'how does your work compare to the literature?' This chapter positions our thesis within the broader research landscape and identifies what makes our approach distinct.

23.1 Research Timeline — Handover Optimization Evolution

Era	Approach	Typical PP Rate	Key Limitation
Pre-2010 (2G/3G)	Fixed thresholds, manual tuning	15-25%	No adaptation to cell density or speed
2010-2015 (4G LTE)	A3 event with configurable TTT/offset	5-15%	Still purely reactive, no prediction
2015-2018 (4G ML era)	Decision trees, SVM, shallow ML	4-10%	No temporal modeling — ignores sequence
2018-2020 (5G mmWave)	LSTM, RNN for HO prediction	3-8%	High model complexity, slow inference
2020-2022 (O-RAN era)	xApp-based ML, GRU, attention models	2-6%	Often not integrated with real O-RAN stack
Our work (2025-2026)	GRU xApp + O-RAN E2 + cooldown	1.72-3.65%	Only 3 simulation runs, no real network test

23.2 Key Related Papers and How We Differ

The following papers are representative of the related work space:

- **Balasubramanian et al. (2020) — LSTM-based HO in LTE:** Used LSTM to predict HO targets in simulated LTE networks. Achieved ~4% PP rate. Key difference from our work: (1) LTE sub-6 GHz, not mmWave — channel behaves differently. (2) No O-RAN integration — they simulated the decision in MATLAB, not a real RIC. (3) No cooldown mechanism.
- **Polese et al. (2021) — NS-3 mmWave + FlexRIC integration:** Demonstrated that NS-3 mmWave can interface with FlexRIC via E2. Key difference: they focused on the integration infrastructure, not on ML-based handover optimization. We built on their integration work and added the GRU decision layer.
- **Dryjanski et al. (2021) — O-RAN xApp survey:** Survey of xApp architectures and use cases. Documents the xApp lifecycle we implement in best2.c. Our work provides a concrete implementation example for the handover use case they describe.
- **Kim et al. (2022) — GRU for 5G mmWave HO:** Used GRU in a simulated 5G mmWave environment. Achieved 3-5% PP rate. Key difference: their GRU is standalone (not integrated with O-RAN RIC), and

they do not use a cooldown timer. Our combined GRU + cooldown achieves lower PP rates (1.72%) in favorable mobility scenarios.

23.3 What Makes Our Work Novel

Our work's novelty is primarily in **systems integration** rather than algorithmic novelty:

- 1. End-to-end O-RAN implementation:** We are one of few thesis projects to implement GRU-based handover as a fully functional O-RAN xApp that communicates via the real E2 interface. Most related work evaluates ML algorithms in isolation, not integrated with a standards-compliant RIC.
- 2. Combined GRU + cooldown:** The explicit anti-ping-pong cooldown, combined with GRU prediction, has not been evaluated together in the literature we reviewed. Previous works use one or the other — not both simultaneously.
- 3. Multi-component open platform:** The full stack (NS-3 + FlexRIC + xApp + GRU service + 2D/3D GUIs + orchestration scripts) is open-source and reproducible — serving as a platform for future research in a way that closed simulations cannot.

23.4 Limitations Compared to State of the Art

Honest assessment of where our work falls short:

- **Statistical rigor:** 3 simulation runs vs 20-30 in published papers. Our results are preliminary trends, not statistically validated with confidence intervals.
- **No baseline comparison:** We did not run the same scenarios with pure A3 (no GRU). Without a baseline, we cannot claim 'X% improvement from GRU' — only that our absolute PP rate is in the good range.
- **Synthetic mobility:** Random Walk 2D is less realistic than SUMO urban mobility models or real user traces. Real deployments would see different PP patterns.
- **Single environment:** We tested one cell layout (7 cells, specific positions). Generalizability to different topologies is unknown.
- **No real network validation:** All results are from NS-3 simulation. Real mmWave channels have properties (weather effects, human body blockage, beam switching latency) that are simplified in simulation.

These limitations are standard for a university thesis and do not invalidate the contribution. They are honest starting points for future research.

Appendix — Abbreviations, References, and Full Parameter Index

A.1 List of Abbreviations

3GPP — 3rd Generation Partnership Project	5G NR — 5th Generation New Radio
A3 — Handover measurement event A3 (3GPP TS 38.331)	ASN.1 — Abstract Syntax Notation One
BBU — Baseband Unit	CLI — Command Line Interface
CSV — Comma-Separated Values	DL — Downlink (gNB → UE)
DU — Distributed Unit	E2AP — E2 Application Protocol
E2SM — E2 Service Model	EURECOM — European Research Center for Telecommunications
GRU — Gated Recurrent Unit	HTTP — HyperText Transfer Protocol
HO — Handover	IANA — Internet Assigned Numbers Authority
JSON — JavaScript Object Notation	KPM — Key Performance Measurement
LTE — Long-Term Evolution (4G)	LSTM — Long Short-Term Memory
ML — Machine Learning	mmWave — Millimeter Wave (frequencies > 24 GHz)
NLOS — Non-Line-of-Sight	LOS — Line-of-Sight
NS-3 — Network Simulator 3	O-CU — Open Central Unit
O-DU — Open Distributed Unit	O-RAN — Open Radio Access Network
O-RU — Open Radio Unit	PP — Ping-Pong (handover)
RAN — Radio Access Network	RC — RAN Control (E2 service model)
REST — Representational State Transfer	RF — Radio Frequency
RIC — RAN Intelligent Controller	RNN — Recurrent Neural Network
RNTI — Radio Network Temporary Identifier	RRC — Radio Resource Control
RSRP — Reference Signal Received Power	SCTP — Stream Control Transmission Protocol
SINR — Signal-to-Interference-plus-Noise Ratio	SMO — Service Management and Orchestration
TCP — Transmission Control Protocol	TTT — Time-to-Trigger
UE — User Equipment	UL — Uplink (UE → gNB)
WSGI — Web Server Gateway Interface	xApp — near-RT RIC application

A.2 Key File Paths — Complete Reference

File / Directory	Full Path
Project root	/home/omar_farouk/open-ran-clean/
NS-3 simulation script	.../yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/gru_scenario.cc
xApp source code	.../yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/best2.c
GRU Python service	.../yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/gru_xapp.py
InfluxDB pusher	.../yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/sim_data_pusher.py
Master launcher	.../yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/gru.sh
Kill script	.../yousef_fathy/ns-O-RAN-flexric/mmwave-LENA-oran/kill_sim.sh
2D GUI Docker compose	/home/omar_farouk/open-ran-clean/GUI/docker-compose.yml
3D GUI controller	/home/omar_farouk/open-ran-clean/controller.py (or 3D_GUI/backend/)
Simulation results	/home/omar_farouk/open-ran-clean/3D_GUI_Sim_Results/
FlexRIC runtime log	/tmp/flexric.log
This thesis guide	/home/omar_farouk/open-ran-clean/GRU_ORAN_Thesis_Guide.pdf
PDF generator script	/home/omar_farouk/open-ran-clean/generate_thesis_pdf.py

A.3 O-RAN Alliance Standards Referenced

Standard	Title	Relevance
O-RAN.WG3.E2AP-R003	E2 Application Protocol	Protocol used by FlexRIC and NS-3 to communicate
O-RAN.WG3.E2SM-KPM-R003	E2 Service Model: KPM	Format of SINR/RSRP measurement reports
O-RAN.WG3.E2SM-RC-R003	E2 Service Model: RC	Format of handover control commands
3GPP TR 38.901	5G NR Channel Model	Channel propagation model used in NS-3 mmWave
3GPP TS 38.331	5G NR RRC Protocol	RRC Reconfiguration procedure for handovers
3GPP TS 38.321	5G NR MAC Protocol	RNTI assignment to UEs

The GRU Model — Step-by-Step Mathematics (Simplified)

N1.1 Why We Need to See the Equations

Understanding the GRU equations is not about memorizing formulas — it is about being able to explain to your committee what the model is actually doing at each time step when it processes a UE's SINR history. Each equation corresponds to a real decision: 'how much should I remember?', 'is the past relevant right now?', 'what new thing should I consider?', and 'what is my final memory state?' These four decisions, applied 10 times in sequence (once per KPM sample), produce the hidden state that the dense output layer uses to predict the best handover target cell.

N1.2 The Four Core GRU Equations

At each timestep t (each new KPM report for a UE), the GRU cell computes four quantities. The notation below uses: $\mathbf{x}_t$ = new input vector (SINR, RSRP, velocity at time t), $\mathbf{h}_{\{t-1\}}$ = hidden state from previous timestep (the 'memory'), $\mathbf{W}$ and $\mathbf{U}$ = learned weight matrices, $\mathbf{b}$ = learned bias vectors, **sigmoid** = $1/(1+e^{-x})$ which squashes output to range $[0, 1]$, **tanh** = hyperbolic tangent which squashes output to range $[-1, +1]$, and $\odot$ = element-wise (Hadamard) multiplication.

Equation 1: Update Gate

$$\mathbf{z}_t = \text{sigmoid}(\mathbf{W}_z \cdot \mathbf{x}_t + \mathbf{U}_z \cdot \mathbf{h}_{\{t-1\}} + \mathbf{b}_z)$$

The **update gate** $\mathbf{z}_t$ is a vector of values between 0 and 1. Each element asks the question: 'how much should this dimension of the hidden state be updated with new information versus kept from before?' When $\mathbf{z}_t$ is close to 1, the hidden state is almost entirely replaced with the new candidate. When $\mathbf{z}_t$ is close to 0, the old hidden state is preserved almost unchanged. The sigmoid activation ensures $\mathbf{z}_t$ is always in $[0, 1]$, making it interpretable as a blending ratio.

***Note:** Analogy: The update gate is like an inbox filter that decides how much of today's new mail to actually read and act on versus keeping yesterday's notes unchanged. If SINR is rapidly dropping (clear trend), $\mathbf{z}_t$ stays high — update aggressively. If SINR is stable, $\mathbf{z}_t$ falls low — preserve existing memory.*

Equation 2: Reset Gate

$$\mathbf{r}_t = \text{sigmoid}(\mathbf{W}_r \cdot \mathbf{x}_t + \mathbf{U}_r \cdot \mathbf{h}_{\{t-1\}} + \mathbf{b}_r)$$

The **reset gate** $\mathbf{r}_t$ is also a vector in $[0, 1]$. Each element asks: 'when computing the new candidate hidden state, how relevant is the previous hidden state $\mathbf{h}_{\{t-1\}}$?' When $\mathbf{r}_t$ is close to 0, the previous hidden state is completely ignored for this timestep — the cell acts as if starting fresh. When $\mathbf{r}_t$ is close to 1, the previous state is fully taken into account. This gate allows the GRU to forget irrelevant past context when the current input changes dramatically — for example, when a UE suddenly enters a blockage and SINR drops to a new regime.

Note: Analogy: The reset gate is like a researcher deciding 'are my old notes still relevant to this new question?' If the UE has completely moved to a new area, old SINR history from its previous location is irrelevant — r_t goes to 0 and the GRU starts fresh.

Equation 3: Candidate Hidden State

$$h_tilde_t = \tanh(W_h \cdot x_t + U_h \cdot (r_t \blacksquare h_{t-1}) + b_h)$$

The **candidate hidden state** h_tilde_t is what the new information 'wants' the hidden state to be. Notice that h_{t-1} is multiplied by r_t element-wise before being used — this is how the reset gate filters how much of the old memory participates in forming the candidate. The tanh activation squashes the result to $[-1, +1]$, providing a normalized representation of the new candidate information.

Note: Analogy: The candidate is the committee's initial proposal for a new decision, already filtered by the reset gate to include only the relevant parts of past experience. It is not the final answer — it still needs to be blended with old memory.

Equation 4: Final Hidden State (Output)

$$h_t = (1 - z_t) \blacksquare h_{t-1} + z_t \blacksquare h_tilde_t$$

The **final hidden state** h_t is a linear interpolation between the old hidden state h_{t-1} and the new candidate h_tilde_t , controlled by the update gate z_t . When $z_t = 0$: $h_t = h_{t-1}$ (no update, preserve perfectly). When $z_t = 1$: $h_t = h_tilde_t$ (full update, replace completely). In practice, z_t takes values between 0 and 1, so h_t is always a blend. This is the key mechanism that prevents vanishing gradients: the ' $(1 - z_t) \blacksquare h_{t-1}$ ' term provides a direct highway from old state to new state, bypassing the squashing activation that would otherwise shrink the gradient.

N1.3 Plain-English Summary of All Four Equations

Gate / State	Equation Summary	Analogy
Update gate z_t	How much do I update? (0=keep old, 1=replace with new)	How much of today's news do I need to act on?
Reset gate r_t	Is the past relevant NOW? (0=ignore past, 1=use past fully)	Are my old notes still relevant to this question?
Candidate h_tilde_t	What new info should I absorb, given the filtered past?	Draft proposal from filtered experience + new data
Final state h_t	Blend old memory with new candidate using z_t as weight	Final decision: part old memory, part new insight

N1.4 Numeric Example: SINR Drops from 15 dB to 8 dB

Consider UE #7 whose serving cell SINR was stable at about 15 dB for the first 6 samples, then drops to 8 dB over the next 4 samples (perhaps the UE walked behind a building). Here is what the GRU 'feels' as it processes this 10-sample window:

Sample	Serving SINR	GRU Behavior
1	15.2 dB	h_0 initialized to zeros — GRU has no memory yet
2	15.0 dB	Stable input. Update gate z low — $h_1 \approx h_0$. No significant change.
3	14.8 dB	Slight drop. z still low — small update. GRU remembers ~15dB regime.
4	14.5 dB	Continuing drop. Reset gate starts opening — past less 'perfect'.
5	13.0 dB	Clear trend emerging. Update gate rises — 'something is changing'.
6	11.2 dB	Strong drop. z rises further. Candidate h_{tilde} strongly represents falling trend.
7	9.8 dB	Rapid decline. GRU hidden state now represents 'degrading fast'.
8	9.1 dB	h_t now strongly encodes the 'SINR declining' pattern in its dimensions.
9	8.5 dB	GRU memory fully adapted to new regime — 8-9dB is the learned context.
10	8.0 dB	Final hidden state fed to Dense+softmax. Output: neighbor cell with rising SINR gets high probability.

The key insight: a pure threshold rule looking only at sample 10 might say 'SINR=8dB, serving cell is still OK, no handover.' But the GRU, having processed all 10 samples, encodes the TREND of decline and predicts: this UE is moving away, switch now to avoid connection drop.

N1.5 Input Feature Table — What the GRU Receives at Each Timestep

Feature Name	Unit	Typical Range	Why It Matters for Handover
serving_sinr	dB	-10 to +30	Primary signal quality — falling trend indicates need for HO
serving_rsrp	dBm	-140 to -44	Signal power level — absolute strength from current cell
velocity	m/s	0 to 5 (pedestrian)	Fast UE = needs predictive HO; slow UE = can wait for clear signal
serving_cell_id	integer (0-6)	0, 1, 2, 3, 4, 5, 6	Current cell context — model learns cell-specific behavior
neighbor_sinr_0	dB	-20 to +25	SINR from Cell 0 (even if not serving) — rising = approaching Cell 0
neighbor_sinr_1	dB	-20 to +25	SINR from Cell 1 — rising trend = Cell 1 becoming best candidate
neighbor_sinr_2	dB	-20 to +25	SINR from Cell 2 — best candidate if this rises while serving falls
neighbor_sinr_3	dB	-20 to +25	SINR from Cell 3 — same as above for Cell 3

Feature Name	Unit	Typical Range	Why It Matters for Handover
neighbor_sinr_4	dB	-20 to +25	SINR from Cell 4
neighbor_sinr_5	dB	-20 to +25	SINR from Cell 5
neighbor_sinr_6	dB	-20 to +25	SINR from Cell 6 — full picture of all cells in each timestep

With 11 features per timestep and 10 timesteps, the GRU processes an input tensor of shape (1, 10, 11) for each prediction. The 10 timesteps allow it to see TRENDS — not just the current value, but whether each cell's signal is rising or falling over the past 0.5 simulation-seconds.

The E2 Interface — How NS-3 Talks to FlexRIC

N2.1 The E2 Interface in Context

The E2 interface is the central nervous system of our simulation. Without it, NS-3 would be a closed world — running a virtual network with no external intelligence. The E2 interface opens a bi-directional channel between NS-3 (the E2 Node) and FlexRIC (the near-RT RIC), allowing: (1) measurement reports to flow from NS-3 to the xApp, and (2) control commands to flow from the xApp back to NS-3. This two-way communication is what makes our system a genuine O-RAN implementation rather than just a simulation with hardcoded decisions.

The E2 interface is defined by the O-RAN Alliance Working Group 3 (WG3) and consists of two layers: the **E2AP (E2 Application Protocol)** which is the signaling envelope, and the **E2 Service Models (E2SMs)** which define the content of the messages. Think of E2AP as the envelope and E2SMs as the letters inside.

N2.2 E2AP — The Envelope Protocol

E2AP (E2 Application Protocol) is the outer protocol layer that handles all E2 interface management. It is responsible for:

- **E2 Setup:** When NS-3 connects, it sends an E2 Setup Request identifying itself as an E2 Node. FlexRIC responds with an E2 Setup Response that confirms the connection and negotiates which service models are supported.
- **E2 Subscription:** The xApp registers interest in receiving specific measurements (KPM service) at a specific interval. This subscription flows through FlexRIC to NS-3.
- **E2 Indication:** NS-3 sends measurement reports wrapped in E2AP Indication messages. The payload (the actual SINR/RSRP data) is formatted according to E2SM-KPM.
- **E2 Control:** The xApp sends handover commands wrapped in E2AP Control messages. The payload is formatted according to E2SM-RC.
- **E2 Reset and Teardown:** Used to cleanly close the connection when the simulation ends.

All E2AP messages are encoded using **ASN.1 (Abstract Syntax Notation One)** — a binary encoding format designed for telecom protocols. ASN.1 is compact and standardized, but computationally expensive to encode and decode, which contributes to the simulation's slow pace.

N2.3 E2SM-KPM — The Monitoring Service Model

E2SM-KPM (E2 Service Model for Key Performance Measurements) defines the structure of monitoring reports. It is the 'monitoring language' of the E2 interface. When the xApp subscribes to KPM, it specifies: which measurements it wants (SINR, RSRP), which granularity period (0.05 sim-seconds), and which cells and UEs to monitor. NS-3 then encodes these measurements into KPM format and sends them periodically.

The KPM Indication message contains a **PM Container** (Performance Measurement Container) with a list of UE measurement records. Each record identifies the UE by RNTI (Radio Network Temporary Identifier) and provides the measurement values for the serving cell and all neighbor cells.

N2.4 E2SM-RC — The Control Service Model

E2SM-RC (E2 Service Model for RAN Control) defines the format of control commands. It is the 'command language' of the E2 interface. When the xApp wants to trigger a handover, it constructs an RC Control message specifying:

- The UE to be handed over (identified by RNTI)
- The target cell ID (which of the 7 cells the UE should move to)
- The control style (handover command type as defined in 3GPP TS 38.331)

NS-3 receives this RC Control message, decodes the target cell ID, and executes the RRC Reconfiguration procedure to move the UE. From NS-3's perspective, the handover command arrives 'from the RIC' — exactly as it would in a real O-RAN deployment.

N2.5 Message Sequence — Step by Step

Here is the complete sequence of E2 messages from simulation startup to the first handover, numbered for reference:

1. **NS-3 starts:** `gru_scenario.cc` initializes and creates the `E2TermInterface` object. This object implements the E2 Node side of the E2AP protocol.
2. **SCTP connection:** `E2TermInterface` opens a SCTP socket and connects to 127.0.0.1:36421 (FlexRIC's E2 listening port). SCTP provides ordered, reliable delivery with support for multiple streams — ideal for signaling traffic.
3. **E2 Setup Request →:** NS-3 sends E2 Setup Request to FlexRIC. This message contains the E2 Node ID (gNB identifier), the global cell IDs (0-6), and the list of E2 Service Models supported (KPM and RC).
4. **← E2 Setup Response:** FlexRIC responds with E2 Setup Response, confirming the E2 Node is registered. FlexRIC now knows about all 7 cells.
5. **xApp subscribes to KPM:** The xApp (already running) sends a KPM Subscription Request through FlexRIC's internal API. FlexRIC translates this into an E2 Subscription Request message to NS-3.
6. **← E2 Subscription Response:** NS-3 acknowledges the subscription and configures its internal KPM report generator to fire every 0.05 sim-seconds.
7. **KPM Indications (every 0.05 sim-seconds) →:** NS-3 sends E2AP Indication messages, each containing one KPM report with all UE measurements. FlexRIC extracts the payload and delivers it to the xApp callback function.
8. **xApp processes report:** A3 check + GRU call. Decision: UE 5 should move from Cell 1 to Cell 3.
9. **xApp sends RC Control →:** The xApp sends an E2SM-RC Control message through FlexRIC to NS-3, specifying UE RNTI for UE 5 and target cell ID = 3.

10. **NS-3 executes handover:** NS-3 receives the RC Control message, initiates RRC Reconfiguration, UE 5 disconnects from Cell 1 and connects to Cell 3. The event is logged to handover.csv.
11. **← RC Control Acknowledgment:** NS-3 sends an E2AP Control Acknowledge back to FlexRIC confirming the handover was executed.

N2.6 What is SCTP and Why Not TCP?

SCTP (Stream Control Transmission Protocol) is a transport layer protocol designed specifically for telecom signaling traffic. The O-RAN Alliance mandates SCTP for the E2 interface because:

Property	TCP	SCTP (Used by E2)
Streams	Single ordered stream	Multiple independent streams — one message delay does not block others
Head-of-line blocking	Yes — one lost packet blocks all	No — streams are independent
Message boundaries	No — TCP is a byte stream	Yes — SCTP preserves message boundaries natively
Multi-homing	No — one IP per connection	Yes — multiple IPs per endpoint (failover)
Connection setup	3-way handshake	4-way handshake (cookie-echo, more secure against SYN flood)
Port (E2 interface)	N/A	36421 (IANA registered for S1AP/E2AP)

In practice, for our loopback simulation (127.0.0.1 → 127.0.0.1 on a single machine), the differences between TCP and SCTP are invisible. But the E2 specification requires SCTP, so FlexRIC implements it, and NS-3's E2 module speaks SCTP on port 36421.

N2.7 What Happens If the Connection Drops Mid-Simulation

If the SCTP connection between NS-3 and FlexRIC is interrupted (e.g., FlexRIC is killed while NS-3 is still running):

1. NS-3's E2TermInterface detects the SCTP connection failure via socket error
2. NS-3 logs the disconnection and stops sending KPM Indication messages
3. NS-3 continues the simulation internally (UEs keep moving, channel model keeps running) but without E2 interface — no more handover commands will arrive
4. FlexRIC, if restarted, would need NS-3 to also restart for a fresh E2 Setup
5. Result: the rest of the simulation runs with UEs staying on their current cells — possible degradation and no HOs recorded after disconnection point

This is why the kill sequence in kill_sim.sh is important: NS-3 is killed before FlexRIC to avoid this partial-run state.

KPM Reports — What Data Flows Every 50ms

N3.1 The KPM Report as a Data Packet

Every 0.05 simulation-seconds (50 milliseconds in simulated time), NS-3 assembles a KPM Indication message and sends it to FlexRIC. This message is a complete snapshot of the radio environment for all UEs at that instant. Think of it as a photograph of the entire network taken every 50ms — frozen in time, containing every metric the xApp needs to make a handover decision.

For 20 UEs and 7 cells, each KPM report contains 20 UE measurement records. At 0.05s intervals over a 120-second simulation, that is: $120 / 0.05 = 2,400$ reports $\times$ 20 UEs = 48,000 individual UE measurement snapshots. This is why the lstm_features.csv file (which stores these features) grows to 7.8 MB.

N3.2 Complete Field-by-Field Breakdown

Field Name	Type	Unit	Example Value	Why the xApp Needs It
UE ID (RNTI)	Integer	None	5	Identifies which of the 20 UEs this measurement belongs to
Serving Cell ID	Integer (0-6)	None	1	Current serving cell — needed to know what to hand over FROM
sim_time	Float	seconds	10.05	Timestamp — used to check cooldown timer per UE
serving_sinr	Float	dB	15.8	Signal quality of current cell — primary A3 comparison value
serving_rsrp	Float	dBm	-82.3	Signal power of current cell — secondary quality indicator
neighbor_sinr_0	Float	dB	8.2	SINR from Cell 0 (if not serving) — candidate comparison
neighbor_sinr_1	Float	dB	18.1	SINR from Cell 1 — if $>$ serving + 2dB, A3 fires
neighbor_sinr_2	Float	dB	12.7	SINR from Cell 2
neighbor_sinr_3	Float	dB	7.4	SINR from Cell 3
neighbor_sinr_4	Float	dB	4.1	SINR from Cell 4
neighbor_sinr_5	Float	dB	3.8	SINR from Cell 5
neighbor_sinr_6	Float	dB	6.9	SINR from Cell 6

Field Name	Type	Unit	Example Value	Why the xApp Needs It
ue_x	Float	meters	142.3	UE x-coordinate (if position reporting enabled)
ue_y	Float	meters	87.6	UE y-coordinate
velocity	Float	m/s	2.1	Estimated UE speed — derived from position change

N3.3 SINR — The Primary Quality Metric

SINR (Signal-to-Interference-plus-Noise Ratio) measures how well the desired signal stands out from background interference and noise. It is measured in **decibels (dB)**, which is a logarithmic scale. A 10 dB increase means the ratio is 10× better in linear terms. A 3 dB increase means approximately 2× better.

SINR Range (dB)	Quality Description	What It Means for the UE
> 20 dB	Excellent	Near center of cell. Very high data rates possible. HO very unlikely.
10 to 20 dB	Good	Comfortable coverage zone. Moderate data rates. Some distance from cell edge.
0 to 10 dB	Marginal	Approaching cell edge. Data rates limited. HO likely soon.
< 0 dB	Poor	At cell edge or in shadow. Very low data rates. HO urgently needed.
< -10 dB	Critical	Deep shadow or far from all cells. Connection likely to drop if no HO.

In our simulation with 28 GHz mmWave, the SINR range is typically -10 dB to +25 dB. Values above 25 dB are rare because the high path loss at 28 GHz limits the peak received power. Values below -10 dB usually trigger an immediate A3 evaluation.

N3.4 RSRP — The Signal Strength Metric

RSRP (Reference Signal Received Power) measures the absolute power level of the cell's reference signal at the UE antenna. It is measured in **dBm** (decibels relative to 1 milliwatt). RSRP tells you 'how strong is the signal?' — but not 'how usable is it?' (that is what SINR does). A cell may have high RSRP but low SINR if there is heavy interference from other cells.

RSRP Range (dBm)	Quality Description	Context in 5G mmWave
> -80 dBm	Excellent	Very close to gNB. Strong beamforming alignment.
-80 to -90 dBm	Good	Comfortable coverage. Typical mid-cell position.
-90 to -100 dBm	Moderate	Approaching cell edge. Beam may be misaligned.

RSRP Range (dBm)	Quality Description	Context in 5G mmWave
-100 to -110 dBm	Poor	Near cell edge. Beam alignment difficult.
< -110 dBm	Very poor	Beyond cell edge or in blockage. HO essential.

The key difference between SINR and RSRP: RSRP measures one cell's signal in isolation. SINR accounts for ALL cells — a cell with high RSRP may still have low SINR if its neighbors are also transmitting strongly (high interference). For handover decisions, SINR is more important than RSRP.

N3.5 Neighbor Cell Measurements — The Full Radio Picture

A distinctive feature of our KPM reports is that they include **neighbor cell measurements** — SINR and RSRP values not just from the serving cell, but from all 6 other cells. This is critical for handover decisions because:

- Without neighbor measurements, the xApp only knows 'current cell is getting worse' but not 'which other cell is better'
- The A3 comparison directly uses neighbor_sinr vs serving_sinr — impossible without neighbor measurements
- The GRU model uses neighbor SINR trends to predict which cell the UE is moving toward
- A UE near the boundary of Cells 2 and 3 might show Cell 2 = 15dB, Cell 3 = 14.5dB — the GRU uses the rising/falling trend over 10 samples to decide

N3.6 Why 50ms is 'Near-Real-Time'

The O-RAN architecture defines three control loops by their latency:

Control Loop	Latency	Example Application	Our KPM Period
Non-Real-Time	> 1 second	Model training, policy generation, configuration	N/A — too slow for HO
Near-Real-Time	10ms – 1s	Handover optimization, interference management, load balancing	50ms — fits perfectly
Real-Time (embedded)	< 10ms	Radio scheduler, beam management, HARQ retransmission	Too fast for our system

The 50ms KPM period places our system firmly in the near-RT RIC domain. It is fast enough to detect channel trends before the UE moves too far (at 3 m/s pedestrian speed, the UE moves 0.15 meters in 50ms — imperceptible), but not so fast that we overwhelm the processing pipeline. The 50ms granularity was also chosen because the GRU window covers $10 \times 50\text{ms} = 500\text{ms}$ of history, which is enough to see a clear SINR trend while remaining within the near-RT latency budget.

The Ping-Pong Effect — Deep Dive

N4.1 Precise Definition of Ping-Pong

A **ping-pong handover event** is formally defined as: a UE performs a handover from cell A to cell B, and then within a short time window, performs another handover from cell B back to cell A. The 'bouncing' between two cells is reminiscent of a ping-pong ball, hence the name.

The 3GPP standard (3GPP TS 36.331) defines ping-pong as a return handover within **1-2 seconds** in 4G LTE networks. In our system, we use a more conservative definition: a handover is marked as ping-pong if the UE returns to its previous cell within **5.0 simulation-seconds**. This stricter definition catches more potential instability while accounting for the slower-paced mmWave environment.

Mathematically: HO at time T (cell A $\rightarrow$ B) is ping-pong if there exists another HO at time T_2 (cell B $\rightarrow$ A) where $T_2 - T \leq 5.0$ sim-seconds.

N4.2 Root Causes of Ping-Pong Handovers

Root Cause 1: Cell Edge Location

When a UE is exactly between two cells — equidistant from Cell A and Cell B — both cells produce similar SINR values (e.g., Cell A = 12.3 dB, Cell B = 12.5 dB). A small random fluctuation (as small as 0.5 dB due to small-scale fading) can push Cell B above the A3 threshold, triggering a handover. Then another fluctuation brings Cell A back above Cell B, triggering a reverse handover. The UE 'oscillates' because its physical position genuinely does not favor either cell strongly.

Root Cause 2: Shadow Fading Fluctuation

At 28 GHz, even a person walking between the UE and a cell can cause 10-20 dB of signal attenuation (blockage). Shadow fading at mmWave is rapid and severe. A momentary blockage of Cell A's signal by a pedestrian can make Cell B appear 12 dB better — triggering an A3 event and handover. When the blocking pedestrian moves away (within 1-2 seconds), Cell A recovers, and the UE hands back.

Root Cause 3: A3 Threshold Too Low

If $A3_OFFSET + A3_HYSTERESIS$ is set too small (e.g., 0.5 dB instead of 2.0 dB), even tiny measurement noise (~ 0.3 dB standard deviation in NS-3) can repeatedly trigger A3 events back and forth. This is why we use 2.0 dB total threshold: it is large enough to ignore noise but small enough to catch real coverage differences.

Root Cause 4: UE Moving Parallel to Cell Boundary

A UE walking along a street that runs parallel to the boundary between Cell 2 and Cell 3 will maintain nearly equal SINR from both cells for an extended period. In a random walk mobility model, such trajectories occur regularly. During this parallel-path movement, any small deviation closer to Cell 2 or Cell 3 triggers a handover, and the subsequent return deviation triggers another.

N4.3 Real-World Impact of Ping-Pong

- **Wasted radio resources:** Each handover requires signaling messages (RRC Reconfiguration) that consume radio bandwidth on both the old and new cell.
- **Temporary connection interruption:** During a hard handover, there is a brief period (typically 10-50ms) where the UE is detached from both cells — data packets may be delayed or dropped.
- **Increased signaling load:** The RRC Reconfiguration messages propagate through the control plane. In a real network with thousands of UEs, frequent ping-pong events generate measurable signaling overhead.
- **User experience degradation:** An active VoIP call or video stream may be briefly interrupted during each ping-pong cycle — the user hears a 'blip' or sees a frame drop.
- **Battery consumption:** Handovers require the UE's radio module to re-synchronize with the new cell, consuming extra power. A UE in a ping-pong zone drains its battery faster than a UE with a stable connection.

N4.4 The Cooldown Mechanism — Step-by-Step

Our primary defense against ping-pong is the **cooldown timer**, implemented in best2.c. Here is the exact step-by-step behavior:

1. **UE 5 completes handover at sim_time = 10.0s** (Cell 1 → Cell 3). The xApp records: `last_ho_time[5] = 10.0`
2. **At sim_time = 12.0s:** A new KPM report arrives for UE 5. The A3 check shows Cell 1 SINR = 16.0 dB, Cell 3 (now serving) SINR = 14.5 dB — Cell 1 is 1.5 dB better. BUT: the xApp first checks: $12.0 - 10.0 = 2.0 < 5.0 \rightarrow$ COOLDOWN ACTIVE $\rightarrow$ skip evaluation entirely.
3. **At sim_time = 14.9s:** Another KPM report. Same check: $14.9 - 10.0 = 4.9 < 5.0 \rightarrow$ still in cooldown $\rightarrow$ skip.
4. **At sim_time = 15.1s:** $15.1 - 10.0 = 5.1 > 5.0 \rightarrow$ cooldown expired. A3 evaluation now allowed. If Cell 1 is still 2.0+ dB better, handover fires. If not, UE stays with Cell 3.
5. **If handover fires at 15.1s** (Cell 3 → Cell 1): this IS a ping-pong (Cell 3 → Cell 1 within 5.1s of the Cell 1 → Cell 3 HO at 10.0s). The cooldown allowed it because $5.1 > 5.0$. This is a boundary artifact.

Cooldown condition: `(current_sim_time - last_ho_time[ue_id]) > COOLDOWN_TIME`
COOLDOWN_TIME = 5.0 sim-seconds (defined in best2.c)

N4.5 Why 5.0 Seconds? Physical Reasoning

The 5.0 second cooldown was not chosen arbitrarily. Consider a UE moving at the typical pedestrian speed in our simulation (~3 m/s). In 5 simulation-seconds, the UE moves approximately $3 \times 5 = 15$ meters. At 28 GHz, the typical cell radius in a dense urban deployment is 100-200 meters. A UE that moves 15 meters in 5 seconds has moved enough to establish whether it is truly heading away from the old cell or merely experiencing a momentary fluctuation.

A shorter cooldown (e.g., 2 seconds = 6 meters of movement) is insufficient — the UE might genuinely still be at the cell boundary after 2 seconds. A longer cooldown (e.g., 10 seconds = 30 meters) might block necessary

handovers for UEs that genuinely need to change cells quickly. The 5-second value was validated empirically: our PP rate of 3.24% with 5s cooldown vs higher rates observed in preliminary tests with 2s cooldown.

N4.6 Why Ping-Pongs Still Happen Despite Cooldown

The 10 ping-pong events observed in sim011 (despite the 5.0s cooldown) are not a failure of the mechanism — they are a mathematical boundary artifact. Consider:

- The cooldown blocks HOs for exactly 5.0 seconds. A return HO at exactly $T+5.01s$ passes the cooldown check ($5.01 > 5.0$) but is still within 5.01 seconds of the original — classified as ping-pong.
- With 20 UEs and 309 HOs over 120 sim-seconds, the probability of at least one UE being at a cell boundary just as the cooldown expires is significant.
- Random walk mobility produces reversals — a UE can genuinely need to return to the previous cell because it randomly changed direction.

The 10 PPs out of 309 HOs represent cases where the UE's trajectory reversed direction shortly after a handover, or where the cooldown expired at exactly the wrong moment (UE still near the cell boundary). Eliminating these entirely would require either an infinite cooldown (which would block all handovers) or perfect future trajectory prediction (physically impossible for a random walk model).

N4.7 Detailed Calculation: What 3.24% Means

$$\text{PP Rate} = (\text{Ping-Pong HOs} / \text{Total HOs}) \times 100\%$$

$$\text{sim011: PP Rate} = (10 / 309) \times 100\% = 3.236\% \approx 3.24\%$$

Unpacking what this means in practical terms: in a 120-second simulation with 20 UEs, 309 handover decisions were made. Of these, 299 were 'good' decisions — the UE moved to the new cell and stayed there for at least 5 simulation-seconds. Only 10 were 'bad' decisions where the UE bounced back. The system made the right call 96.76% of the time.

Compared to the literature baseline: traditional A3-only handover in 5G mmWave systems typically shows 8-20% PP rate. Our system's 3.24% represents a 60-84% reduction in ping-pong events compared to the baseline approach.

N4.8 How GRU Reduces Ping-Pong Versus Pure A3

The GRU's contribution to ping-pong reduction works through two mechanisms:

- **Trend validation:** When the A3 condition fires (neighbor is 2+ dB better), the GRU checks whether this improvement is part of a sustained trend or just a momentary spike. If the neighbor's SINR only recently crossed the threshold (one sample above, rest below), the GRU may recommend a different cell or the serving cell — preventing a premature HO.
- **Stable neighbor selection:** When multiple neighbors all pass the A3 threshold, the GRU selects not the momentarily strongest one, but the one whose SINR has been consistently rising over the 10-sample window. A cell that just had a sudden spike is less likely to be recommended than a cell showing a steady upward trend.

Step-by-Step Simulation Walkthrough — sim011

N5.1 Overview

This chapter walks through a complete simulation run — sim011 — from the moment the user types 'bash gru.sh' to the moment the final results are saved. Understanding this flow end-to-end demonstrates that you know not just the theory but the actual system operation. sim011 ran for 120 simulation-seconds and produced 309 handovers with 10 ping-pong events (3.24% PP rate, 96.76% accuracy).

N5.2 The Complete Step-by-Step Flow

Phase 1: Startup (Steps 1-8)

1. **User runs: bash gru.sh** — the master startup script begins executing in the terminal. The user will see minimal output since most services run in the background.
2. **gru.sh starts Docker services:** `docker -compose -f GUI/docker-compose.yml up -d` — InfluxDB, Grafana, and nginx containers launch. The script waits 3 seconds for containers to initialize.
3. **gru.sh starts FlexRIC:** `./nearRT-RIC &` — FlexRIC binds to SCTP port 36421 and waits for E2 connections. All output goes to `/tmp/flexric.log`. The script waits 2 seconds.
4. **gru.sh starts gru_xapp.py:** `python3 gru_xapp.py &` — Flask server starts on port 5000. The GRU model is loaded from disk (takes ~1 second). The script waits 2 seconds to ensure Flask is ready.
5. **gru.sh starts xapp_handover_gru:** The compiled C xApp binary starts, registers with FlexRIC, and sends a KPM Subscription Request: 'please send me SINR/RSRP for all UEs every 0.05 sim-seconds.' FlexRIC holds this subscription, waiting for NS-3.
6. **gru.sh waits 3 more seconds** — allowing the xApp to complete its subscription registration with FlexRIC before NS-3 arrives.
7. **gru.sh starts NS-3:** `./ns3 run gru_scenario --simTime=120 --RngRun=11` — NS-3 begins. It creates 7 gNBs, 20 UEs, sets up the 28 GHz mmWave channel model, and opens a SCTP connection to FlexRIC at 127.0.0.1:36421.
8. **gru.sh starts sim_data_pusher.py** in background — begins monitoring NS-3 output files and pushing data to InfluxDB every 5 real-seconds for Grafana visualization.

Phase 2: Simulation Begins (Steps 9-11)

9. **E2 Setup handshake:** NS-3 sends E2 Setup Request → FlexRIC responds with E2 Setup Response → FlexRIC forwards the xApp's KPM subscription to NS-3 → NS-3 acknowledges. Simulation clock starts at `sim_time = 0.000s`. All 20 UEs begin moving in random directions.
10. **sim_time = 0.05s: First KPM report arrives.** NS-3 has computed SINR/RSRP for all 20 UEs at `t=0.05s` and packs them into a KPM Indication message → FlexRIC → xApp. The xApp receives 20 UE

measurement records. Rolling window: each UE has 1/10 samples. Cooldown check: all UEs have $\text{last_ho_time} = 0 \rightarrow \text{cooldown} = 0 \rightarrow \text{NOT blocking}$. A3 check: at $t=0.05\text{s}$, all UEs are at their starting positions, none near a cell boundary. Result: No handovers triggered at this step.

11. **sim_time = 0.10s through 0.50s:** Reports 2-10 arrive. Rolling windows build up. At 0.50s, all 20 UEs have exactly 10 samples — the GRU window is now full. From this point, every A3 trigger will use the GRU for prediction rather than pure A3 fallback.

Phase 3: First Handover Event (Steps 12-17)

12. **sim_time = 10.05s: First A3 event fires.** UE 5's KPM report shows: serving cell 1 SINR = 15.8 dB, neighbor cell 3 SINR = 18.0 dB. $\Delta = 18.0 - 15.8 = 2.2 \text{ dB} > \text{A3_COMBINED threshold of } 2.0 \text{ dB} \rightarrow \text{A3 fires}$. Cooldown check: $\text{last_ho_time}[5] = 0$, $\text{current_time} = 10.05 \rightarrow 10.05 - 0 > 5.0 \rightarrow \text{NOT in cooldown}$. Window check: $\text{window_full}[5] = \text{True}$ (has 10 samples). Proceed to GRU.
13. **xApp calls GRU service:** HTTP POST to `http://127.0.0.1:5000/predict` with the 10-sample window for UE 5. The JSON payload includes 10 timesteps of SINR, RSRP, velocity, and all neighbor cell SINRs. Latency: <5ms on localhost.
14. **GRU service responds:** The GRU processes the 10-sample sequence, outputs 7 cell probabilities via softmax. Response JSON: `{"cell_id": 3, "confidence": 0.87, "all_probs": [0.01, 0.05, 0.87, 0.03, 0.02, 0.01, 0.01]}`. Cell 3 wins with 87% confidence.
15. **xApp validates GRU recommendation:** Confirms that Cell 3 also passes the A3 threshold (it does: $18.0 > 15.8 + 2.0 = 17.8 \text{ dB}$). Recommendation accepted.
16. **xApp sends RC handover command:** Constructs E2SM-RC Control message specifying UE RNTI for UE 5, target cell ID = 3. Sends via FlexRIC to NS-3.
17. **NS-3 executes handover:** UE 5 disconnects from Cell 1, attaches to Cell 3. Event logged to `handover.csv`: `10.05, 5, 1, 3, A3_HO, True`. xApp updates state: $\text{last_ho_time}[5] = 10.05$, $\text{serving_cell}[5] = 3$.

Phase 4: Long Simulation Run (Steps 18-20)

18. **sim_time = 10.05s to 120.0s:** The simulation continues. KPM reports arrive every 0.05 sim-seconds (2,400 total rounds). For each round, the xApp evaluates all 20 UEs — checking cooldown, checking A3, calling GRU when appropriate. Over 120 sim-seconds, 309 handovers are executed and logged.
19. **sim_time = 120.0s: Simulation ends.** NS-3 prints 'Simulation complete' and exits. The E2 SCTP connection closes gracefully. FlexRIC logs the disconnection.
20. **controller.py auto-saves results** to: `3D_GUI_Sim_Results/sim011_20260505_210259_gru_scenario/`
This directory contains: `decision_log.csv`, `summary.json`, `sinr_history.csv`, `metadata.json`.

Phase 5: Shutdown and Post-Processing (Steps 21-23)

21. **User runs: bash kill_sim.sh** — kills all background processes in order: NS-3 (already dead), `sim_data_pusher.py`, `xapp_handover_gru`, `nearRT-RIC`, `gru_xapp.py`, Docker.
22. **generate_plots.py runs automatically** (called by `kill_sim.sh` or `controller.py`): reads `decision_log.csv`, computes PP rate and accuracy, generates 4 plots in the `plots/` subdirectory.

23. Final result: 309 total handovers, 10 ping-pong events, 3.24% PP rate, 96.76% accuracy. Results visible in 3D GUI at port 3001 and Grafana at port 8000.

N5.3 Timeline Summary

sim_time	Real time (approx.)	Event
0.000s	0 min	Simulation starts. 20 UEs begin moving. E2 handshake complete.
0.050s	10 sec	First KPM report. Windows start filling (1/10 samples each UE).
0.500s	100 sec (~1.7 min)	All windows full (10/10 samples). GRU now active for all UEs.
10.05s	~33 min	First handover: UE 5, Cell 1 → Cell 3. GRU confidence 87%.
~30s	~100 min	Approx. 80 handovers completed. Simulation in full swing.
~60s	~200 min (~3.3 hr)	Halfway point. ~155 HOs, ~5 PPs so far.
120.0s	~400 min (~6.7 hr)	Simulation ends. Final: 309 HOs, 10 PPs.

What Each Output File Contains — Complete Reference

N6.1 Overview of Output Files

Each simulation run produces a collection of output files that capture different aspects of the simulation results. Understanding what each file contains is important for analyzing results, debugging problems, and presenting findings to your committee. All per-run files are saved to a timestamped directory in 3D_GUI_Sim_Results/:

```
3D_GUI_Sim_Results/sim011_20260505_210259_gru_scenario/
■■■ handover.csv
■■■ decision_log.csv
■■■ decision_summary.json
■■■ lstm_features.csv (gitignored – up to 7.8 MB)
■■■ summary.txt
■■■ plots/
■■■ decision_quality.png
■■■ handovers_over_time.png
■■■ ho_per_ue.png
■■■ ho_activity.png
```

N6.2 handover.csv — The Raw Event Log

This is the primary output file written directly by NS-3 during the simulation. Every handover execution is appended to this file in real time. NS-3 does NOT buffer this file — each row is written immediately when the handover executes. This means you can monitor the simulation progress by watching this file grow: `watch -n 5 wc -l handover.csv`

Column	Type	Example Value	Description
time_sec	float	10.05	Simulation time when handover was executed
ue_id	integer	5	Which UE performed the handover (0-19)
from_cell	integer	1	Cell the UE was attached to before handover
to_cell	integer	3	Cell the UE attached to after handover
event	string	A3_HO	Type of event: A3_HO (GRU-assisted) or FALLBACK_HO (pure A3)
executed_ok	boolean	True	Whether NS-3 successfully completed the handover procedure

Before each new simulation run, `kill_sim.sh` resets `handover.csv` to a header-only file (containing only the column names). This prevents data from previous runs from being mixed with the new run's data. Forgetting to reset this file is a common mistake that inflates HO counts.

N6.3 `decision_log.csv` — The Analysis Log

This file is created by the post-processing step (`generate_plots.py` or `controller.py`) by reading `handover.csv` and adding the ping-pong classification. Unlike `handover.csv`, `decision_log.csv` is written after the simulation completes, not during it.

Column	Type	Example	Description
<code>time</code>	float	10.05	Simulation time of handover decision
<code>ue_id</code>	integer	5	Which UE
<code>from_cell</code>	integer	1	Source cell
<code>to_cell</code>	integer	3	Target cell
<code>is_pingpong</code>	boolean	False	True if UE returned to <code>from_cell</code> within 5.0 sim-seconds

The **`is_pingpong`** column is determined by the following logic: for each handover from cell A to cell B at time T, search forward in time for any handover from cell B back to cell A for the same UE within T + 5.0 sim-seconds. If found, mark both HOs as `is_pingpong = True`. The accuracy calculation uses this column: $\text{Accuracy} = \text{count}(\text{is_pingpong}=\text{False}) / \text{total_rows} \times 100\%$.

N6.4 `decision_summary.json` — Machine-Readable Results

This JSON file contains all key metrics in a structured format, readable by both humans and the 3D GUI frontend. The frontend reads this file to populate the results panel shown after each simulation.

```
{
  "total_handovers": 309,
  "ping_pong_count": 10,
  "pp_rate_percent": 3.24,
  "accuracy_percent": 96.76,
  "sim_time_seconds": 120,
  "num_ues": 20,
  "num_cells": 7,
  "gru_model": "gru_model.h5",
  "a3_offset_db": 1.0,
  "cooldown_sec": 5.0
}
```

N6.5 `Istm_features.csv` — The Raw GRU Input Log (Gitignored)

This file records every feature vector that was sent to the GRU service, along with the GRU's response. It is primarily used for debugging and post-hoc model analysis. Each row represents one GRU prediction call: 10 SINR values, 10 RSRP values, velocity, and the GRU's cell recommendation and confidence score. At 309 HO across 120 sim-seconds, this file grows to approximately 7.8 MB — too large to commit to Git (hence gitignored). If you need to retrain the model with new data, this file is the source.

Important: *lstm_features.csv must be cleared before each simulation run, or it will accumulate data from multiple runs, making retraining impossible without knowing which rows belong to which run. kill_sim.sh handles this reset automatically.*

N6.6 summary.txt — Human-Readable Quick Overview

A plain-text file with 5-7 lines summarizing the key results. Designed for quick inspection without opening CSV files or a GUI. Example content:

```
Simulation: sim011 | Duration: 120s | UEs: 20 | Cells: 7
Total Handovers: 309
Ping-Pong Events: 10
PP Rate: 3.24%
GRU Accuracy: 96.76%
Generated: 2026-05-05 21:02:59
```

N6.7 plots/ Directory — The Four Visualization Files

File	Chart Type	X-Axis	Y-Axis	What It Shows
decision_quality.png	Pie chart	N/A	N/A	Proportion of correct HOs (blue, 96.76%) vs ping-pong HOs (red, 3.24%)
handovers_over_time.png	Line chart	sim_time (seconds)	Cumulative HO count	How HOs accumulate over the 120s simulation — slope shows HO rate
ho_per_ue.png	Bar chart	UE ID (0-19)	HO count	Which UEs performed the most handovers — reveals mobility patterns
ho_activity.png	Heatmap	sim_time bins	UE ID	When each UE was active in handovers — color intensity = HO frequency

N6.8 flexric.log — The Runtime Diagnostic Log

The FlexRIC log at /tmp/flexric.log is NOT saved in the simulation results directory — it is ephemeral and is overwritten each time FlexRIC starts. It contains low-level diagnostic messages that are essential for debugging but not for result analysis:

- **xApp registration messages:** 'xApp registered with ID X'

-
- **E2 Setup messages:** 'E2 Setup Request received from gNB ID Y'
 - **Subscription messages:** 'KPM subscription accepted for report period 50ms'
 - **Indication delivery logs:** 'KPM Indication delivered to xApp, UE count: 20'
 - **RC Control logs:** 'RC Control sent to gNB for UE RNTI Z, target cell C'
 - **Error messages:** 'E2 connection reset by peer' (appears when NS-3 is killed)

To monitor FlexRIC in real time during a simulation: `tail -f /tmp/flexric.log`

Why This System Is Novel — Research Contribution

N7.1 The State of the Art Before This Work

Before this thesis, handover management in cellular networks was dominated by the **A3 event with fixed threshold** approach. This approach, defined by 3GPP in TS 38.331, works as follows: whenever a neighboring cell's SINR exceeds the serving cell's SINR by a fixed margin (the A3 offset), a handover is triggered. This is purely reactive — it responds to what is happening NOW, with no memory of what happened in the last second, and no prediction of what will happen next.

The limitations of threshold-only A3 handover are well-documented in the literature:

- It cannot distinguish between a sustained SINR improvement (UE genuinely moving toward new cell) and a momentary spike (shadow fading artifact)
- It produces ping-pong rates of 8-15% in mmWave environments where channel conditions vary rapidly
- The optimal threshold value depends on the specific deployment and UE speed — a single fixed value cannot be optimal for all scenarios
- It makes the same decision whether the UE is moving toward the candidate cell or parallel to the boundary

N7.2 What This Work Adds

Our thesis makes the following novel contributions:

- 1. GRU-based proactive handover prediction:** We integrate a trained GRU neural network into the handover decision loop as a validator. The GRU processes the 10-sample SINR/RSRP history to predict not just which cell is currently stronger, but which cell will be best given the UE's movement trend. This converts the handover decision from reactive (current snapshot) to proactive (trend-aware prediction).
- 2. Real O-RAN near-RT RIC integration:** Unlike most ML-handover research papers that simulate the entire system in a closed environment (no real O-RAN stack), our system uses FlexRIC — a real near-RT RIC implementation — as the control plane. The xApp runs inside FlexRIC's framework and communicates via the actual E2AP/E2SM protocols. This makes our results directly transferable to real O-RAN deployments.
- 3. Cooldown + GRU combination:** We demonstrate that combining a hard cooldown timer (preventing re-evaluation within 5 seconds) with GRU prediction achieves lower PP rates than either mechanism alone.
- 4. Fully open-source and reproducible:** All components (FlexRIC, NS-3, GRU xApp, Python service, visualization stack) are open-source. The exact commands and parameters needed to reproduce our results are documented in `gru.sh` and `MANUAL_COMMANDS.txt`.

N7.3 Comparison with Related Work

Paper / System	Method	Typical PP Rate	Platform	O-RAN Integration
Traditional 3GPP A3 only	Fixed threshold	8-15%	Real LTE networks	N/A (pre-O-RAN)
Banoula et al. (2020)	DQN-based HO	~5%	Simulation only	None
Khosravi et al. (2021)	LSTM + A3 hybrid	~4-6%	NS-3 simulation	None (custom controller)
Liu et al. (2022)	Transformer attention HO	~3.5%	Simulation	None
This work (GRU + O-RAN)	GRU + A3 + cooldown	1.72-3.65%	NS-3 + FlexRIC	Full E2AP/E2SM-KP M/RC

The key differentiator of our work is the last column: full O-RAN integration via FlexRIC. Most papers in this area simulate a simplified controller without any real O-RAN stack. Our system demonstrates that GRU-based handover optimization can run within the actual near-RT RIC framework at real O-RAN timing constraints.

N7.4 Why GRU Over LSTM for This Application

Criterion	LSTM	GRU (Our Choice)	Why GRU Wins Here
Parameters	More (3 gates)	Fewer (2 gates)	Faster training and inference — critical for near-RT
Inference latency	~5-8ms	~2-4ms	GRU fits within near-RT 10ms latency budget more comfortably
Sequence length	Very long (100+)	Moderate (10-30)	Our 10-sample window is GRU's sweet spot
Accuracy on our task	Comparable	96.76% (our result)	No measurable accuracy gap for short sequences
Memory capacity	Higher	Sufficient	10-sample window does not need LSTM's extra capacity

N7.5 Why Near-RT RIC (Not Non-RT RIC)

A common committee question is: why run the handover xApp in the near-RT RIC and not the non-RT RIC? The answer is timing. A handover decision has a latency budget determined by how fast channel conditions change:

- At 3 m/s pedestrian speed, a UE moves ~3 meters per second. At 28 GHz, the cell edge covers perhaps 10-20 meters. The UE can cross the cell edge in 3-7 seconds.

-
- A non-RT RIC (>1 second latency) could miss the optimal handover moment if the UE crosses the boundary and the SINR drops rapidly.
 - The near-RT RIC (10ms-1s latency) can react within one SINR sample period (50ms) of the A3 trigger — well within the handover execution window.
 - The real-time embedded controller (<10ms) is theoretically faster, but it is deeply integrated into the O-DU hardware and cannot run ML inference without specialized hardware acceleration (FPGAs, dedicated ML chips).

The near-RT RIC is the correct home for handover ML because it provides the right latency range: fast enough for real-time decisions, slow enough to run Python ML inference on a general-purpose server CPU.

Common Mistakes and How to Avoid Them

These are real issues that occurred during the development and testing of this system. Each mistake is documented with: what symptom it causes, how to detect it, and how to fix it. Understanding these issues demonstrates depth of knowledge about the system internals and is likely to impress your committee if such questions arise.

Mistake 1: xApp Starts AFTER NS-3

What happens: When NS-3 starts before the xApp, FlexRIC has not yet received the KPM subscription request. NS-3 connects to FlexRIC and waits for subscriptions to be registered. If the xApp connects 10-20 seconds later, the first 10-20 seconds of simulation time produce NO KPM reports to the xApp. The rolling window for all 20 UEs remains empty. For the first several hundred KPM cycles, every A3 evaluation falls back to pure A3 (no GRU) because `window_full[ue_id] = False` for all UEs.

How to detect it: Check `/tmp/flexric.log`. If you see NS-3 E2 Setup messages before any xApp registration messages, the order was wrong. Also: if the first HO in `handover.csv` has `event = FALLBACK_HO` instead of `A3_HO`, the GRU was not ready.

How to fix it: Follow the exact startup order in `gru.sh`: GRU service first, then FlexRIC, then xApp, then NS-3. Each step has a mandatory wait (2-3 seconds) before the next. Never change this order without thorough testing.

Mistake 2: kill_sim.sh Kills Processes Before Save Completes

What happens: When `kill_sim.sh` sends `SIGTERM` to the processes, `controller.py` may be in the middle of writing the `decision_summary.json` or generating plots. If the process is killed before the file write completes, the JSON file is truncated (invalid JSON) and the results are lost. The 3D GUI then shows no results for that run.

How to detect it: After running `kill_sim.sh`, check if `decision_summary.json` is valid: `python3 -c "import json; json.load(open('decision_summary.json'))"`. If it raises an exception, the file was corrupted.

How to fix it: The `generate_plots.py` fallback script can regenerate all output files from `handover.csv` alone. Run: `python3 generate_plots.py handover.csv`. `kill_sim.sh` was updated with a 2-second delay before killing `controller.py`.

Mistake 3: handover.csv Not Reset Before New Run

What happens: NS-3 opens `handover.csv` in append mode. If the file already contains 309 rows from `sim011` when `sim012` starts, the new run will append to the existing data. The result file for `sim012` will contain 309 + `new_HOs` rows, making the HO counts appear inflated and the statistics incorrect.

How to detect it: Compare the timestamp of the first row in `handover.csv` with the simulation start time. If the first row's `sim_time` is 10.05 but the simulation was just started fresh, the file was not reset.

How to fix it: `kill_sim.sh` resets `handover.csv` to header-only before each run. Manually: `echo "time_sec,ue_id,from_cell,to_cell,event,executed_ok" > handover.csv`

Mistake 4: Port 5000 Already in Use

What happens: If another process is already using port 5000 (another Flask app, a previous `gru_xapp.py` that was not killed, etc.), `gru_xapp.py` fails to bind and exits immediately. The xApp attempts to connect to port 5000 and gets 'Connection refused.' The xApp falls back to pure A3 for ALL handover decisions — GRU is never used.

How to detect it: Run: `ss -ltnp | grep 5000`. If another process is listed, kill it first. Also check `gru_xapp.py`'s terminal output — it will show an `OSError` if port binding fails.

How to fix it: Run: `kill $(lsof -t -i:5000)` to kill whatever is using port 5000. Or: `kill_sim.sh` handles this automatically.

Mistake 5: FlexRIC Log Grows Between Runs Without Clearing

What happens: FlexRIC appends to `/tmp/flexric.log` every time it runs. After 5-6 simulation runs, the log file may contain thousands of lines from previous sessions. When diagnosing a problem, the relevant messages from the current run are buried under old messages. Timestamps help, but not if you are scrolling through thousands of lines.

How to fix it: Clear the log before each run: `echo '' > /tmp/flexric.log`. `kill_sim.sh` does this automatically as part of cleanup.

Mistake 6: lstm_features.csv Not Cleared Between Runs

What happens: `lstm_features.csv` accumulates GRU input/output data across multiple runs. A single 120-second simulation with 309 HOs produces approximately 7.8 MB of data. After three runs, the file is 23 MB. After ten runs, 78 MB. This makes IO slower (Python CSV reads take longer) and makes retraining ambiguous (which rows belong to which run?).

How to fix it: `kill_sim.sh` truncates `lstm_features.csv` to header-only. Also gitignored to prevent accidental large commits.

Mistake 7: Docker Not Started Before FlexRIC

What happens: `sim_data_pusher.py` attempts to write to InfluxDB via HTTP. If the InfluxDB Docker container is not running, every write attempt fails with 'Connection refused.' The pusher logs errors and may crash. Without InfluxDB, Grafana has no data to display — the 2D visualization dashboard shows empty panels.

How to detect it: Run: `docker ps`. You should see influxdb, grafana, and nginx containers listed as 'Up'.
Also test: `curl http://localhost:8086/health` should return `{"status": "pass"}`.

How to fix it: `gru.sh` starts Docker first and waits 3 seconds before starting FlexRIC. If Docker was already stopped, run: `docker-compose -f GUI/docker-compose.yml up -d` and wait 10 seconds.

Extended Glossary — 80+ Terms Defined

This extended glossary defines every technical term used in this thesis and its surrounding fields. Each definition is written in plain language with enough context to allow you to use the term correctly in your defense. Terms are organized alphabetically.

A3 Event: A handover measurement event defined in 3GPP TS 38.331. Named 'A3' in the measurement event classification. Fires when a neighboring cell's SINR exceeds the serving cell's SINR by $A3_OFFSET + A3_HYSTERESIS$ (2.0 dB in our system). It is the standard trigger for handover in 4G/5G networks.

A3 Offset: The minimum SINR gain required for an A3 event to fire. In our system: 1.0 dB. Prevents handovers for negligible improvements.

A3 Hysteresis: An additional margin on top of $A3_OFFSET$ that prevents oscillation due to measurement noise. In our system: 1.0 dB. The total A3 threshold = $A3_OFFSET + A3_HYSTERESIS = 2.0$ dB.

Accuracy: In our context: the percentage of handover decisions that did NOT result in ping-pong. Accuracy = 100% - PP_Rate. Our sim011 accuracy = 96.76%.

Activation Function: A mathematical function applied to the output of each neural network neuron. Common choices: sigmoid (outputs 0-1), tanh (outputs -1 to +1), ReLU (outputs 0 to infinity). GRU uses sigmoid for gates and tanh for candidates.

A1 Interface: The O-RAN interface connecting the non-RT RIC to the near-RT RIC. Used to send AI policies and ML model updates downward. Not used in our current system (GRU model loaded from disk at startup).

Backpropagation Through Time (BPTT): The algorithm used to train recurrent neural networks (RNNs, GRUs, LSTMs). It unrolls the network through time steps and computes gradients by working backward from the output error. BPTT is where the vanishing gradient problem manifests.

Bandwidth: The range of radio frequencies used for communication. Our simulation uses 400 MHz bandwidth at 28 GHz. More bandwidth means higher maximum data rates.

Baseband Unit (BBU): The traditional 4G/5G hardware box that handles all digital signal processing. O-RAN splits the BBU into O-DU and O-CU to enable openness and flexibility.

Beam: A narrow, focused radio transmission in a specific direction, produced by phased array antennas. mmWave 5G uses beamforming to concentrate signal energy and overcome high path loss.

Beamforming: The technique of combining multiple antenna elements to create a focused beam toward a specific UE. Essential for mmWave 5G where signals lose power rapidly with distance.

Cell: The coverage area served by one base station (gNB) and its antenna. Our simulation has 7 cells. UEs connect to the cell with the best SINR.

Cell ID: An integer identifier for each cell in the simulation. Our cells are numbered 0 through 6.

Serving Cell: The cell that a UE is currently connected to and receiving service from. The UE's serving cell SINR is the primary signal quality metric.

Neighbor Cell: Any cell other than the serving cell. Neighbor cell SINRs are measured and reported in KPM to identify better handover candidates.

Channel Model: A mathematical model describing how radio signals propagate from a transmitter to a receiver, accounting for distance, obstacles, and multipath effects. Our system uses the 3GPP TR 38.901 model.

Confidence Score: The probability output by GRU's softmax for the recommended cell. A confidence of 0.87 means the GRU assigns 87% probability to cell 3 being the best choice.

Control Plane: The part of the network responsible for signaling, configuration, and control messages (as opposed to the data plane which carries user data). Handover commands flow through the control plane.

COOLDOWN: See 'Cooldown Period'.

Cooldown Period: A 5.0 simulation-second window after each handover during which no new handover is evaluated for the same UE. Prevents ping-pong by blocking premature re-evaluation.

Cooldown Timer: A per-UE timestamp variable (`last_ho_time[ue_id]`) in the xApp that records when the UE's last handover occurred. Used to enforce the 5-second cooldown.

dB (Decibel): A logarithmic unit used to express ratios of power or signal strength. +3 dB means approximately double the power. +10 dB means 10× the power. Used for SINR.

dBm (Decibels-milliwatt): A logarithmic unit for absolute power level, referenced to 1 milliwatt. 0 dBm = 1 mW. -30 dBm = 0.001 mW. Used for RSRP. Lower (more negative) values indicate weaker signal.

Dense Layer: A fully connected neural network layer where every input neuron connects to every output neuron. The GRU model has a Dense + softmax final layer that maps the hidden state to 7 cell probabilities.

Discrete-Event Simulation: A simulation paradigm where the model jumps from one event to the next (packet arrival, timer expiry, UE movement) rather than advancing in fixed time steps. NS-3 uses this approach.

Docker: A container platform that packages software with all its dependencies. In our system: InfluxDB, Grafana, and nginx run as Docker containers managed by docker-compose.

E2 Interface: The O-RAN interface connecting the near-RT RIC (FlexRIC) to E2 Nodes (NS-3 gNBs). Uses SCTP port 36421. Carries KPM monitoring reports and RC control commands.

E2AP (E2 Application Protocol): The signaling protocol layer of the E2 interface. Handles E2 Setup, Subscription, Indication, and Control message types. Encoded in ASN.1 format.

E2SM-KPM: E2 Service Model for Key Performance Measurements. Defines the format of measurement reports (SINR, RSRP, etc.) from NS-3 to FlexRIC.

E2SM-RC: E2 Service Model for RAN Control. Defines the format of control commands (handover triggers) from FlexRIC to NS-3.

FlexRIC: An open-source near-RT RIC implementation from EURECOM. Implements E2AP and provides an SDK for writing xApps in C, C++, or Python. Runs on standard Linux servers.

Flask: A lightweight Python web framework. Used in `gru_xapp.py` to create the HTTP inference service on port 5000.

FastAPI: A modern Python web framework optimized for APIs. Used in `controller.py` to expose the 3D GUI backend at port 8001.

Gate: In GRU and LSTM, a gate is a vector of values in $[0,1]$ that controls information flow. The GRU has two gates: update gate (z_t) and reset gate (r_t).

gNB: Next Generation Node B — the 5G base station that connects UEs to the network core. In our simulation, we have 7 gNBs forming the cell layout.

Grafana: An open-source data visualization platform. Reads from InfluxDB and displays real-time dashboards. Runs as a Docker container, accessible at port 8000 in our system.

GRU (Gated Recurrent Unit): A type of recurrent neural network with two gates (update and reset) that solves the vanishing gradient problem. Processes sequential data efficiently. Used in our system to predict handover targets from 10-sample SINR history.

Handover (HO): The process of transferring an active UE connection from one cell to another without dropping the session. Necessary when the UE moves away from its serving cell.

Hard Handover: A handover where the UE disconnects from the old cell BEFORE connecting to the new cell. Brief interruption (~10-50ms). Used in 5G NR.

Soft Handover: A handover where the UE connects to the new cell BEFORE disconnecting from the old cell. No interruption. Used in some 3G systems. Not used in 5G.

Hidden State: The internal memory vector of a GRU cell, updated at each timestep. Encodes information learned from all previous timesteps in the sequence. Passed to the Dense output layer after all 10 samples are processed.

HTTP: HyperText Transfer Protocol — the standard web protocol used for communication between the xApp (C) and the GRU service (Python Flask). POST requests to /predict, JSON responses.

InfluxDB: A time-series database optimized for timestamped data. Runs as a Docker container on port 8086. Receives SINR/RSRP/HO event data from sim_data_pusher.py.

Indication Message: An E2AP message type where the E2 Node (NS-3) sends measurement data to the RIC (FlexRIC). Contains a KPM or other service model payload.

JSON (JavaScript Object Notation): A lightweight, human-readable data format used for API communication. The xApp sends JSON POST requests to gru_xapp.py, which responds with JSON.

KPM (Key Performance Measurement): The O-RAN measurement service model for monitoring. Defines which metrics (SINR, RSRP, cell load) are reported and at what frequency.

KPM Report: A measurement snapshot sent from NS-3 to FlexRIC every 0.05 sim-seconds. Contains SINR/RSRP for all 20 UEs across all 7 cells.

Latency: The time delay between an event occurring and a response being sent. In O-RAN: non-RT latency >1s, near-RT 10ms-1s, real-time <10ms.

LSTM (Long Short-Term Memory): A type of RNN with three gates (input, forget, output) that solves the vanishing gradient problem. More complex than GRU with more parameters. GRU was chosen over LSTM for our system due to lower inference latency.

Loss Function: A mathematical function that measures how wrong the neural network's prediction is compared to the ground truth. We use categorical cross-entropy. Lower loss = more accurate model.

Learning Rate: A hyperparameter controlling how large each weight update step is during training. Too high: unstable training. Too low: very slow convergence. Adam optimizer adapts learning rates automatically.

mmWave: Millimeter Wave — radio frequencies above 24 GHz (wavelength: 1-10mm). Used in 5G NR for high bandwidth. Shorter range and more easily blocked than sub-6 GHz.

MIMO (Multiple-Input Multiple-Output): Using multiple antennas at both transmitter and receiver to increase data rate and reliability. Essential for mmWave 5G to compensate for high path loss.

Mobility Model: The mathematical model describing how UEs move. Our system uses Random Walk 2D: UEs move in random directions at random speeds, bouncing off boundaries.

Near-RT RIC: Near Real-Time RAN Intelligent Controller. Operates on 10ms-1s timescales. Runs xApps. In our system: FlexRIC.

Non-RT RIC: Non-Real-Time RIC. Operates on >1 second timescales. Runs rApps for long-term optimization. Not directly used in our system.

NS-3 (Network Simulator 3): A free, open-source discrete-event network simulator written in C++. Used to simulate our 7-cell mmWave network with 20 UEs. NS-3's mmWave module implements the 28 GHz channel model.

O-RAN (Open Radio Access Network): A mobile network architecture with open, standardized interfaces that enables multi-vendor deployments and custom software intelligence. The O-RAN Alliance defines the specifications.

O-CU (Open Central Unit): The O-RAN component handling higher-layer protocols (RRC, PDCP). Can run in a central data center.

O-DU (Open Distributed Unit): The O-RAN component handling real-time lower-layer processing, including the radio scheduler. Runs close to the antenna.

O-RU (Open Radio Unit): The O-RAN component handling radio frequency processing. Physically located at the antenna tower.

Ping-Pong: A handover event where a UE moves from Cell A to Cell B, then back from B to A within 5.0 simulation-seconds. Indicates an unnecessary handover.

Port Number: A 16-bit number identifying a specific service on a machine. Key ports in our system: 5000 (GRU service), 8086 (InfluxDB), 8000 (Grafana), 8001 (FastAPI), 3001 (Vite), 36421 (FlexRIC SCTP).

PP Rate: Ping-Pong Rate = $(\text{Ping-Pong HOs} / \text{Total HOs}) \times 100\%$. Our sim011 PP rate = 3.24%.

RAN (Radio Access Network): The part of the mobile network that connects UEs to the core network via radio links. Consists of base stations (gNBs) and their associated infrastructure.

rApp: An application running in the non-RT RIC, performing long-term optimization and policy management. Different from xApps which run in the near-RT RIC.

RIC (RAN Intelligent Controller): A software platform for running intelligence (xApps, rApps) on top of the radio network. O-RAN defines near-RT RIC and non-RT RIC.

RSRP (Reference Signal Received Power): The absolute power level of a cell's reference signal at the UE, in dBm. Indicates signal strength but not quality (unlike SINR which also accounts for interference).

RSRQ (Reference Signal Received Quality): A quality metric combining RSRP and interference level. Less commonly used than SINR in our system.

SCTP (Stream Control Transmission Protocol): A transport protocol used for the E2 interface. Supports multiple streams and message boundaries. Port 36421. Better than TCP for signaling traffic.

SINR (Signal-to-Interference-plus-Noise Ratio): The ratio of desired signal power to interference plus noise power, measured in dB. The primary metric for channel quality and handover decisions.

Softmax: A mathematical function that converts raw output scores to probabilities summing to 1.0. The GRU model's final activation, producing 7 cell probabilities.

State Machine: A computational model with defined states and transitions. Our xApp maintains a state per UE: IDLE, COOLDOWN, or HO_REQUESTED.

TTT (Time-to-Trigger): In 3GPP standards: the time that a measurement event must remain true before triggering a handover. In our system, we use a cooldown instead. 3GPP recommends TTT of 40-1280ms for various scenarios.

Training Data: The dataset of historical examples used to adjust the GRU model's weights. Our training data was collected by Fares from preliminary simulation runs.

UE (User Equipment): Any device connecting to the mobile network: smartphone, laptop, IoT sensor. In NS-3, we simulate 20 UEs moving with Random Walk 2D mobility.

Update Gate: One of the two GRU gates. Controls how much of the new candidate hidden state replaces the old hidden state. Update gate close to 1 = fully update; close to 0 = preserve old state.

Vanishing Gradient: A problem in training RNNs where gradients become very small as they propagate backward through many timesteps, preventing the network from learning long-range dependencies. GRU solves this with its gating mechanism.

Vite: A fast JavaScript build tool and development server. Used to serve the 3D visualization React/Three.js frontend on port 3001.

Weight Matrix: In neural networks: a matrix of learned parameters (W or U in GRU equations) that transforms input or hidden state vectors. Updated during training to minimize the loss function.

xApp: An application running inside the near-RT RIC. Subscribes to E2 service models (KPM, RC) and performs real-time network optimization. Our xApp (xapp_handover_gru, built from best2.c) handles handover decisions.

X2 Interface: The 4G LTE interface between neighboring eNBs for inter-cell handover coordination. The 5G equivalent is the Xn interface. Different from the E2 interface (which connects RIC to gNBs).

Self-Test Questions — 30 Questions Before Your Defense

Answer these 30 questions without looking at your notes. If you can answer all of them correctly and confidently, you are ready for your defense. Cover the answers with a piece of paper, read the question, formulate your answer mentally, then reveal and compare. Pay special attention to any question where your mental answer differs from the written answer — these are your weak spots to review.

Questions 1-10: System Architecture and Parameters

Q: What does O-RAN stand for?

A: Open Radio Access Network. An architecture that breaks vendor lock-in in mobile networks by defining open, standardized interfaces so equipment from different manufacturers can interoperate.

Q: What is the E2 interface?

A: The O-RAN interface connecting the near-RT RIC (FlexRIC in our system) to E2 Nodes (NS-3 gNBs). Carries KPM measurement reports and RC control commands over SCTP protocol on port 36421.

Q: What port does FlexRIC use?

A: SCTP port 36421. This is the IANA-registered port for S1AP/E2AP protocols. All E2 communication (both KPM reports from NS-3 and RC commands to NS-3) passes through this port.

Q: What is A3_OFFSET in our system?

A: 1.0 dB. This is the minimum SINR gain required for the A3 event to fire. Combined with A3_HYSTERESIS (also 1.0 dB), the total A3 trigger threshold is 2.0 dB above the serving cell SINR.

Q: What is the cooldown time?

A: 5.0 simulation-seconds. After any handover, the xApp will not evaluate any new handover for the same UE for 5.0 sim-seconds. This prevents rapid ping-pong oscillation.

Q: How many UEs in our simulation?

A: 20 UEs. All 20 use the Random Walk 2D mobility model and move at approximately pedestrian speed (~3 m/s).

Q: How many cells?

A: 7 gNBs (cells), numbered 0 through 6. They are arranged in a pattern covering the simulation area. Each cell uses 28 GHz mmWave frequencies with 400 MHz bandwidth.

Q: What is the GRU window size?

A: 10 samples. This rolling window contains the last 10 KPM reports for each UE. At 0.05 sim-seconds per report, the window covers 0.5 simulation-seconds of recent channel history.

Q: What port does gru_xapp.py listen on?

A: Port 5000. This is the standard Flask development server port. The xApp communicates with gru_xapp.py via HTTP POST requests to `http://127.0.0.1:5000/predict`.

Q: What is the ping-pong rate in sim011?

A: 3.24%. Calculated as $(10 \text{ ping-pong HOs} / 309 \text{ total HOs}) \times 100\% = 3.236\%$, rounded to 3.24%.

Questions 11-20: GRU and Machine Learning

Q: What is GRU accuracy in sim011?

A: 96.76%. Calculated as $(299 \text{ correct HOs} / 309 \text{ total HOs}) \times 100\% = 96.76\%$. This equals 100% minus the PP rate of 3.24%.

Q: What does KPM stand for?

A: Key Performance Measurement. The E2 service model (E2SM-KPM) defines the format and content of measurement reports sent from NS-3 to FlexRIC every 0.05 sim-seconds.

Q: What does RC stand for in E2SM-RC?

A: RAN Control. The E2 service model (E2SM-RC) defines the format of control commands, including handover triggers, sent from the xApp via FlexRIC to NS-3.

Q: How often does NS-3 send KPM reports?

A: Every 0.05 simulation-seconds (50 milliseconds in simulated time), or 20 reports per simulation-second. Each report contains measurements for all 20 UEs.

Q: What database stores real-time metrics for visualization?

A: InfluxDB. It is a time-series database running as a Docker container on port 8086. `sim_data_pusher.py` pushes SINR/RSRP/HO data to InfluxDB, which Grafana then queries for dashboards.

Q: What is the simulation pace?

A: Approximately 0.005 simulation-seconds per real-second. This means a 60 sim-second simulation takes about 12,000 real seconds (~3.3 hours). The 120s sim011 took approximately 6-7 hours.

Q: Why is lstm_features.csv gitignored?

A: Because it grows to approximately 7.8 MB per simulation run. Git repositories should not store large binary/data files that change every run. The file can be regenerated by running the simulation again.

Q: What tool generates the final plots?

A: generate_plots.py. It reads decision_log.csv (or handover.csv) and produces 4 matplotlib plots: decision_quality.png, handovers_over_time.png, ho_per_ue.png, and ho_activity.png.

Q: What does the reset gate in GRU do?

A: The reset gate r_t decides how much of the previous hidden state h_{t-1} is relevant when computing the new candidate hidden state. r_t close to 0: ignore past history. r_t close to 1: use past history fully. It allows the GRU to selectively forget irrelevant context.

Q: Why use near-RT RIC instead of non-RT RIC for handover decisions?

A: Handover decisions must be made within 10ms-1s of an A3 trigger to avoid signal degradation. The near-RT RIC operates in this 10ms-1s window. The non-RT RIC (>1 second) is too slow — the UE might already be disconnected by the time the decision arrives.

Questions 21-30: Protocols, Files, and Operations

Q: What is SCTP and why is it used instead of TCP for E2?

A: Stream Control Transmission Protocol. It is used instead of TCP because: (1) SCTP supports multiple independent streams — one lost message does not block others. (2) SCTP preserves message boundaries natively (TCP is a byte stream). (3) SCTP supports multi-homing for redundancy. Port 36421 is IANA-registered for this use.

Q: What happens if the GRU service (gru_xapp.py) is down?

A: The xApp falls back to pure A3 decision: it picks the neighbor cell with the highest SINR among those that pass the A3 threshold (neighbor SINR > serving SINR + 2.0 dB). The fallback ensures handovers continue working even without the GRU model.

Q: What is the A3 trigger formula?

A: $\text{neighbor_SINR} > \text{serving_SINR} + \text{A3_OFFSET} + \text{A3_HYSTERESIS}$, which in numbers is: $\text{neighbor_SINR} > \text{serving_SINR} + 1.0 + 1.0 = \text{neighbor_SINR} > \text{serving_SINR} + 2.0 \text{ dB}$.

Q: How many total handovers in sim011?

A: 309 total handover events were logged in handover.csv. The simulation ran for 120 simulation-seconds with 20 UEs. On average: $309 / 120 / 20 = 0.129$ HO per UE per sim-second.

Q: How many ping-pong events in sim011?

A: 10 ping-pong events. These are handovers where the UE returned to the previous cell within 5.0 simulation-seconds. $10 \text{ out of } 309 = 3.24\%$ PP rate.

Q: What is SINR?

A: Signal-to-Interference-plus-Noise Ratio. The ratio of desired signal power to (interference from other cells + thermal noise), measured in dB. The primary metric for channel quality. Values above 10 dB are good; below 0 dB is poor.

Q: What is RSRP?

A: Reference Signal Received Power. The absolute power level of a cell's reference signal at the UE, measured in dBm. Indicates signal strength but not quality. Values above -80 dBm are excellent; below -110 dBm are very poor.

Q: What language is the xApp written in?

A: C. The xApp source file is best2.c, compiled into the executable xapp_handover_gru. C was chosen because FlexRIC's SDK is C-native and provides the best performance for real-time message processing.

Q: What framework is the 3D GUI built with?

A: Vite + React + Three.js for the frontend (port 3001), and FastAPI for the backend controller.py (port 8001). React handles the UI components, Three.js handles the 3D visualization, and FastAPI exposes the REST API.

Q: What does 'bash kill_sim.sh' do?

A: It kills all simulation processes in reverse order: NS-3 first, then sim_data_pusher.py, then xapp_handover_gru, then nearRT-RIC (FlexRIC), then gru_xapp.py, and finally the Docker containers. It also resets handover.csv and lstm_features.csv to header-only files, and generates the final plots.

Key Point: If you answered all 30 questions correctly without looking at notes, you are ready to defend. Focus extra review time on any questions where you hesitated or gave an incomplete answer. Remember: the committee is more impressed by an honest 'I am not sure, but my best understanding is...' followed by a reasonable answer, than by a confident wrong answer.

GRU Model Training — How the Model Was Built

N11.1 The Training Pipeline Overview

The GRU model used in our system was trained by Fares as a companion component to the xApp infrastructure built by Omar. The training pipeline follows the standard supervised machine learning workflow: (1) collect labeled training data from simulation runs, (2) preprocess and normalize features, (3) define the model architecture, (4) train the model on labeled examples, (5) evaluate on held-out test data, (6) save the trained model for deployment.

The critical insight that makes supervised training possible: in a simulation environment, we KNOW the ground truth. When NS-3 moves a UE to Cell 3 based on an external handover command, and the UE then stays in Cell 3 for the next 10+ sim-seconds without ping-pong, we can label that handover as 'correct — Cell 3 was the right choice.' This labeled data becomes the training signal for the GRU model.

N11.2 Training Data Collection

Training data was collected from multiple preliminary simulation runs using a pure A3 handover policy (no GRU, just threshold-based decisions). For each handover event, the system recorded:

- **Features:** The 10-sample window of SINR/RSRP/velocity values immediately BEFORE the handover decision (this is the GRU input)
- **Label:** Which cell was 'correct' — defined as the cell that the UE remained attached to for the longest subsequent period
- **Negative examples:** Cases where the pure A3 decision caused a ping-pong — these were labeled with the ACTUAL correct cell (the one the UE eventually settled at) rather than the cell chosen by A3

Training Data Property	Value	Explanation
Total training samples	~5,000 labeled HO events	From multiple 120s simulation runs
Input shape per sample	(10, 11)	10 timesteps × 11 features
Output classes	7	One per cell (cells 0-6)
Train/Validation split	80% / 20%	Standard practice for generalization testing
Class balance	Approximately uniform	All 7 cells roughly equally represented as optimal choices
Data augmentation	None applied	NS-3 random seeds provide sufficient diversity

N11.3 Model Architecture Choices

The final GRU architecture was chosen after experimentation with several configurations. The key design decisions:

Input Layer

Shape (10, 11): 10 timesteps (the rolling window) × 11 features (serving SINR, serving RSRP, velocity, serving cell ID encoded as integer, and 7 neighbor cell SINRs). The input is masked to handle sequences shorter than 10 samples during startup.

GRU Layer

One GRU layer with 64 hidden units was found sufficient for this task. Increasing to 128 units improved validation accuracy by less than 0.5% while roughly doubling inference latency. The 64-unit model achieves <3ms inference time on a standard Intel Core i7 CPU, well within the near-RT latency budget.

Dense Output Layer

A single Dense layer with 7 units followed by softmax activation. No hidden Dense layers were needed — the GRU's hidden state already provides a rich representation. Adding intermediate Dense layers did not improve accuracy on the validation set.

Dropout

A dropout layer with rate 0.2 after the GRU was used during training to prevent overfitting. Dropout is disabled during inference. Without dropout, the model showed signs of memorizing the training data (high training accuracy, lower validation accuracy).

N11.4 Training Hyperparameters

Hyperparameter	Value	Justification
Optimizer	Adam	Adaptive learning rates — standard for GRU/LSTM training
Learning rate	0.001 (initial)	Standard Adam default; reduced by scheduler if loss plateaus
Batch size	32	Balances GPU utilization and gradient noise
Epochs	50 (early stopping)	Training stops when validation loss does not improve for 5 epochs
Loss function	Categorical cross-entropy	Standard for multi-class classification
Metrics	Accuracy, top-2 accuracy	Top-2 checks if correct cell is in top 2 predictions
Regularization	Dropout 0.2 on GRU output	Prevents overfitting to specific trajectories

N11.5 Feature Normalization Details

Neural networks require normalized input features. Without normalization, features with large absolute values (like RSRP in the range -140 to -44 dBm) would dominate over features with small ranges (like velocity at 0-5 m/s), causing the model to ignore the smaller-range features.

We use **StandardScaler normalization**: for each feature, subtract the mean and divide by the standard deviation computed from the training set.

$$\text{normalized_feature} = (\text{raw_value} - \text{mean_feature}) / \text{std_feature}$$

Feature	Typical Mean	Typical Std Dev	After Normalization
serving_sinr	8.0 dB	7.0 dB	Values roughly in range [-2, +3]
serving_rsrp	-95 dBm	12 dBm	Values roughly in range [-3, +2]
velocity	2.5 m/s	1.2 m/s	Values roughly in range [-2, +2]
serving_cell_id	3.0 (center)	2.0	Values roughly in range [-1.5, +1.5]
neighbor_sinr_k	5.0 dB	8.0 dB	Values roughly in range [-2, +3]

The mean and std_dev values are saved alongside the model file as a JSON or NumPy file. When gru_xapp.py loads the model at startup, it also loads these normalization parameters. Every incoming KPM sample is normalized before being fed to the GRU.

N11.6 Model Evaluation Results

After training on 80% of the data and evaluating on the held-out 20% validation set:

Metric	Value	Interpretation
Training accuracy	~98%	Model fits training data well (but check for overfitting)
Validation accuracy	~94-96%	Generalizes to unseen trajectories
Top-2 accuracy	~99%	Correct cell is almost always in the top 2 predictions
Average inference latency	~2-4ms	Per prediction on CPU — well within near-RT budget
Model file size	~500KB	Small enough to load quickly at startup

The slight gap between training accuracy (~98%) and validation accuracy (~94-96%) is acceptable and expected. It indicates mild overfitting that the dropout layer partially mitigates. The 96.76% accuracy observed in sim011 matches the validation set performance, confirming the model generalizes well.

Experimental Design — How the Simulations Were Structured

N12.1 Why Three Simulations?

We ran three distinct simulations (sim006, sim010, sim011) to provide evidence that our results are not simply a lucky artifact of one particular random seed. Each simulation uses a different NS-3 random seed (RngRun parameter), which causes the UE mobility patterns to differ. By showing consistent low PP rates across three independent runs with different UE trajectories, we provide stronger evidence of the system's robustness.

Simulation	RngRun seed	simTime	UE movement pattern	Result
sim006	6	60s	UEs traversed many cell boundaries (334 HOs)	PP rate: 1.72%
sim010	10	60s	UEs spent more time in cell centers (137 HOs)	PP rate: 3.65%
sim011	11	120s	Mixed patterns, longer duration (309 HOs)	PP rate: 3.24%

N12.2 What 'RngRun' Controls

NS-3 uses a pseudo-random number generator (PRNG) to produce all random elements in the simulation: UE initial positions, movement directions, movement speed variations, and shadow fading values. The **RngRun** parameter seeds this PRNG. Setting RngRun = 6 always produces exactly the same sequence of random numbers — making the simulation perfectly repeatable. If you run sim006 twice with RngRun=6, you get identical results. Different seeds produce different UE trajectories, different channel realizations, and different handover patterns — but the system parameters (thresholds, cooldown, GRU model) remain constant.

Note: Reproducibility note: To exactly reproduce sim011, run: `.ns3 run gru_scenario --simTime=120 --RngRun=11` with the same `gru_model.h5` and `best2.c` constants as documented in this thesis.

N12.3 Controlled Variables

In a rigorous experiment, we fix all variables except the one being tested. The controlled (constant) variables across all three simulations:

Variable	Fixed Value	Why Fixed
A3_OFFSET	1.0 dB	Consistent handover sensitivity across runs
A3_HYSTERESIS	1.0 dB	Consistent anti-oscillation margin
COOLDOWN_TIME	5.0 sim-seconds	Consistent anti-ping-pong protection

Variable	Fixed Value	Why Fixed
GRU model weights	Same gru_model.h5	Same ML decision quality across runs
GRU window size	10 samples	Same prediction horizon
KPM interval	0.05 sim-seconds	Same measurement frequency
Number of UEs	20	Same system load
Number of cells	7 gNBs	Same cell layout
Carrier frequency	28 GHz	Same channel characteristics

The only variable that differs between sim006, sim010, and sim011 is the NS-3 random seed (and sim011 has twice the simulation duration). This controlled design means that differences in results (1.72% vs 3.24% vs 3.65% PP rate) reflect the natural variance introduced by different UE mobility patterns — not different system configurations.

N12.4 Limitations of the Experimental Design

An intellectually honest evaluation of our experimental design must acknowledge its limitations:

- **Small sample size:** Three simulation runs is a minimal dataset for statistical analysis. Ideally, 20-50 runs with different seeds would allow computing confidence intervals and testing statistical significance. Our results should be interpreted as preliminary evidence, not definitive proof.
- **No baseline comparison:** We do not have a control run with pure A3 (no GRU) under identical conditions. Our comparison to literature values (8-15% PP rate for A3-only) is valid but less rigorous than an in-system comparison. A matched baseline run would definitively quantify the GRU's contribution.
- **Single mobility model:** All three runs use Random Walk 2D. Real users follow streets, avoid buildings, and have purposeful destinations. Results may differ with more realistic mobility models like SUMO or Gauss-Markov.
- **Single cell layout:** Our 7-cell layout is fixed. Different cell densities, different inter-cell distances, or different numbers of cells would change the handover dynamics. The GRU model would need retraining for a different layout.

These limitations are honest and should be included in your thesis discussion section. Committees respect students who can critically analyze their own work's limitations rather than overselling the results.

N12.5 How to Reproduce Any Simulation Run

The complete reproduction procedure for sim011:

1. Ensure all prerequisites are installed: NS-3 mmWave module, FlexRIC, Python 3.8+, TensorFlow 2.x, Flask, Docker.
2. Compile the xApp: `cd mmwave-LENA-oran && make xapp_handover_gru`
3. Compile NS-3 with gru_scenario: `./ns3 build`
4. Ensure gru_model.h5 and its normalization parameters are in the mmwave-LENA-oran/ directory.

-
5. Run: `bash gru.sh --simTime=120 --RngRun=11`
 6. Wait approximately 6-7 hours for the simulation to complete.
 7. Run: `bash kill_sim.sh` to collect results.
 8. Results appear in: `3D_GUI_Sim_Results/sim011_*/`

Defense Preparation — How to Present Your Results

N13.1 The Opening Minute of Your Defense

The first minute of your defense sets the tone. Your opening should answer three questions in simple terms: What problem did you solve? How did you solve it? What did you achieve? Here is a template for your opening:

***Note:** 'Handover management in 5G millimeter-wave networks is challenging because UEs frequently move between cell coverage areas and traditional threshold-based algorithms produce ping-pong effects — unnecessary back-and-forth handovers that waste network resources. My thesis proposes integrating a GRU neural network into the near-RT RIC handover decision process. The GRU processes a 10-sample SINR history to predict the best target cell before committing to a handover. We implemented this in a real O-RAN framework (FlexRIC + NS-3) and achieved a ping-pong rate of 3.24% — compared to 8-15% for traditional A3-only approaches in the literature.'*

Practice this opening until you can say it in under 90 seconds without notes. This single paragraph covers: the problem, your solution, the technical approach, the platform used, and the key result. It answers what every committee member wants to know first.

N13.2 Explaining the Results (When Asked)

When a committee member asks 'tell me about your results,' use this structure:

- 1. State the headline number:** 'We ran three simulations and achieved ping-pong rates between 1.72% and 3.65%, with 96.76% accuracy in our longest run (sim011).'
- 2. Contextualize with literature:** 'The literature reports 8-15% PP rate for traditional A3-only approaches in mmWave systems, so our system represents a 60-80% reduction.'
- 3. Explain the variance:** 'The variation between simulations is due to different random seeds — different UE movement patterns naturally affect how often UEs are at cell boundaries.'
- 4. Acknowledge limitations:** 'Three simulation runs is a starting point — a production evaluation would require 20+ runs for statistical significance.'

N13.3 How to Handle a Question You Don't Know the Answer To

Every defense includes at least one question where you are uncertain. The worst response is to make something up confidently. The best responses:

- **Partial knowledge:** 'I'm not certain of the exact value, but I know the general principle is... and it works in our system because...'

- **Honest admission:** 'That's an interesting limitation I haven't analyzed in detail. My initial thinking is... but this would be worth investigating further.'
- **Redirect to what you know:** 'I don't have the exact specification in mind, but in our implementation we handle this case by... which achieves the goal because...'

Never say 'I don't know' without following up with what you DO know. The committee wants to see how you think under pressure, not just whether you memorized every fact.

N13.4 Key Diagrams to Have Ready

If your presentation includes slides, these are the most important diagrams to prepare:

Diagram	What It Should Show	Why Important
O-RAN Architecture	Non-RT RIC → Near-RT RIC → E2 Node stack with interfaces labeled	Establishes the context for your work
System Component Diagram	NS-3 ↔ FlexRIC ↔ xApp ↔ GRU service ↔ InfluxDB ↔ Grafana	Shows the full technical stack you built
E2 Message Flow	Sequence diagram: NS-3 → E2 Setup → KPM Subscription → KPM Reports → RC Command	Demonstrates E2 protocol understanding
GRU Architecture	Input (10, 11) → GRU(64) → Dense(7) → Softmax → 7 cell probs	Shows ML model design
Ping-Pong Timeline	Timeline showing HO at T, blocked HO at T+3s (cooldown), allowed HO at T+6s	Explains cooldown mechanism visually
Results Table/Chart	Bar chart comparing 3 simulations: PP rate and accuracy side by side	The key result of your thesis

N13.5 Anticipated Tough Questions and Model Answers

'Your PP rate is 3.24%. Is this actually better than just using a longer TTT?'

A longer Time-to-Trigger (TTT) also reduces ping-pong by requiring the A3 condition to be sustained for longer before firing. However, a longer TTT introduces delay — the handover happens later, meaning the UE may experience degraded SINR while waiting. Our GRU approach does not delay the handover — it happens at the first A3 trigger, but the GRU validates that the improvement is sustainable before committing. This gives low PP rate WITHOUT the latency penalty of TTT extension.

'Your system is tested only in simulation. How do you know it works in reality?'

Simulation is a necessary first step in network research — testing unproven algorithms in production networks risks disrupting real users. The NS-3 mmWave module uses the 3GPP TR 38.901 channel model, which is the same model used by equipment vendors for 5G NR design. FlexRIC implements the actual O-RAN E2AP protocol, not a simplified mock. The path to real deployment would involve: field trials with a small number of test UEs, retraining the GRU on real SINR measurements, and gradual rollout with the pure A3 fallback as a

safety net.

'What is the computational overhead of your GRU service?'

The GRU inference latency on a standard Intel Core i7 CPU is approximately 2-4ms per prediction request. The xApp makes at most one GRU call per UE per A3 event — not on every KPM report. In sim011 with 309 HOs over 120 sim-seconds, the total GRU calls were at most 309 (one per HO decision). The CPU overhead of gru_xapp.py is approximately 5-10% during peak handover periods — negligible compared to NS-3's 100% CPU usage for channel computation.

N13.6 The Poster Defense Alternative

Some faculties conduct poster defenses. If yours does, adapt the structure:

- **Top section of poster:** Title, abstract, O-RAN architecture diagram
- **Left column:** Problem statement + literature review table
- **Center column:** System architecture diagram + E2 message flow
- **Right column:** GRU model diagram + results table + conclusions
- **Bottom:** References and acknowledgments

For a poster defense, prepare a 3-minute verbal walk-through that covers the problem, solution, and key result. This mirrors the oral defense opening described in Section N13.1 above.

QUICK REFERENCE — All Key Facts at a Glance

Numbers You Must Know By Heart

Parameter	Value	Remember Because
Cells in simulation	7 gNBs	One central + 6 surrounding
UEs in simulation	20	All random-walk pedestrians
Carrier frequency	28 GHz	mmWave band
KPM report interval	0.05 sim-seconds	= 20 reports per sim-second
GRU window size	10 samples	= 0.5 sim-seconds of history
A3 threshold total	2.0 dB	= 1.0 (offset) + 1.0 (hysteresis)
Cooldown time	5.0 sim-seconds	Per UE after any handover
GRU service port	5000	Flask default, localhost only
E2 interface port	36421 SCTP	Standard O-RAN port
FlexRIC log location	/tmp/flexric.log	ALWAYS here, never stdout
2D GUI port	8000	nginx → Grafana
3D GUI port	3001 + 8001	Vite frontend + FastAPI
Sim pace	0.005 sim-s/real-s	60 sim-s ≈ 3 real hours

Simulation Results Summary

Sim	Time	HOs	PPs	PP Rate	Accuracy	Note
sim006	60s	334	4	1.72%	98.28%	Most HOs (many boundary crossings)
sim010	60s	137	5	3.65%	96.35%	Fewest HOs (UEs stayed in centers)
sim011	120s	309	10	3.24%	96.76%	Longest run — most reliable rate

Key Formulas

A3 trigger: $\text{neighbor_SINR} > \text{serving_SINR} + \text{A3_OFFSET}(1.0) + \text{A3_HYSTERESIS}(1.0) = +2.0 \text{ dB}$

Cooldown: $(\text{current_sim_time} - \text{last_HO_time}[\text{ue}]) > 5.0 \text{ sim-seconds}$

Accuracy = $(\text{Total_HOs} - \text{Ping-Pong_HOs}) / \text{Total_HOs} \times 100\%$

sim011 accuracy = $(309 - 10) / 309 \times 100\% = 96.76\%$

Sim pace: 60 sim-s / 0.005 = 12,000 real-s $\approx$ 3.3 hours

Component Responsibilities — Who Does What

Component	File	Role
NS-3 simulation	gru_scenario.cc	Virtual network: 7 cells, 20 UEs, mmWave channel
FlexRIC RIC	nearRT-RIC binary	O-RAN near-RT RIC: routes E2 messages between NS-3 and xApp
xApp	best2.c → xapp_handover_gru	Handover logic: A3 check, GRU call, RC command
GRU service	gru_xapp.py	ML inference: receives window, returns best cell_id
InfluxDB	Docker container	Time-series database for visualization data
Grafana	Docker container	Dashboard visualization at port 8000
sim_data_pusher	sim_data_pusher.py	Reads NS-3 output, pushes to InfluxDB
FastAPI controller	controller.py	3D GUI backend API at port 8001
Vite frontend	3D GUI source	3D visualization at port 3001
Master launcher	gru.sh	Starts all components in correct order

Good luck with your defense, Omar! You built this entire system — you know it better than anyone in that room.
